

RISC-V "V" Vector Extension

Version 1.0-draft-20200703

Table of Contents

Changes from v0.9

SLLEN=VLEN layout mandatory.

Support ELEN > VLEN for LMUL > 1

Defined vector FP exception behavior

Defined interaction of `misa.v` and `mstatus.vs`

Defined integer narrowing pseudo-instruction `vncvt.x.x.v vd,vs,vm`

Added reciprocal and reciprocal square-root estimate instructions

Added EEW encoding to whole register moves and load/stores to support microarchitectures with internal data rearrangement.

Added `vrgatherei16` instruction

Rearranged bits in `vtype` to make `vlmul` bits into a contiguous field

Moved EDIV to appendix and removed instruction encoding for dot instructions to make clear not part of v1.0

Moved quad-widening `mulacc` to appendix and removed instruction encodings to make clear not part of v1.0

Added `vneg` and `vneg.pseudo`-instructions.

Changed mandatory illegal instruction exceptions on unsupported encodings to reserved, to allow future use of these encodings. Implementations are still recommended to trap on unsupported reserved encodings.

Clarified that support for vector misaligned memory accesses is independent of support for scalar misaligned memory accesses.

Clarified support for precise/imprecise traps. Require `vstart` to report location of faulting element.

Define whole register move instructions in terms of SEW instead of EEW=8 to allow definition to apply to machines with no support for EEW=8, and to help describe how these provide a hint to implementations with internal rearrangement of data. Functionality has not changed.

Specified minimum VLEN=128 for the base "V" vector extension.

Clarified whole register moves, including existence of register group save/restore, and allowed misaligned to be signaled for smallest supported size instead of encoded EEW to support implementations without SEW=8 load/store.

Changed default for vector memory instructions to be "unordered", and added ordered/unordered vector indexed instructions.

Clarified that for precise traps, vector store instructions can update idempotent memory for elements past the location of an element raising a trap.

Reserved `vstart` values greater than maximum element index.

Allow vector strided instruction to perform fewer memory accesses when `rs2=x0`.

Clarified that `vsetvl{i}_x0, x0, ...` form instructions may set `vll` when new SEW/LMUL ratio would change VLMAX.

Made clear that implementations must support LMUL 1,2,4,8.

1. Introduction

2. Implementation-defined Constant Parameters

3. Vector Extension Programmer's Model

3.1. Vector Registers

3.2. Vector Context Status in `mstatus`

3.3. Vector type register, `vtype`

3.4. Vector Length Register `vl`

3.5. Vector Byte Length `vlmb`

3.6. Vector Start Index CSR `vstart`

3.7. Vector Fixed-Point Rounding Mode Register `vxrm`

3.8. Vector Fixed-Point Saturation Flag `vxsat`

3.9. Vector Control and Status Register `vcsr`

3.10. State of Vector Extension at Reset

4. Mapping of Vector Elements to Vector Register State

- [4.1. Mapping for LMUL ≤ 1](#)
- [4.2. Mapping with LMUL > 1](#)
- [4.3. Mapping across Mixed-Width Operations](#)
- [4.4. Mapping with LMUL > 1 and ELEN > VLEN](#)
- [4.5. Mask Register Layout](#)
- [5. Vector Instruction Formats](#)
 - [5.1. Scalar Operands](#)
 - [5.2. Vector Operands](#)
 - [5.3. Vector Masking](#)
 - [5.4. Prestart, Active, Inactive, Body, and Tail Element Definitions](#)
- [6. Configuration-Setting Instructions](#)
 - [6.1. vsetvli/vsetvl instructions](#)
 - [6.2. Constraints on Setting vl](#)
 - [6.3. Example of stripmining and changes to SEW](#)
- [7. Vector Loads and Stores](#)
 - [7.1. Vector Load/Store Instruction Encoding](#)
 - [7.2. Vector Load/Store Addressing Modes](#)
 - [7.3. Vector Load/Store Width Encoding](#)
 - [7.4. Vector Unit-Stride Instructions](#)
 - [7.5. Vector Strided Instructions](#)
 - [7.6. Vector Indexed Instructions](#)
 - [7.7. Unit-stride Fault-Only-First Loads](#)
 - [7.8. Vector Load/Store Segment Instructions](#)
 - [7.9. Vector Load/Store Whole Register Instructions](#)
- [8. Vector AMO Operations](#)
- [9. Vector Memory Alignment Constraints](#)
- [10. Vector Memory Consistency Model](#)
- [11. Vector Arithmetic Instruction Formats](#)
 - [11.1. Vector Arithmetic Instruction encoding](#)
 - [11.2. Widening Vector Arithmetic Instructions](#)
 - [11.3. Narrowing Vector Arithmetic Instructions](#)
- [12. Vector Integer Arithmetic Instructions](#)
 - [12.1. Vector Single-Width Integer Add and Subtract](#)
 - [12.2. Vector Widening Integer Add/Subtract](#)
 - [12.3. Vector Integer Extension](#)
 - [12.4. Vector Integer Add-with-Carry / Subtract-with-Borrow Instructions](#)
 - [12.5. Vector Bitwise Logical Instructions](#)
 - [12.6. Vector Single-Width Bit Shift Instructions](#)
 - [12.7. Vector Narrowing Integer Right Shift Instructions](#)
 - [12.8. Vector Integer Comparison Instructions](#)
 - [12.9. Vector Integer Min/Max Instructions](#)
 - [12.10. Vector Single-Width Integer Multiply Instructions](#)
 - [12.11. Vector Integer Divide Instructions](#)
 - [12.12. Vector Widening Integer Multiply Instructions](#)
 - [12.13. Vector Single-Width Integer Multiply-Add Instructions](#)
 - [12.14. Vector Widening Integer Multiply-Add Instructions](#)
 - [12.15. Vector Integer Merge Instructions](#)
 - [12.16. Vector Integer Move Instructions](#)
- [13. Vector Fixed-Point Arithmetic Instructions](#)
 - [13.1. Vector Single-Width Saturating Add and Subtract](#)

- [13.2. Vector Single-Width Averaging Add and Subtract](#)
- [13.3. Vector Single-Width Fractional Multiply with Rounding and Saturation](#)
- [13.4. Vector Single-Width Scaling Shift Instructions](#)
- [13.5. Vector Narrowing Fixed-Point Clip Instructions](#)
- [14. Vector Floating-Point Instructions](#)
 - [14.1. Vector Floating-Point Exception Flags](#)
 - [14.2. Vector Single-Width Floating-Point Add/Subtract Instructions](#)
 - [14.3. Vector Widening Floating-Point Add/Subtract Instructions](#)
 - [14.4. Vector Single-Width Floating-Point Multiply/Divide Instructions](#)
 - [14.5. Vector Widening Floating-Point Multiply](#)
 - [14.6. Vector Single-Width Floating-Point Fused Multiply-Add Instructions](#)
 - [14.7. Vector Widening Floating-Point Fused Multiply-Add Instructions](#)
 - [14.8. Vector Floating-Point Square-Root Instruction](#)
 - [14.9. Vector Floating-Point Reciprocal Square-Root Estimate Instruction](#)
 - [14.10. Vector Floating-Point Reciprocal Estimate Instruction](#)
 - [14.11. Vector Floating-Point MIN/MAX Instructions](#)
 - [14.12. Vector Floating-Point Sign-Injection Instructions](#)
 - [14.13. Vector Floating-Point Compare Instructions](#)
 - [14.14. Vector Floating-Point Classify Instruction](#)
 - [14.15. Vector Floating-Point Merge Instruction](#)
 - [14.16. Vector Floating-Point Move Instruction](#)
 - [14.17. Single-Width Floating-Point/Integer Type-Convert Instructions](#)
 - [14.18. Widening Floating-Point/Integer Type-Convert Instructions](#)
 - [14.19. Narrowing Floating-Point/Integer Type-Convert Instructions](#)
- [15. Vector Reduction Operations](#)
 - [15.1. Vector Single-Width Integer Reduction Instructions](#)
 - [15.2. Vector Widening Integer Reduction Instructions](#)
 - [15.3. Vector Single-Width Floating-Point Reduction Instructions](#)
 - [15.4. Vector Widening Floating-Point Reduction Instructions](#)
- [16. Vector Mask Instructions](#)
 - [16.1. Vector Mask-Register Logical Instructions](#)
 - [16.2. Vector mask population count vpopc](#)
 - [16.3. vfirst find-first-set mask bit](#)
 - [16.4. vmsbf .m set-before-first mask bit](#)
 - [16.5. vmsif .m set-including-first mask bit](#)
 - [16.6. vmsof .m set-only-first mask bit](#)
 - [16.7. Example using vector mask instructions](#)
 - [16.8. Vector Iota Instruction](#)
 - [16.9. Vector Element Index Instruction](#)
- [17. Vector Permutation Instructions](#)
 - [17.1. Integer Scalar Move Instructions](#)
 - [17.2. Floating-Point Scalar Move Instructions](#)
 - [17.3. Vector Slide Instructions](#)
 - [17.4. Vector Register Gather Instructions](#)
 - [17.5. Vector Compress Instruction](#)
 - [17.6. Whole Vector Register Move](#)
- [18. Exception Handling](#)
 - [18.1. Precise vector traps](#)
 - [18.2. Imprecise vector traps](#)
 - [18.3. Selectable precise/imprecise traps](#)

18.4. Swappable traps

19. Vector Instruction Listing

Appendix A: Vector Assembly Code Examples

A.1. Vector-vector add example

A.2. Example with mixed-width mask and compute.

A.3. Malloc example

A.4. Conditional example

A.5. SAXPY example

A.6. SGEMM example

A.7. Division approximation example

A.8. Square root approximation example

Appendix B: Calling Convention

Appendix C: Vector Quad-Widening Integer Multiply-Add Instructions (Extension Zvqmac).

Appendix D: Divided Element Extension (Extension Zvediv).

D.1. Instructions not affected by EDIV

D.2. Instructions Affected by EDIV

D.3. Vector Integer Dot-Product Instruction

D.4. Vector Floating-Point Dot Product Instruction

Contributors include: Alon Amid, Krste Asanovic, Allen Baum, Alex Bradbury, Tony Brewer, Chris Celio, Aliaksei Chapyzhenka, Silviu Chiricescu, Ken Dockser, Bob Dreyer, Roger Espasa, Sean Halle, John Hauser, David Horner, Bruce Houtt, Bill Huffman, Nicholas Knight, Constantine Korikov, Ben Korpan, Hanna Kruppe, Yunsup Lee, Guy Lemieux, Grigorios Magklis, Filip Moc, Rich Newell, Albert Ou, David Patterson, Colin Schmidt, Alex Solomatnikov, Steve Wallach, Andrew Waterman, Jim Wilson.

Known issues with current version:

- encoding needs better formatting
- vector memory consistency model needs to be clarified
- interaction with privileged architectures

Changes from v0.9

SLEN=VLEN layout mandatory

The group has decided to make the SLEN=VLEN layout mandatory. In-register layout of the bytes of a vector matches in-memory layout of bytes in a vector. Many of the optimizations possible with the earlier SLEN<VLEN layouts can be achieved with microarchitectural techniques on wide datapath machines, and SLEN=VLEN provides a much simpler specification and interface to software.

Support ELEN > VLEN for LMUL > 1

Specification was loosened to allow elements wider than a single vector register to be supported using a vector register group, but profiles can still mandate a minimum ELEN when LMUL = 1.

Defined vector FP exception behavior

Defined interaction of `misa.v` and `mstatus.vs`

Defined integer narrowing pseudo-instruction `vncvt.x.x.v vd, vs, vm`

Added reciprocal and reciprocal square-root estimate instructions

Added EEW encoding to whole register moves and load/stores to support microarchitectures with internal data rearrangement.

Added `vrgatheri16` instruction

Rearranged bits in `vtype` to make `vlmul` bits into a contiguous field

Moved EDIV to appendix and removed instruction encoding for dot instructions to make clear not part of v1.0

Moved quad-widening mulacc to appendix and removed instruction encodings to make clear not part of v1.0

Added `vneg` and `vfneg` pseudo-instructions.

Changed mandatory illegal instruction exceptions on unsupported encodings to reserved, to allow future use of these encodings. Implementations are still recommended to trap on unsupported reserved encodings.

Clarified that support for vector misaligned memory accesses is independent of support for scalar misaligned memory accesses.

Clarified support for precise/imprecise traps. Require `vstart` to report location of faulting element.

Define whole register move instructions in terms of SEW instead of EEW=8 to allow definition to apply to machines with no support for EEW=8, and to help describe how these provide a hint to implementations with internal rearrangement of data. Functionality has not changed.

Specified minimum VLEN=128 for the base "V" vector extension.

Clarified whole register moves, including existence of register group save/restore, and allowed misaligned to be signaled for smallest supported size instead of encoded EEW to support implementations without

SEW=8 load/store.

Changed default for vector memory instructions to be "unordered", and added ordered/unordered vector indexed instructions.

Clarified that for precise traps, vector store instructions can update idempotent memory for elements past the location of an element raising a trap.

Reserved vstart values greater than maximum element index.

Allow vector strided instruction to perform fewer memory accesses when $rs2=x0$.

Clarified that `vsetvl{i} x0, x0, ...` form instructions may set `vll` when new SEW/LMUL ratio would change VLMAX.

Made clear that implementations must support LMUL 1,2,4,8.

1. Introduction

This document describes the draft of version 1.0 of the RISC-V vector extension.

This is a draft of the stable proposal for the vector extension specification to be used for implementation and evaluation. Once the draft label is removed, version 1.0 is intended to be sent out for public review as part of the RISC-V International ratification process. Version 1.0 is also considered stable enough to begin developing toolchains, functional simulators, and initial implementations, including in upstream software projects, and is not expected to have major functionality changes except if serious issues are discovered during ratification. Once ratified, the spec will be given version 2.0.

This draft spec is intended to capture how the complete set of currently defined vector instructions, but is not intended to determine what set of vector instructions and which supported element widths are mandatory for a given platform profile.

The term *base vector extension* is used informally to describe the standard set of vector ISA components that will be required for the single-letter "V" extension, which is intended for use in standard server and application-processor platform profiles. The set of mandatory instructions and supported element widths will vary with the base ISA (RV32I, RV64I) as described below.

Other profiles, including embedded profiles, may choose to mandate only subsets of these extensions. The exact set of mandatory supported instructions for an implementation to be compliant with a given profile will only be determined when each profile spec is ratified. For convenience in defining subset profiles, vector instruction subsets are given ISA string names beginning with the "Zv" prefix.

The document describes all the individual features to be included in the base vector extension.

The set of instructions to be included or not in the base "V" extension, and the naming of all the vector instruction subsets and extensions is still under review in this draft.

The base vector extension is designed to act as a base for additional vector extensions in various domains, including cryptography and machine learning.

2. Implementation-defined Constant Parameters

Each hart supporting the vector extension defines two parameters:

1. The maximum size of a vector element that any operation can produce or consume in bits, $ELEN \geq 8$, which must be a power of 2.
2. The number of bits in a single vector register, $VLEN$, which must be a power of 2.

Profiles may set further constraints on these parameters, for example, requiring that $ELEN \geq \max(XLEN, FLEN)$, or requiring a minimum $VLEN$ value.

The base "V" vector extension requires that $VLEN \geq 128$.

The value of 128 was chosen as a compromise for application processors. Providing a larger $VLEN$ allows stripmining code to be elided in some cases for short vectors, but also increases the size of the minimum implementation. Note that larger $LMUL$ can be used to avoid stripmining for longer known-size application vectors at the cost of having fewer available vector register groups. For example, an $LMUL$ of 8 allows vectors of up to sixteen 64-bit elements to be processed without stripmining using four vector register groups.

The ISA supports writing binary code that under certain constraints will execute portably on harts with different values for these parameters.

Code can be written that will expose differences in implementation parameters.

In general, thread contexts with active vector state cannot be migrated during execution between harts that have any difference in $VLEN$ or $ELEN$ parameters.

3. Vector Extension Programmer's Model

The vector extension adds 32 vector registers, and seven unprivileged CSRs (`vstart`, `vxsat`, `vxrm`, `vcsr`, `vtype`, `vl`, `vlenb`) to a base scalar RISC-V ISA.

Table 1. New vector CSRs

Address	Privilege	Name	Description
0x008	URW	<code>vstart</code>	Vector start position
0x009	URW	<code>vxsat</code>	Fixed-Point Saturate Flag
0x00A	URW	<code>vxrm</code>	Fixed-Point Rounding Mode
0x00F	URW	<code>vcsr</code>	Vector control and status register
0xC20	URO	<code>vl</code>	Vector length
0xC21	URO	<code>vtype</code>	Vector data type register
0xC22	URO	<code>vlenb</code>	VLEN/8 (vector register length in bytes)

3.1. Vector Registers

The vector extension adds 32 architectural vector registers, `v0-v31` to the base scalar RISC-V ISA.

Each vector register has a fixed VLEN bits of state.

Zfinx ("F in X") is a new ISA option under consideration where floating-point instructions take their arguments from the integer register file. The 1.0 vector extension is also compatible with Zfinx.

3.2. Vector Context Status in `mstatus`

A vector context status field, VS, is added to `mstatus[10:9]` and shadowed in `sstatus[10:9]`. It is defined analogously to the floating-point context status field, FS.

Attempts to execute any vector instruction, or to access the vector CSRs, raise an illegal-instruction exception when the VS field is set to Off.

When the VS field is set to Initial or Clean, executing any instruction that changes vector state, including the vector CSRs, will change VS to Dirty.

Implementations may also change VS field to Dirty at any time, even when there is no change in vector state. Accurate setting of the VS field is an optimization.

Implementations may have a writable `misa.v` field. Analogous to the way in which the floating-point unit is handled, the `mstatus.vs` field may exist even if `misa.v` is clear.

Allowing `mstatus.vs` to exist when `misa.v` is clear, enables vector emulation and simplifies handling of `mstatus.vs` in systems with writable `misa.v`.

3.3. Vector type register, `vtype`

The read-only XLEN-wide *vector type* CSR, `vtype` provides the default type used to interpret the contents of the vector register file, and can only be updated by `vsetvl{i}` instructions. The vector type also determines the organization of elements in each vector register, and how multiple vector registers are grouped.

Earlier drafts allowed the `vtype` register to be written using regular CSR writes. Allowing updates only via the `vsetvl{i}` instructions simplifies maintenance of the `vtype` register state.

In the base vector extension, the `vtype` register has five fields, `vill`, `vma`, `vta`, `vsew[2:0]`, and `vlmul[2:0]`.

Table 2. vtype register layout

Bits	Name	Description
XLEN-1	vill	Illegal value if set
XLEN-2:8		Reserved (write 0)
7	vma	Vector mask agnostic
6	vta	Vector tail agnostic
5:3	vsew[2:0]	Selected element width (SEW) setting
2:0	vlmul[2:0]	Vector register group multiplier (LMUL) setting

A small implementation supporting ELEN=32 requires only seven bits of state in vtype: two bits for ma and ta, two bits for vsew[1:0] and three bits for vmul[2:0]. The illegal value represented by vill can be internally encoded using the illegal 64-bit combination in vsew[1:0] without requiring an additional storage bit to hold vill.

Further standard and custom extensions to the vector base will extend these fields to support a greater variety of data types.

It is anticipated that an extended 64-bit instruction encoding would allow these fields to be specified statically in the instruction encoding.

3.3.1. Vector selected element width vsew[2:0]

The value in vsew sets the dynamic *selected element width* (SEW). By default, a vector register is viewed as being divided into VLEN/SEW elements.

In the base vector "V" extension, only SEW up to ELEN = max(XLEN,FLEN) are required to be supported. Other profiles may impose different constraints on ELEN.

Table 3. vsew[2:0] (selected element width) encoding

vsew[2:0]			SEW
0	0	0	8
0	0	1	16
0	1	0	32
0	1	1	64
1	0	0	128
1	0	1	256
1	1	0	512
1	1	1	1024

Table 4. Example VLEN = 128 bits

SEW	Elements per vector register
64	2
32	4
16	8
8	16

The supported element width may vary with LMUL, but profiles may mandate the minimum SEW that must be supported with LMUL=1.

The standard base V vector extension requires that SEW=max(XLEN,FLEN) is supported with LMUL=1.

Some implementations may support larger SEWs only when bits from multiple vector registers are combined. Software that relies on large SEW should attempt to use the largest LMUL, and hence the fewest vector register groups, to increase the number of implementations on which the code will run. The vill bit in vtype should be checked after setting vtype to see if the configuration is supported, and an alternate code path should be provided if it is not. Alternatively, a profile can mandate the minimum SEW at each LMUL setting.

3.3.2. Vector Register Grouping (vlmul[2:0])

Multiple vector registers can be grouped together, so that a single vector instruction can operate on multiple vector registers. The term *vector register group* is used herein to refer to one or more vector registers used as a single operand to a vector

instruction. Vector register groups allow double-width or larger elements to be operated on with the same vector length as selected-width elements. Vector register groups also provide greater execution efficiency for longer application vectors.

The vector length multiplier, $LMUL$, when greater than 1, represents the default number of vector registers that are combined to form a vector register group. Implementations must support $LMUL$ integer values of 1,2,4,8.

$LMUL$ can also be a fractional value, reducing the number of bits used in a vector register. $LMUL$ can have fractional values 1/2, 1/4, 1/8. Fractional $LMUL$ is used to increase the number of usable architectural registers when operating on mixed-width values, by not requiring that larger-width vectors occupy multiple vector registers. Instead, wider values can occupy a single vector register and narrower values can occupy a fraction of a vector register.

Implementations must support fractional $LMUL$ settings for $LMUL \geq SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$, where $SEW_{LMUL1MIN}$ is the narrowest supported SEW value at $LMUL=1$ and $SEW_{LMUL1MAX}$ is the widest supported SEW value at $LMUL=1$. An attempt to set an unsupported SEW and $LMUL$ configuration sets the `vill` bit in `vtype`.

Requiring $LMUL \geq SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$ allows software operating on mixed-width elements to only use a single vector register to hold the wider elements, with fractional $LMUL$ used to hold narrower elements. When $LMUL < SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$, there is no guarantee an implementation would have enough bits in the fractional vector register to store at least one element, as $VLEN=SEW_{LMUL1MAX}$ is a valid implementation choice.

The constraint is written using $SEW_{LMUL1MAX}$ and not $ELEN$ because some systems might only support larger SEW values for $LMUL>1$. Note that in these cases, the constraint ensures that no more than a single vector register is needed to hold the widest-supported element that can be held in a single vector register, when code is also performing operations on narrower widths.

The use of `vtype` encodings with $LMUL < SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$ is *reserved*, but implementations can set `vill` if they do not support these configurations.

Requiring all implementations to set `vill` in this case would prohibit future use of this case in an extension, so to allow for a future definition of $LMUL < SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$ behavior, we consider the use of this case to be *reserved*.

It is recommended that assemblers provide a warning (not an error) if a `vsetvli` instruction attempts to write an $LMUL < SEW_{LMUL1MIN}/SEW_{LMUL1MAX}$.

$LMUL$ is set by the signed `vlmul` field in `vtype` ($LMUL = 2^{v1mul[2:0]}$).

The derived value $VLMAX = LMUL * VLEN / SEW$ represents the maximum number of elements that can be operated on with a single vector instruction given the current SEW and $LMUL$ settings as shown in the table below.

v1mul			LMUL	#groups	VLMAX	Registers grouped with register <i>n</i>
1	0	0	-	-	-	reserved
1	0	1	1/8	32	$VLEN/SEW/8$	$v\ n$ (single register in group)
1	1	0	1/4	32	$VLEN/SEW/4$	$v\ n$ (single register in group)
1	1	1	1/2	32	$VLEN/SEW/2$	$v\ n$ (single register in group)
0	0	0	1	32	$VLEN/SEW$	$v\ n$ (single register in group)
0	0	1	2	16	$2 * VLEN/SEW$	$v\ n, v\ n+1$
0	1	0	4	8	$4 * VLEN/SEW$	$v\ n, \dots, v\ n+3$
0	1	1	8	4	$8 * VLEN/SEW$	$v\ n, \dots, v\ n+7$

When $LMUL=2$, the vector register group contains vector register $v\ n$ and vector register $v\ n+1$, providing twice the vector length in bits. Instructions specifying an $LMUL=2$ vector register group with an odd-numbered vector register are reserved.

When $LMUL=4$, the vector register group contains four vector registers, and instructions specifying an $LMUL=4$ vector register group using vector register numbers that are not multiples of four are reserved.

When $LMUL=8$, the vector register group contains eight vector registers, and instructions specifying an $LMUL=8$ vector register group using register numbers that are not multiples of eight are reserved.

Mask registers are always contained in a single vector register, regardless of $LMUL$.

3.3.3. Vector Tail Agnostic and Vector Mask Agnostic `vta` and `vma`

These two bits modify the behavior of destination tail elements and destination inactive masked-off elements respectively during the execution of vector instructions. The tail and inactive sets contain element positions that are not receiving new results during a vector operation, as defined in Section [Prestart, Active, Inactive, Body, and Tail Element Definitions](#).

All systems must support all four options:

vta	vma	Tail Elements	Inactive Elements
0	0	undisturbed	undisturbed
0	1	undisturbed	agnostic
1	0	agnostic	undisturbed
1	1	agnostic	agnostic

When a set is marked undisturbed, the corresponding set of destination elements in any vector or mask destination operand retain the value they previously held.

When a set is marked agnostic, the corresponding set of destination elements in any vector or mask destination operand can either retain the value they previously held, or are overwritten with 1s. Within a single vector instruction, each destination element can be either left undisturbed or overwritten with 1s, in any combination, and the pattern of undisturbed or overwritten with 1s is not required to be deterministic when the instruction is executed with the same inputs.

The agnostic policy was added to accommodate machines with vector register renaming, and/or that have deeply temporal vector registers. With an undisturbed policy, all elements would have to be read from the old physical destination vector register to be copied into the new physical destination vector register. This causes an inefficiency when these inactive or tail values are not required for subsequent calculations.

The intent is for software to select the option that reduces microarchitectural work by selecting agnostic when the value in the respective set does not matter.

The value of all 1s instead of all 0s was chosen for the overwrite value to discourage software developers from depending on the value written.

A simple in-order implementation can ignore the setting and simply execute all vector instructions using the undisturbed policy. The vta and vma state bits must still be provided in vtype for compatibility and to support thread migration.

An out-of-order implementation can choose to implement tail-agnostic + mask-agnostic using tail-agnostic + mask-undisturbed to reduce implementation complexity.

The definition of agnostic result policy is left loose to accommodate migrating application threads between harts on a small in-order core (which probably leaves agnostic regions undisturbed) and harts on a larger out-of-order core with register renaming (which probably overwrites agnostic elements with 1s). As it might be necessary to restart in the middle, we allow arbitrary mixing of agnostic policies within a single vector instruction. This allowed mixing of policies also enables implementations that might change policies for different granules of a vector register, for example, using undisturbed within a granule that is actively operated on but renaming to all 1s for granules in the tail.

The assembly syntax adds two flags to the vsetvli instruction:

```

ta  # Tail agnostic
tu  # Tail undisturbed
ma  # Mask agnostic
mu  # Mask undisturbed

vsetvli t0, a0, e32,m4,ta,ma  # Tail agnostic, mask agnostic
vsetvli t0, a0, e32,m4,tu,ma  # Tail undisturbed, mask agnostic
vsetvli t0, a0, e32,m4,ta,mu  # Tail agnostic, mask undisturbed
vsetvli t0, a0, e32,m4,tu,mu  # Tail undisturbed, mask undisturbed

```

To maintain backward compatibility in the short term and reduce software churn in the move to 0.9, when these flags are not specified on a vsetvli, they should default to mask-undisturbed/tail-undisturbed. The use of vsetvli without these flags should be deprecated, however, such that the specifying a flag setting becomes mandatory. If anything, the default should be tail-agnostic/mask-agnostic, so software has to specify when it cares about the non-participating elements, but given the historical meaning of the instruction prior to introduction of these flags, it is safest to always require them in future assembly code.

3.3.4. Vector Type Illegal vill

The vill bit is used to encode that a previous vsetvli instruction attempted to write an unsupported value to vtype.

The `vill` bit is held in bit `XLEN-1` of the CSR to support checking for illegal values with a branch on the sign bit.

If the `vill` bit is set, then any attempt to execute a vector instruction that depends upon `vtype` will raise an illegal-instruction exception.

`vsetvl{i}` and whole-register loads, stores, and moves do not depend upon `vtype`.

When the `vill` bit is set, the other `XLEN-1` bits in `vtype` shall be zero.

3.4. Vector Length Register `v1`

The `XLEN`-bit-wide read-only `v1` CSR can only be updated by the `vsetvli` and `vsetvl` instructions, and the *fault-only-first* vector load instruction variants.

The `v1` register holds an unsigned integer specifying the number of elements to be updated by a vector instruction. Elements in any destination vector register group with indices $\geq v1$ are unmodified during execution of a vector instruction. When $vstart \geq v1$, no elements are updated in any destination vector register group.

As a consequence, when $v1=0$, no elements are updated in the destination vector register group, regardless of `vstart`.

Instructions that write an `x` register or `f` register do so even when $vstart \geq v1$.

The number of bits implemented in `v1` depends on the implementation's maximum vector length of the smallest supported type. The smallest vector implementation, RV32IV, would need at least six bits in `v1` to hold the values 0-32 (with $VLEN=32$, $LMUL=8$ and $SEW=8$ results in $VLMAX$ of 32).

3.5. Vector Byte Length `vlenb`

The `XLEN`-bit-wide read-only CSR `vlenb` holds the value $VLEN/8$, i.e., the vector register length in bytes.

The value in `vlenb` is a design-time constant in any implementation.

Without this CSR, several instructions are needed to calculate $VLEN$ in bytes, and the code has to disturb current `v1` and `vtype` settings which require them to be saved and restored.

3.6. Vector Start Index CSR `vstart`

The `vstart` read-write CSR specifies the index of the first element to be executed by a vector instruction.

Normally, `vstart` is only written by hardware on a trap on a vector instruction, with the `vstart` value representing the element on which the trap was taken (either a synchronous exception or an asynchronous interrupt), and at which execution should resume after a resumable trap is handled.

All vector instructions are defined to begin execution with the element number given in the `vstart` CSR, leaving earlier elements in the destination vector undisturbed, and to reset the `vstart` CSR to zero at the end of execution.

All vector instructions, including `vsetvl{i}`, reset the `vstart` CSR to zero.

`vstart` is not modified by vector instructions that raise illegal-instruction exceptions.

For instructions where the number of elements to be performed is set by `v1`, if the value in the `vstart` register is greater than or equal to the vector length `v1` then no element operations are performed. The `vstart` register is then reset to zero.

The `vstart` CSR is defined to have only enough writable bits to hold the largest element index (one less than the maximum $VLMAX$) or $\lg_2(VLEN)$ bits.

The maximum vector length is obtained with the largest $LMUL$ setting (8) and the smallest SEW setting (8), so $VLMAX_max = 8 * VLEN / 8 = VLEN$. For example, for $VLEN=256$, `vstart` would have 8 bits to represent indices from 0 through 255.

The use of `vstart` values greater than the largest element index is reserved.

The `vstart` CSR is writable by unprivileged code, but non-zero `vstart` values may cause vector instructions to run substantially slower on some implementations, so `vstart` should not be used by application programmers. A few vector

instructions cannot be executed with a non-zero `vstart` value and will raise an illegal instruction exception as defined below.

Making `vstart` visible to unprivileged code supports user-level threading libraries.

Implementations are permitted to raise illegal instruction exceptions when attempting to execute a vector instruction with a value of `vstart` that the implementation can never produce when executing that same instruction with the same `vtype` setting.

For example, some implementations will never take interrupts during execution of a vector arithmetic instruction, instead waiting until the instruction completes to take the interrupt. Such implementations are permitted to raise an illegal instruction exception when attempting to execute a vector arithmetic instruction when `vstart` is nonzero.

When migrating a software thread between two harts with different microarchitectures, the `vstart` value might not be supported by the new hart microarchitecture. The runtime on the receiving hart might then have to emulate instruction execution to a supported `vstart` element position. Alternatively, migration events can be constrained to only occur at mutually supported `vstart` locations.

3.7. Vector Fixed-Point Rounding Mode Register `vxrm`

The vector fixed-point rounding-mode register holds a two-bit read-write rounding-mode field. The vector fixed-point rounding-mode is given a separate CSR address to allow independent access, but is also reflected as a field in `vcsr`.

The fixed-point rounding algorithm is specified as follows. Suppose the pre-rounding result is v , and d bits of that result are to be rounded off. Then the rounded result is $(v \gg d) + r$, where r depends on the rounding mode as specified in the following table.

Table 5. `vxrm` encoding

Bits [1:0]	Abbreviation	Rounding Mode	Rounding increment, r
0 0	<code>rnu</code>	round-to-nearest-up (add +0.5 LSB)	$v[d-1]$
0 1	<code>rne</code>	round-to-nearest-even	$v[d-1] \ \& \ (v[d-2:0] \neq 0 \mid v[d])$
1 0	<code>rdn</code>	round-down (truncate)	0
1 1	<code>rod</code>	round-to-odd (OR bits into LSB, aka "jam")	$!v[d] \ \& \ v[d-1:0] \neq 0$

The rounding functions:

```
roundoff_unsigned(v, d) = (unsigned(v) >> d) + r
roundoff_signed(v, d) = (signed(v) >> d) + r
```

are used to represent this operation in the instruction descriptions below.

Bits[XLEN-1:2] should be written as zeros.

The rounding mode can be set with a single `csrwi` instruction.

3.8. Vector Fixed-Point Saturation Flag `vxsat`

The `vxsat` CSR holds a single read-write bit that indicates if a fixed-point instruction has had to saturate an output value to fit into a destination format.

The `vxsat` bit is mirrored in `vcsr`.

3.9. Vector Control and Status Register `vcsr`

The `vxrm` and `vxsat` separate CSRs can also be accessed via fields in the vector control and status CSR, `vcsr`.

Table 6. `vcsr` layout

Bits	Name	Description
2:1	<code>vxrm[1:0]</code>	Fixed-point rounding mode
0	<code>vxsat</code>	Fixed-point accrued saturation flag

3.10. State of Vector Extension at Reset

The vector extension must have a consistent state at reset. In particular, `vtype` and `vl` must have values that can be read and then restored with a single `vsetvl` instruction.

It is recommended that at reset, `vtype.vill` is set, the remaining bits in `vtype` are zero, and `vl` is set to zero.

The `vstart`, `vxrm`, `vxsat` CSRs can have arbitrary values at reset.

Any use of the vector unit will require an initial `vsetvl{i}`, which will reset `vstart`. The `vxrm` and `vxsat` fields should be reset explicitly in software before use.

The vector registers can have arbitrary values at reset.

4. Mapping of Vector Elements to Vector Register State

The following diagrams illustrate how different width elements are packed into the bytes of a vector register depending on the current SEW and LMUL settings, as well as implementation VLEN. Elements are packed into each vector register with the least-significant byte in the lowest-numbered bits.

4.1. Mapping for LMUL = 1

When LMUL=1, elements are simply packed in order from the least-significant to most-significant bits of the vector register.

To increase readability, vector register layouts are drawn with bytes ordered from right to left with increasing byte address. Bits within an element are numbered in a little-endian format with increasing bit index from right to left corresponding to increasing magnitude.

LMUL=1 examples.

The element index is given in hexadecimal and is shown placed at the least-significant byte of the stored element.

VLEN=32b

Byte 3 2 1 0

SEW=8b 3 2 1 0

SEW=16b 1 0

SEW=32b 0

VLEN=64b

Byte 7 6 5 4 3 2 1 0

SEW=8b 7 6 5 4 3 2 1 0

SEW=16b 3 2 1 0

SEW=32b 1 0

SEW=64b 0

VLEN=128b

Byte F E D C B A 9 8 7 6 5 4 3 2 1 0

SEW=8b F E D C B A 9 8 7 6 5 4 3 2 1 0

SEW=16b 7 6 5 4 3 2 1 0

SEW=32b 3 2 1 0

SEW=64b 1 0

SEW=128b 0

VLEN=256b

Byte 1F 1E 1D 1C 1B 1A 19 18 17 16 15 14 13 12 11 10 F E D C B A 9 8 7 6 5 4 3 2 1 0

SEW=8b 1F 1E 1D 1C 1B 1A 19 18 17 16 15 14 13 12 11 10 F E D C B A 9 8 7 6 5 4 3 2 1 0

SEW=16b F E D C B A 9 8 7 6 5 4 3 2 1 0

SEW=32b 7 6 5 4 3 2 1 0

SEW=64b 3 2 1 0

SEW=128b 1 0

When LMUL < 1, only the first LMUL*VLEN/SEW elements in the vector register are used. The remaining space in the vector register is treated as part of the tail.

Example, VLEN=128b, LMUL=1/4

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0
SEW=8b	-	-	-	-	-	-	-	-	-	-	-	-	-	3	2	1 0
SEW=16b	-	-	-	-	-	-	-	-	-	-	-	-	-	1	0	0
SEW=32b	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	0

4.2. Mapping with LMUL > 1

When vector registers are grouped, the elements of the vector register group are striped across the constituent vector registers. The elements are packed contiguously in element order in each vector register in the group, moving to the next highest-numbered vector register in the group once each vector register is filled.

LMUL > 1 examples

VLEN=32b, SEW=8b, LMUL=2

Byte	3	2	1	0
v2*n	3	2	1	0
v2*n+1	7	6	5	4

VLEN=32b, SEW=16b, LMUL=2

Byte	3	2	1	0
v2*n	1	0		
v2*n+1	3	2		

VLEN=32b, SEW=16b, LMUL=4

Byte	3	2	1	0
v4*n	1	0		
v4*n+1	3	2		
v4*n+2	5	4		
v4*n+3	7	6		

VLEN=32b, SEW=32b, LMUL=4

Byte	3	2	1	0
v4*n				0
v4*n+1				1
v4*n+2				2
v4*n+3				3

VLEN=64b, SEW=32b, LMUL=2

Byte	7	6	5	4	3	2	1	0
v2*n				1				0
v2*n+1				3				2

VLEN=64b, SEW=32b, LMUL=4

Byte	7	6	5	4	3	2	1	0
v4*n				1				0
v4*n+1				3				2
v4*n+2				5				4
v4*n+3				7				6

VLEN=128b, SEW=32b, LMUL=2

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0
v2*n				3				2				1				0
v2*n+1				7				6				5				4

VLEN=128b, SEW=32b, LMUL=4

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0
v4*n					3				2			1				0
v4*n+1					7				6			5				4
v4*n+2					B				A			9				8
v4*n+3					F				E			D				C

4.3. Mapping across Mixed-Width Operations

The vector ISA is designed to support mixed-width operations without requiring explicit additional rearrangement instructions. The recommended software strategy is to modify vtype dynamically to keep SEW/LMUL constant (and hence

VLMAX constant) when operating on vectors of different precision values.

The following example shows four different packed element widths (8b, 16b, 32b, 64b) in a VLEN=128b implementation. The vector register grouping factor (LMUL) is increased by the relative element size such that each group can hold the same number of vector elements (VLMAX=8 in this example) to simplify stripmining code.

Example VLEN=128b, with SEW/LMUL=16

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	
vn	-	-	-	-	-	-	-	7	6	5	4	3	2	1	0		SEW=8b, LMUL=1/2
vn								7	6	5	4	3	2	1	0		SEW=16b, LMUL=1
v2*n																	SEW=32b, LMUL=2
v2*n+1																	
v4*n																	SEW=64b, LMUL=4
v4*n+1																	
v4*n+2																	
v4*n+3																	

The following table shows each possible constant SEW/LMUL operating point for loops with mixed-width operations. Each column represents a constant SEW/LMUL operating point. Entries in table are the LMUL values that yield that column's SEW/LMUL value for the datawidth on that row. In each column, an LMUL setting for a datawidth indicates that it can be aligned with the other datawidths in the same column that also have an LMUL setting, such that all have the same VLMAX.

SEW/LMUL	1	2	4	8	16	32	64	128	256	512	1024	2048	4096	8192
SEW= 8	8	4	2	1	1/2	1/4	1/8							
SEW= 16		8	4	2	1	1/2	1/4	1/8						
SEW= 32			8	4	2	1	1/2	1/4	1/8					
SEW= 64				8	4	2	1	1/2	1/4	1/8				
SEW= 128					8	4	2	1	1/2	1/4	1/8			
SEW= 256						8	4	2	1	1/2	1/4	1/8		
SEW= 512							8	4	2	1	1/2	1/4	1/8	
SEW=1024								8	4	2	1	1/2	1/4	1/8

Larger LMUL settings can also used to simply increase vector length to reduce instruction fetch and dispatch overheads in cases where fewer vector register groups are needed.

The SEW/LMUL values of 2048 and greater are shown in the table for completeness but they do not add a useful operating point in the base architecture as they use less than the full register capacity and do not enable more architectural registers.

4.4. Mapping with LMUL > 1 and ELEN > VLEN

If vector registers are grouped to support larger SEW, with ELEN > VLEN, the vector registers in the group are concatenated to form a single array of bytes, with the lowest-numbered register in the group holding the lowest-addressed bytes from the memory layout.

LMUL > 1 ELEN>VLEN, examples

VLEN=32b, SEW=64b, LMUL=2

Byte	3	2	1	0
v2*n				0
v2*n+1				

VLEN=32b, SEW=64b, LMUL=4

Byte	3	2	1	0
v4*n				0
v4*n+1				
v4*n+2				1
v4*n+3				

VLEN=32b, SEW=64b, LMUL=8

Byte	3	2	1	0
v8*n				0
v8*n+1				
v8*n+2				1
v8*n+3				
v8*n+4				2
v8*n+5				
v8*n+6				3
v8*n+7				

4.5. Mask Register Layout

A vector mask occupies only one vector register regardless of SEW and LMUL. Each element is allocated a single mask bit in a mask vector register.

Earlier designs (pre-0.9) had a varying number of bits per mask value (MLEN). In the 0.9 design, MLEN=1.

4.5.1. Mask Element Locations

The mask bit for element i is located in bit i of the mask register, independent of SEW or LMUL.

VLEN=32b

Byte	3	2	1	0
LMUL=1, SEW=8b				
	3	2	1	0
	[03]	[02]	[01]	[00]
				Element
				Mask bit position in decimal

LMUL=2, SEW=16b

	1	0
	[01]	[00]
	3	2
	[03]	[02]

LMUL=4, SEW=32b

	0
	[00]
	1
	[01]
	2
	[02]
	3
	[03]

LMUL=2, SEW=8b

3	2	1	0
[03]	[02]	[01]	[00]
7	6	5	4
[07]	[06]	[05]	[04]

LMUL=8, SEW=32b

0
[00]
1
[01]
2
[02]
3
[03]
4
[04]
5
[05]
6
[06]
7
[07]

LMUL=8, SEW=8b


3	2	1	0
[03]	[02]	[01]	[00]
7	6	5	4
[07]	[06]	[05]	[04]
B	A	9	8
[11]	[10]	[09]	[08]
F	E	D	C
[15]	[14]	[13]	[12]
13	12	11	10
[19]	[18]	[17]	[16]
17	16	15	14
[23]	[22]	[21]	[20]
1B	1A	19	18
[27]	[26]	[25]	[24]
1F	1E	1D	1C
[31]	[30]	[29]	[28]

5. Vector Instruction Formats


The instructions in the vector extension fit under three existing major opcodes (LOAD-FP, STORE-FP, AMO) and one new major opcode (OP-V).

Vector loads and stores are encoding within the scalar floating-point load and store major opcodes (LOAD-FP/STORE-FP). The vector load and store encodings repurpose a portion of the standard scalar floating-point load/store 12-bit immediate field to provide further vector instruction encoding, with bit 25 holding the standard vector mask bit (see [Mask Encoding](#)).


Format for Vector Load Instructions under LOAD-FP major opcode




Format for Vector Store Instructions under STORE-FP major opcode




Format for Vector AMO Instructions under AMO major opcode



Formats for Vector Arithmetic Instructions under OP-V major opcode



Formats for Vector Configuration Instructions under OP-V major opcode



Vector instructions can have scalar or vector source operands and produce scalar or vector results, and most vector instructions can be performed either unconditionally or conditionally under a mask.

Vector loads and stores move bit patterns between vector register elements and memory. Vector arithmetic instructions operate on values held in vector register elements.

5.1. Scalar Operands

Scalar operands can be immediates, or taken from the x registers, the f registers, or element 0 of a vector register. Scalar results are written to an x or f register or to element 0 of a vector register. Any vector register can be used to hold a scalar regardless of the current LMUL setting.

In a change from v0.6, the floating-point registers no longer overlay the vector registers and scalars can now come from the integer or floating-point registers. Not overlaying the f registers reduces vector register pressure, avoids interactions with the standard calling convention, simplifies high-performance scalar floating-point design, and provides compatibility with the Zfinx ISA option. Overlaying f with v would provide the advantage of lowering the number of state bits in some implementations, but complicates high-performance designs and would prevent compatibility with the Zfinx ISA option.

5.2. Vector Operands

Each vector operand has an *effective element width* (EEW) and an *effective LMUL* (EMUL) that is used to determine the size and location of all the elements within a vector register group. By default, for most operands of most instructions, EEW=SEW and EMUL=LMUL.

Some vector instructions have source and destination vector operands with the same number of elements but different widths, so that EEW and EMUL differ from SEW and LMUL respectively but EEW/EMUL = SEW/LMUL. For example, most widening arithmetic instructions have a source group with EEW=SEW and EMUL=LMUL but destination group with EEW=2*SEW and EMUL=2*LMUL. Narrowing instructions have a source operand that has EEW=2*SEW and EMUL=2*LMUL but destination where EEW=SEW and EMUL=LMUL.

Vector operands or results may occupy one or more vector registers depending on EMUL, but are always specified using the lowest-numbered vector register in the group. Using other than the lowest-numbered vector register to specify a vector register group is reserved..

A destination vector register group can overlap a source vector register group only if one of the following holds:

- The destination EEW equals the source EEW.
- The destination EEW is smaller than the source EEW and the overlap is in the lowest-numbered part of the source register group (e.g., when LMUL=1, vnsr1.wi v0, v0, 3 is legal, but a destination of v1 is not).
- The destination EEW is greater than the source EEW, the source EMUL is at least 1, and the overlap is in the highest-numbered part of the destination register group (e.g., when LMUL=8, vzext.vf4 v0, v6 is legal, but a source of v0, v2, or v4 is not).

For the purpose of register group overlap constraints, mask elements have EEW=1.

The largest vector register group used by an instruction can not be greater than 8 vector registers (i.e., EMUL≤8), and if a vector instruction would require greater than 8 vector registers in a group, the instruction encoding is reserved. For example, a widening operation that produces a widened vector register group result when LMUL=8 is reserved as this would imply a result EMUL=16.

Widened scalar values, e.g., results from widening reduction operations, are held in the first element of a vector register and have EMUL=1.

Current reduction operations are defined to hold input and output values in a single vector register, with implicit EMUL of 1, so cannot accommodate using a vector register group to hold a wide scalar reduction result. This would require an independent parameter to give the EMUL for the scalar reduction element.

5.3. Vector Masking

Masking is supported on many vector instructions. Element operations that are masked off (inactive) never generate exceptions. The destination vector register elements corresponding to masked-off elements are handled with either a mask-undisturbed or mask-agnostic policy depending on the setting of the *vma* bit in *vtype* (Section [Vector Tail Agnostic and Vector Mask Agnostic *vta* and *vma*](#)).

In the base vector extension, the mask value used to control execution of a masked vector instruction is always supplied by vector register *v0*.

Future vector extensions may provide longer instruction encodings with space for a full mask register specifier.

The destination vector register group for a masked vector instruction cannot overlap the source mask register (*v0*), unless the destination vector register is being written with a mask value (e.g., comparisons) or the scalar result of a reduction. These instruction encodings are reserved.

This constraint supports restart with a non-zero *vstart* value.

Some masked instructions that target *v0* which were legal in v0.8 are illegal with the new MLEN=1 mask layout for v1.0. For example, `vadd.vv v0, v1, v2, v0.m` is now always illegal; previously, it was legal for LMUL=1.

Other vector registers can be used to hold working mask values, and mask vector logical operations are provided to perform predicate calculations.

When a mask is written with a compare result, destination mask bits past the end of the current vector length are handled according to the tail policy (undisturbed or agnostic) set by the *vta* bit in *vtype* (Section [Vector Tail Agnostic and Vector Mask Agnostic *vta* and *vma*](#)).

5.3.1. Mask Encoding

Where available, masking is encoded in a single-bit *vm* field in the instruction (`inst[25]`).

vm	Description
0	vector result, only where <code>v0.mask[i] = 1</code>
1	unmasked

In earlier proposals, *vm* was a two-bit field `vm[1:0]` that provided both true and complement masking using *v0* as well as encoding scalar operations.

Vector masking is represented in assembler code as another vector operand, with `.t` indicating if operation occurs when `v0.mask[i]` is 1. If no masking operand is specified, unmasked vector execution (`vm=1`) is assumed.

<code>vop.v*</code>	<code>v1, v2, v3, v0.t</code>	# enabled where <code>v0.mask[i]=1</code> , <code>m=0</code>
<code>vop.v*</code>	<code>v1, v2, v3</code>	# unmasked vector operation, <code>m=1</code>

Even though the base only supports one vector mask register *v0* and only the true form of predication, the assembly syntax writes it out in full to be compatible with future extensions that might add a mask register specifier and supporting both true and complement masking. The `.t` suffix on the masking operand also helps to visually encode the use of a mask.

The `.mask` suffix is not part of the assembly syntax. We only append it in contexts where a mask vector is subscripted, e.g., `v0.mask[i]`.

5.4. Prestart, Active, Inactive, Body, and Tail Element Definitions

The destination element indices operated on during a vector instruction's execution can be divided into three disjoint subsets.

- The *prestart* elements are those whose element index is less than the initial value in the *vstart* register. The prestart elements do not raise exceptions and do not update the destination vector register.

- The *body* elements are those whose element index is greater than or equal to the initial value in the `vstart` register, and less than the current vector length setting in `v1`. The body can be split into two disjoint subsets:
 - The *active* elements during a vector instruction's execution are the elements within the body and where the current mask is enabled at that element position. The active elements can raise exceptions and update the destination vector register group.
 - The *inactive* elements are the elements within the body but where the current mask is disabled at that element position. The inactive elements do not raise exceptions and do not update any destination vector register group unless masked agnostic is specified (`vtype.vma=1`), in which case inactive elements may be overwritten with 1s.
- The *tail* elements during a vector instruction's execution are the elements past the current vector length setting specified in `v1`. The tail elements do not raise exceptions, and do not update any destination vector register group unless tail agnostic is specified (`vtype.vta=1`), in which case tail elements may be overwritten with 1s. When `LMUL < 1`, the tail includes the elements past `VLMAX` that are held in the same vector register.

```

for element index x
prestart    = (0 <= x < vstart)
body(x)     = (vstart <= x < v1)
tail(x)     = (v1 <= x < max(VLMAX, VLEN/SEW))
mask(x)     = unmasked || v0.mask[x] == 1
active(x)   = body(x) && mask(x)
inactive(x) = body(x) && !mask(x)

```

Some instructions such as `vslidtdown` and `vrgather` may read indices past `v1` or even `VLMAX` in source vector register groups. The general policy is to return the value 0 when the index is greater than `VLMAX` in the source vector register group.

6. Configuration-Setting Instructions

One of the common approaches to handling a large number of elements is "stripmining" where each iteration of a loop handles some number of elements, and the iterations continue until all elements have been processed. The RISC-V vector specification provides direct, portable support for this approach. The application specifies the total number of elements to be processed (the application vector length or AVL) as a candidate value for `v1`, and the hardware responds via a general-purpose register with the (frequently smaller) number of elements that the hardware will handle per iteration (stored in `v1`), based on the microarchitectural implementation and the `vtype` setting. A straightforward loop structure, shown in [Example of stripmining and changes to SEW](#), depicts the ease with which the code keeps track of the remaining number of elements and the amount per iteration handled by hardware.

A set of instructions is provided to allow rapid configuration of the values in `vl` and `vtype` to match application needs.

6.1. vsetvli/vsetvl instructions

The `vsetvl{i}` instructions set the `vtype` and `vl` CSRs based on their arguments, and write the new value of `vl` into `rd`.

```
vsetvli rd, rs1, vtypei # rd = new v1, rs1 = AVL, vtypei = new vtype setting
vsetvl rd, rs1, rs2     # rd = new v1, rs1 = AVL, rs2 = new vtype value
```

The new vtype setting is encoded in the immediate fields of vsetvli and in the rs2 register for vsetv1. The new vector length setting is based on AVL, which is encoded in the rs1 and rd fields as follows:

Table 7. AVL used in `vsetvli` and `vsetvl` instructions

rd	rs1	AVL value	Effect on v1
-	!x0	Value in x[rs1]	Normal stripmining
!x0	x0	~0	Set v1 to VLMAX
x0	x0	Value in v1 register	Keep existing v1 (of course, vtype may change)

When `rs1` is not `x0`, the AVL is an unsigned integer held in the x register specified by `rs1`, and the new `v1` value is also written to the x register specified by `rd`.

When `rs1=x0` but `rd!=x0`, the maximum unsigned integer value (`~0`) is used as the AVL, and the resulting VLMAX is written to `v1` and also to the `x` register specified by `rd`.

When `rs1=x0` and `rd=x0`, the instruction operates as if the current vector length in `v1` is used as the AVL, and the resulting value is written to `v1`, but not to a destination register. This form can only be used when VLMAX and hence `v1` is not actually changed by the new SEW/LMUL ratio.

Use of the instruction with a new SEW/LMUL ratio that would result in a change of VLMAX is reserved. Implementations may set `vll` in this case.

This last form of the instructions allows the `vtype` register to be changed while maintaining the current `v1`, provided `VLMAX` is not reduced. This design was chosen to ensure `v1` would always hold a legal value for current `vtype` setting. The current `v1` value can be read from the `v1 CSR`. The `v1` value could be reduced by this instruction if the new `SEW/LMUL` ratio causes `VLMAX` to shrink, and so this case has been reserved as it is not clear this is a generally useful operation, and implementations can otherwise assume `v1` is not changed by this instruction to optimize their microarchitecture.

Formats for Vector Configuration Instructions under OP-V major opcode




Table 8. vtype register layout

Bits	Name	Description
XLEN-1	vill	Illegal value if set
XLEN-2:8		Reserved (write 0)
7	vma	Vector mask agnostic
6	vta	Vector tail agnostic
5:3	vsew[2:0]	Selected element width (SEW) setting
2:0	vlmul[2:0]	Vector register group multiplier (LMUL) setting

Suggested assembler names used for vsetvli immediate

```
e8    # SEW=8b
e16   # SEW=16b
e32   # SEW=32b
e64   # SEW=64b
e128  # SEW=128b
e256  # SEW=256b
e512  # SEW=512b
e1024 # SEW=1024b
```

```
mf8   # LMUL=1/8
mf4   # LMUL=1/4
mf2   # LMUL=1/2
m1    # LMUL=1, assumed if m setting absent
m2    # LMUL=2
m4    # LMUL=4
m8    # LMUL=8
```

Examples:

```
vsetvli t0, a0, e8          # SEW= 8, LMUL=1
vsetvli t0, a0, e8,m2       # SEW= 8, LMUL=2
vsetvli t0, a0, e32,mf2     # SEW=32, LMUL=1/2
```

The vsetvl variant operates similarly to vsetvli except that it takes a vtype value from rs2 and can be used for context restore.

If the vtype setting is not supported by the implementation, then the vill bit is set in vtype, the remaining bits in vtype are set to zero, and the vl register is also set to zero.

Earlier drafts required a trap when setting vtype to an illegal value. However, this would have added the first data-dependent trap on a CSR write to the ISA. The current scheme supports light-weight runtime interrogation of the supported vector unit configurations by checking if vill is clear for a given setting.

Implementations may choose to trap when illegal values are written to vtype instead of setting vill, to allow emulation to support new configurations for forward-compatibility.

6.2. Constraints on Setting vl

The vsetvl{i} instructions first set VLMAX according to the vtype argument, then set vl obeying the following constraints:

1. $vl = AVL$ if $AVL \leq VLMAX$
2. $ceil(AVL / 2) \leq vl \leq VLMAX$ if $AVL < (2 * VLMAX)$
3. $vl = VLMAX$ if $AVL \geq (2 * VLMAX)$
4. Deterministic on any given implementation for same input AVL and VLMAX values
5. These specific properties follow from the prior rules:

- a. $v1 = 0$ if $AVL = 0$
- b. $v1 > 0$ if $AVL > 0$
- c. $v1 \leq VLMAX$
- d. $v1 \leq AVL$
- e. a value read from $v1$ when used as the AVL argument to $vsetv1\{i\}$ results in the same value in $v1$, provided the resultant $VLMAX$ equals the value of $VLMAX$ at the time that $v1$ was read

The $v1$ setting rules are designed to be sufficiently strict to preserve $v1$ behavior across register spills and context swaps for $AVL \leq VLMAX$, yet flexible enough to enable implementations to improve vector lane utilization for $AVL > VLMAX$.

For example, this permits an implementation to set $v1 = \text{ceil}(AVL / 2)$ for $VLMAX < AVL < 2*VLMAX$ in order to evenly distribute work over the last two iterations of a stripmine loop. Requirement 2 ensures that the first stripmine iteration of reduction loops uses the largest vector length of all iterations, even in the case of $AVL < 2*VLMAX$. This allows software to avoid needing to explicitly calculate a running maximum of vector lengths observed during a stripmined loop.

6.3. Example of stripmining and changes to SEW

The SEW and LMUL settings can be changed dynamically to provide high throughput on mixed-width operations in a single loop.

```
# Example: Load 16-bit values, widen multiply to 32b, shift 32b result
# right by 3, store 32b values.
# On entry:
#  a0 holds the total number of elements to process
#  a1 holds the address of the source array
#  a2 holds the address of the destination array

loop:
    vsetvli a3, a0, e16,m4,ta,ma # vtype = 16-bit integer vectors;
                                # also update a3 with v1 (# of elements this iteration)
    vle16.v v4, (a1)           # Get 16b vector
    slli t1, a3, 1              # Multiply # elements this iteration by 2 bytes/source element
    add a1, a1, t1              # Bump pointer
    vwmul.vx v8, v4, x10        # Widening multiply into 32b in <v8--v15>

    vsetvli x0, x0, e32,m8,ta,ma # Operate on 32b values
    vsrl.vi v8, v8, 3
    vse32.v v8, (a2)           # Store vector of 32b elements
    slli t1, a3, 2              # Multiply # elements this iteration by 4 bytes/destination element
    add a2, a2, t1              # Bump pointer
    sub a0, a0, a3              # Decrement count by v1
    bnez a0, loop              # Any more?
```


7. Vector Loads and Stores

Vector loads and stores move values between vector registers and memory. Vector loads and stores are masked and do not raise exceptions on inactive elements. Masked vector loads do not update inactive elements in the destination vector register group. Masked vector stores only update active memory elements. All vector loads and stores may generate and accept a non-zero vstart value.


7.1. Vector Load/Store Instruction Encoding

Vector loads and stores are encoded within the scalar floating-point load and store major opcodes (LOAD-FP/STORE-FP). The vector load and store encodings repurpose a portion of the standard scalar floating-point load/store 12-bit immediate field to provide further vector instruction encoding, with bit 25 holding the standard vector mask bit (see [Mask Encoding](#)).

Format for Vector Load Instructions under LOAD-FP major opcode



Format for Vector Store Instructions under STORE-FP major opcode



Field	Description
rs1[4:0]	specifies x register holding base address
rs2[4:0]	specifies x register holding stride
vs2[4:0]	specifies v register holding address offsets
vs3[4:0]	specifies v register holding store data
vd[4:0]	specifies v register destination of load
vm	specifies whether vector masking is enabled (0 = mask enabled, 1 = mask disabled)
width[2:0]	specifies size of memory elements, and distinguishes from FP scalar
mew	extended memory element width. See Vector Load/Store Width Encoding
mop[1:0]	specifies memory addressing mode
nf[2:0]	specifies the number of fields in each segment, for segment load/stores
lumop[4:0]/sumop[4:0]	are additional fields encoding variants of unit-stride instructions

Vector memory unit-stride and constant-stride operations directly encode EEW of the data to be transferred statically in the instruction to reduce the number of vtype changes when accessing memory in a mixed-width routine. Indexed operations use the explicit EEW encoding in the instruction to set the size of the indices used, and use SEW/LMUL to specify the data width.

7.2. Vector Load/Store Addressing Modes

The base vector extension supports unit-stride, strided, and indexed (scatter/gather) addressing modes. Vector load/store base registers and strides are taken from the GPR x registers.

The base effective address for all vector accesses is given by the contents of the x register named in rs1.

Vector unit-stride operations access elements stored contiguously in memory starting from the base effective address.

Vector constant-strided operations access the first memory element at the base effective address, and then access subsequent elements at address increments given by the byte offset contained in the x register specified by rs2.

Vector indexed operations add the contents of each element of the vector offset operand specified by vs2 to the base effective address to give the effective address of each element. The data vector register group has EEW=SEW, EMUL=LMUL, while the offset vector register group has EEW encoding in the instruction and EMUL=(EEW/SEW)*LMUL.

The vector offset operand is treated as a vector of byte-address offsets.

The indexed operations can also be used to access fields within a vector of objects, where the vs2 vector holds pointers to the base of the objects and the scalar x register holds the offset of the member field in each object. Supporting this case is why the indexed operations were not defined to scale the element indices by the data EEW.

If the vector offset elements are narrower than XLEN, they are zero-extended to XLEN before adding to the base effective address. If the vector offset elements are wider than XLEN, the least-significant XLEN bits are used in the address calculation.

A profile may place a upper limit on the maximum required index EEW (e.g., only up to XLEN) smaller than ELEN, in which case an illegal instruction exception can be raised if the EEW is not supported.

The vector addressing modes are encoded using the 2-bit mop[1:0] field.

Table 9. encoding for loads

mop [1:0]		Description	Opcodes
0	0	unit-stride	VLE<EEW>
0	1	indexed-unordered	VLUXEI<EEW>
1	0	strided	VLSE<EEW>
1	1	indexed-ordered	VLOXEI<EEW>

Table 10. encoding for stores

mop [1:0]		Description	Opcodes
0	0	unit-stride	VSE<EEW>
0	1	indexed-unordered	VSUXEI<EEW>
1	0	strided	VSSE<EEW>
1	1	indexed-ordered	VSOXEI<EEW>

Vector unit-stride and constant-stride memory accesses do not guarantee ordering between individual element accesses. The vector indexed load and store memory operations have two forms, ordered and unordered. The indexed-ordered variants preserve element ordering on memory accesses.

For unordered instructions (mop!=11) there is no guarantee on element access order. If the accesses are to a strongly ordered IO region, the element accesses can be initiated in any order.

To provide ordered vector accesses to a strongly ordered IO region, the ordered indexed instructions should be used.

For implementations with precise vector traps, exceptions on indexed-unordered stores are precise.

Additional unit-stride vector addressing modes are encoded using the 5-bit lumop and sumop fields in the unit-stride load and store instruction encodings respectively.

Table 11. lumop

lumop[4:0]					Description
0	0	0	0	0	unit-stride
0	0	x	x	x	reserved, x !=0
0	1	0	0	0	unit-stride, whole registers
0	1	x	x	x	reserved, x !=0
1	0	0	0	0	unit-stride fault-only-first
1	x	x	x	x	reserved, x!=0

Table 12. sumop

sumop[4:0]					Description
0	0	0	0	0	unit-stride
0	0	x	x	x	reserved, x !=0
0	1	0	0	0	unit-stride, whole registers
0	1	x	x	x	reserved, x !=0
1	x	x	x	x	reserved

The $nf[2:0]$ field encodes the number of fields in each segment. For regular vector loads and stores, $nf=0$, indicating that a single value is moved between a vector register group and memory at each element position. Larger values in the nf field are used to access multiple contiguous fields within a segment as described below in Section [Vector Load/Store Segment Instructions](#).

The nf field for segment load/stores has replaced the use of the same bits for an address offset field. The offset can be replaced with a single scalar integer calculation, while segment load/stores add more powerful primitives to move items to and from memory.

The $nf[2:0]$ field also encodes the number of whole vector registers to transfer for the whole vector register load/store instructions.

7.3. Vector Load/Store Width Encoding

Vector loads and stores have an EEW encoded directly in the instruction. The corresponding EMUL is calculated as $EMUL = (EEW/SEW)*LMUL$. If the EMUL would be out of range ($EMUL > 8$ or $EMUL < 1/8$), the instruction encoding is reserved. The vector register groups must have legal register specifiers for the selected EMUL, else the instruction encoding is reserved.

Vector unit-stride and constant-stride use the EEW/EMUL encoded in the instruction for the data values, while vector indexed loads and stores use the EEW/EMUL encoded in the instruction for the index values and the SEW/LMUL encoded in $vtype$ for the data values.

Vector loads and stores are encoded using width values that are not claimed by the standard scalar floating-point loads and stores. The mew bit ($inst[28]$) encodes expanded memory sizes of 128 bits and above.

Vector loads and stores up to $EEW=ELEN$ must be supported in an implementation. Using a vector load/store with an unsupported EEW is reserved.

Table 13. Width encoding for vector loads and stores.

	mew	width [2:0]			Mem bits	Data Reg bits	Index bits	Opcodes
Standard scalar FP	x	0	0	1	16	FLEN	-	FLH/FSH
Standard scalar FP	x	0	1	0	32	FLEN	-	FLW/FSW
Standard scalar FP	x	0	1	1	64	FLEN	-	FLD/FSD
Standard scalar FP	x	1	0	0	128	FLEN	-	FLQ/FSQ
Vector 8b element	0	0	0	0	8	8	-	VLxE8/VSxE8
Vector 16b element	0	1	0	1	16	16	-	VLxE16/VSxE16
Vector 32b element	0	1	1	0	32	32	-	VLxE32/VSxE32
Vector 64b element	0	1	1	1	64	64	-	VLxE64/VSxE64
Vector 128b element	1	0	0	0	128	128	-	VLxE128/VSxE128
Vector 256b element	1	1	0	1	256	256	-	VLxE256/VSxE256
Vector 512b element	1	1	1	0	512	512	-	VLxE512/VSxE512
Vector 1024b element	1	1	1	1	1024	1024	-	VLxE1024/VSxE1024
Vector 8b index	0	0	0	0	SEW	SEW	8	VLxEI8/VSxEI8
Vector 16b index	0	1	0	1	SEW	SEW	16	VLxEI16/VSxEI16
Vector 32b index	0	1	1	0	SEW	SEW	32	VLxEI32/VSxEI32
Vector 64b index	0	1	1	1	SEW	SEW	64	VLxEI64/VSxEI64

Mem bits is the size of each element accessed in memory.

Data reg bits is the size of each data element accessed in register.

In base V extension, only data elements widths up to max(XLEN,FLEN) must be supported.

Index bits is the size of each index accessed in register.

In base V extension, only index widths up to XLEN must be supported.

Index bit EEW encodings larger than 64b are currently reserved.

RV128 will require index EEW of 128.

7.4. Vector Unit-Stride Instructions

Vector unit-stride loads and stores

```
# vd destination, rs1 base address, vm is mask encoding (v0.t or <missing>)
vle8.v    vd, (rs1), vm # 8-bit unit-stride load
vle16.v   vd, (rs1), vm # 16-bit unit-stride load
vle32.v   vd, (rs1), vm # 32-bit unit-stride load
vle64.v   vd, (rs1), vm # 64-bit unit-stride load
vle128.v  vd, (rs1), vm # 128-bit unit-stride load
vle256.v  vd, (rs1), vm # 256-bit unit-stride load
vle512.v  vd, (rs1), vm # 512-bit unit-stride load
vle1024.v vd, (rs1), vm # 1024-bit unit-stride load

# vs3 store data, rs1 base address, vm is mask encoding (v0.t or <missing>)
vse8.v    vs3, (rs1), vm # 8-bit unit-stride store
vse16.v   vs3, (rs1), vm # 16-bit unit-stride store
vse32.v   vs3, (rs1), vm # 32-bit unit-stride store
vse64.v   vs3, (rs1), vm # 64-bit unit-stride store
vse128.v  vs3, (rs1), vm # 128-bit unit-stride store
vse256.v  vs3, (rs1), vm # 256-bit unit-stride store
vse512.v  vs3, (rs1), vm # 512-bit unit-stride store
vse1024.v vs3, (rs1), vm # 1024-bit unit-stride store
```

7.5. Vector Strided Instructions

Vector strided loads and stores

```
# vd destination, rs1 base address, rs2 byte stride
vlse8.v   vd, (rs1), rs2, vm # 8-bit strided load
vlse16.v  vd, (rs1), rs2, vm # 16-bit strided load
vlse32.v  vd, (rs1), rs2, vm # 32-bit strided load
vlse64.v  vd, (rs1), rs2, vm # 64-bit strided load
vlse128.v vd, (rs1), rs2, vm # 128-bit strided load
vlse256.v vd, (rs1), rs2, vm # 256-bit strided load
vlse512.v vd, (rs1), rs2, vm # 512-bit strided load
vlse1024.v vd, (rs1), rs2, vm # 1024-bit strided load

# vs3 store data, rs1 base address, rs2 byte stride
vsse8.v   vs3, (rs1), rs2, vm # 8-bit strided store
vsse16.v  vs3, (rs1), rs2, vm # 16-bit strided store
vsse32.v  vs3, (rs1), rs2, vm # 32-bit strided store
vsse64.v  vs3, (rs1), rs2, vm # 64-bit strided store
vsse128.v vs3, (rs1), rs2, vm # 128-bit strided store
vsse256.v vs3, (rs1), rs2, vm # 256-bit strided store
vsse512.v vs3, (rs1), rs2, vm # 512-bit strided store
vsse1024.v vs3, (rs1), rs2, vm # 1024-bit strided store
```

Negative and zero strides are supported.

Element accesses within a strided instruction are unordered with respect to each other.

When $rs2=x0$, then an implementation is allowed, but not required, to perform fewer memory operations than the number of active elements, and may perform different numbers of memory operations across different dynamic executions of the same static instruction.

Compilers must be aware to not use the $x0$ form for $rs2$ when the immediate stride is 0 if the intent is to require all memory accesses are performed.

When $rs2 \neq x0$ and the value of $x[rs2]=0$, the implementation must perform one memory access for each active element (but these accesses will not be ordered).

When repeating ordered vector accesses to the same memory address are required, then an ordered indexed operation can be used.

7.6. Vector Indexed Instructions

```
# Vector indexed loads and stores

# Vector unordered indexed load instructions
# vd destination, rs1 base address, vs2 indices
vluxei8.v    vd, (rs1), vs2, vm # unordered 8-bit indexed load of SEW data
vluxei16.v   vd, (rs1), vs2, vm # unordered 16-bit indexed load of SEW data
vluxei32.v   vd, (rs1), vs2, vm # unordered 32-bit indexed load of SEW data
vluxei64.v   vd, (rs1), vs2, vm # unordered 64-bit indexed load of SEW data

# Vector ordered indexed load instructions
# vd destination, rs1 base address, vs2 indices
vloxei8.v    vd, (rs1), vs2, vm # ordered 8-bit indexed load of SEW data
vloxei16.v   vd, (rs1), vs2, vm # ordered 16-bit indexed load of SEW data
vloxei32.v   vd, (rs1), vs2, vm # ordered 32-bit indexed load of SEW data
vloxei64.v   vd, (rs1), vs2, vm # ordered 64-bit indexed load of SEW data

# Vector unordered-indexed store instructions
# vs3 store data, rs1 base address, vs2 indices
vsuxei8.v    vs3, (rs1), vs2, vm # unordered 8-bit indexed store of SEW data
vsuxei16.v   vs3, (rs1), vs2, vm # unordered 16-bit indexed store of SEW data
vsuxei32.v   vs3, (rs1), vs2, vm # unordered 32-bit indexed store of SEW data
vsuxei64.v   vs3, (rs1), vs2, vm # unordered 64-bit indexed store of SEW data

# Vector ordered indexed store instructions
# vs3 store data, rs1 base address, vs2 indices
vsoxei8.v    vs3, (rs1), vs2, vm # ordered 8-bit indexed store of SEW data
vsoxei16.v   vs3, (rs1), vs2, vm # ordered 16-bit indexed store of SEW data
vsoxei32.v   vs3, (rs1), vs2, vm # ordered 32-bit indexed store of SEW data
vsoxei64.v   vs3, (rs1), vs2, vm # ordered 64-bit indexed store of SEW data
```

The assembler syntax for indexed loads and stores uses `eix` instead of `ex` to indicate the statically encoded EEW is of the index not the data.

The indexed operations mnemonics have a "U" or "O" to distinguish between unordered and ordered, while the other vector addressing modes have no character. While this is perhaps a little less consistent, this approach minimizes disruption to existing software, as `VSXEI` previously meant "ordered" - and the opcode can be retained as an alias during transition to help reduce software churn.

7.7. Unit-stride Fault-Only-First Loads

The unit-stride fault-only-first load instructions are used to vectorize loops with data-dependent exit conditions ("while" loops). These instructions execute as a regular load except that they will only take a trap caused by a synchronous exception on element 0. If an element > 0 raises an exception, that element and all following elements in the destination vector register are not modified, and the vector length `vl` is reduced to the index of the element that would have raised an exception.

```
# Vector unit-stride fault-only-first loads and stores

# vd destination, rs1 base address, vm is mask encoding (v0.t or <missing>)
vle8ff.v     vd, (rs1), vm # 8-bit unit-stride fault-only-first load
vle16ff.v    vd, (rs1), vm # 16-bit unit-stride fault-only-first load
vle32ff.v    vd, (rs1), vm # 32-bit unit-stride fault-only-first load
vle64ff.v    vd, (rs1), vm # 64-bit unit-stride fault-only-first load
vle128ff.v   vd, (rs1), vm # 128-bit unit-stride fault-only-first load
vle256ff.v   vd, (rs1), vm # 256-bit unit-stride fault-only-first load
vle512ff.v   vd, (rs1), vm # 512-bit unit-stride fault-only-first load
vle1024ff.v  vd, (rs1), vm # 1024-bit unit-stride fault-only-first load
```

strlen example using unit-stride fault-only-first instruction

link:example/strlen.s[]

Strided and scatter/gather fault-only-first instructions are not provided due to lack of encoding space, and they can also represent a larger security hole, allowing software to check multiple random pages for accessibility without experiencing a trap. The unit-stride versions only allow probing a region immediately contiguous to a known region, and so do not appreciably impact security. It is possible that security mitigations can be implemented to allow fault-only-first variants of non-contiguous accesses in future vector extensions.

Even when an exception is not raised, implementations are permitted to process fewer than `v1` elements and reduce `v1` accordingly, but if `vstart=0` and `v1>0`, then at least one element must be processed.

`v1` is not modified if element 0 raises an exception.

Implementations should not reduce `v1` and instead set a non-zero `vstart` value when the fault-only-first instruction takes a trap due to an interrupt.

7.8. Vector Load/Store Segment Instructions

This instruction subset is given the ISA string name `Zv1sseg`.

This set of instructions is included in the base "V" extension used for the Unix profile.

The vector load/store segment instructions move multiple contiguous fields in memory to and from consecutively numbered vector registers.

These operations support operations on "array-of-structures" datatypes by unpacking each field in a structure into separate vector registers.

The three-bit `nf` field in the vector instruction encoding is an unsigned integer that contains one less than the number of fields per segment, *NFIELDS*.

nf[2:0]			NFIELDS
0	0	0	1
0	0	1	2
0	1	0	3
0	1	1	4
1	0	0	5
1	0	1	6
1	1	0	7
1	1	1	8

The EMUL setting must be such that $EMUL * NFIELDS \leq 8$, otherwise the instruction encoding is reserved.

The product $EMUL * NFIELDS$ represents the number of underlying vector registers that will be touched by a segmented load or store instruction. This constraint makes this total no larger than 1/4 of the architectural register file, and the same as for regular operations with $EMUL=8$. This constraint could be weakened in a future draft.

Each field will be held in successively numbered vector register groups. When $EMUL>1$, each field will occupy a vector register group held in multiple successively numbered vector registers, and the vector register group for each field must follow the usual vector register alignment constraints (e.g., when $EMUL=2$ and $NFIELDS=4$, each field's vector register group must start at an even vector register, but does not have to start at a multiple of 8 vector register number).

An earlier version imposed a vector register number constraint, but this decreased ability to make use of all registers when $NFIELDS$ was not a power of 2.

If the vector register numbers accessed by the segment load or store would increment past 31, then the instruction encoding is reserved.

This constraint is to help provide forward-compatibility with a future longer instruction encoding that has more addressable vector registers.

The `v1` register gives the number of structures to move, which is equal to the number of elements transferred to each vector register group. Masking is also applied at the level of whole structures.

For segment loads and stores, the individual memory accesses within each segment are unordered with respect to each other even for ordered indexed segment loads and stores.

If a trap is taken, `vstart` is in units of structures. If a trap occurs partway through accessing a structure, it is implementation-defined whether a subset of the structure access is performed.

7.8.1. Vector Unit-Stride Segment Loads and Stores

The vector unit-stride load and store segment instructions move packed contiguous segments ("array-of-structures") into multiple destination vector register groups.

For segments with heterogeneous-sized fields, software can later unpack fields using additional instructions after the segment load brings the values into the separate vector registers.

The assembler prefixes `vlseg/vsseg` are used for unit-stride segment loads and stores respectively.

```
# Format
vlseg<nf>e<eew>.v vd, (rs1), vm      # Unit-stride segment load template
vsseg<nf>e<eew>.v vs3, (rs1), vm     # Unit-stride segment store template

# Examples
vlseg8e8.v vd, (rs1), vm  # Load eight vector registers with eight byte fields.

vsseg3e32.v vs3, (rs1), vm # Store packed vector of 3*4-byte segments from vs3,vs3+1,vs3+2 to mem
```

For loads, the `vd` register will hold the first field loaded from the segment. For stores, the `vs3` register is read to provide the first field to be stored in each segment.

```
# Example 1
# Memory structure holds packed RGB pixels (24-bit data structure, 8bpp)
vsetvli a1, t0, e8, ta,ma
vlseg3e8.v v8, (a0), vm
# v8 holds the red pixels
# v9 holds the green pixels
# v10 holds the blue pixels

# Example 2
# Memory structure holds complex values, 32b for real and 32b for imaginary
vsetvli a1, t0, e32, ta,ma
vlseg2e32.v v8, (a0), vm
# v8 holds real
# v9 holds imaginary
```

There are also fault-only-first versions of the unit-stride instructions.

```
# Template for vector fault-only-first unit-stride segment loads and stores.
vlseg<nf>e<eew>ff.v vd, (rs1), vm      # Unit-stride fault-only-first segment loads
```

7.8.2. Vector Strided Segment Loads and Stores

Vector strided segment loads and stores move contiguous segments where each segment is separated by the byte-stride offset given in the `rs2` GPR argument.

Negative and zero strides are supported.

```

# Format
vlsseg<nf>e<eew>.v vd, (rs1), rs2, vm      # Strided segment loads
vssseg<nf>e<eew>.v vs3, (rs1), rs2, vm     # Strided segment stores

# Examples
vsetvli a1, t0, e8, ta,ma
vlsseg3e8.v v4, (x5), x6    # Load bytes at addresses x5+i*x6 into v4[i],
                             # and bytes at addresses x5+i*x6+1 into v5[i],
                             # and bytes at addresses x5+i*x6+2 into v6[i].

# Examples
vsetvli a1, t0, e32, ta,ma
vssseg2e32.v v2, (x5), x6   # Store words from v2[i] to address x5+i*x6
                             # and words from v3[i] to address x5+i*x6+4

```

For strided segment loads and stores where the byte stride is such that segments would overlap in memory, the individual element accesses can occur in any order.

7.8.3. Vector Indexed Segment Loads and Stores

Vector indexed segment loads and stores move contiguous segments where each segment is located at an address given by adding the scalar base address in the rs1 field to byte offsets in vector register vs2. Both ordered and unordered forms are provided, where the ordered forms access segments in element order.

The data vector register group has EEW=SEW, EMUL=LMUL, while the index vector register group has EEW encoded in the instruction with EMUL=(EEW/SEW)*LMUL.

```

# Format
vluxseg<nf>ei<eew>.v vd, (rs1), vs2, vm  # Unordered indexed segment loads
vloxseg<nf>ei<eew>.v vd, (rs1), vs2, vm  # Ordered indexed segment loads
vsuxseg<nf>ei<eew>.v vs3, (rs1), vs2, vm  # Unordered indexed segment stores
vsoxseg<nf>ei<eew>.v vs3, (rs1), vs2, vm  # Ordered indexed segment stores

# Examples
vsetvli a1, t0, e8, ta,ma
vluxseg3ei32.v v4, (x5), v3    # Load bytes at addresses x5+v3[i] into v4[i],
                             # and bytes at addresses x5+v3[i]+1 into v5[i],
                             # and bytes at addresses x5+v3[i]+2 into v6[i].

# Examples
vsetvli a1, t0, e32, ta,ma
vsuxseg2ei32.v v2, (x5), v5    # Store words from v2[i] to address x5+v5[i]
                             # and words from v3[i] to address x5+v5[i]+4


```

For vector indexed segment loads, the destination vector register groups cannot overlap the source vector register group (specified by vs2), else the instruction encoding is reserved.

This constraint supports restart of indexed segment loads that raise exceptions partway through loading a structure.

7.9. Vector Load/Store Whole Register Instructions

Format for Vector Load Whole Register Instructions under LOAD-FP major opcode



These instructions load and store whole vector register groups.

These instructions are intended to be used to save and restore vector registers when the type or length of the current contents of the vector register is not known, or where modifying `v1` and `vtype` would be costly. Examples include compiler register spills, vector function calls where values are passed in vector registers, interrupt handlers, and OS context switches. Software can determine the number of bytes transferred by reading the `vlenb` register.

The load instructions have an EEW encoded in the `mew` and `width` fields following the pattern of regular unit-stride loads. Because in-register byte layouts are identical to in-memory byte layouts, these instructions all operate the same as moving vectors of bytes with EEW=8, regardless of actual EEW encoding. Pseudo-instructions are provided for whole register load instructions that correspond to EEW=8.

For the purposes of opaque save and restore of register state, the instructions have been defined as only moving byte vectors (SEW=8) between registers and memory.

The encoded EEW can be used as a HINT to indicate the destination register group will next be accessed with this EEW, which aids implementations that rearrange data internally.

The vector whole register store instructions are encoded similar to unmasked unit-stride store of elements with EEW=8.

The `nf` field encodes how many vector registers to load and store. The encoded number of registers must be a power of 2 and the vector register numbers must be aligned as with a vector register group, otherwise the instruction encoding is reserved. The `nf` field encodes the number of vector registers to transfer numbered successively after the base. Only `nf` values of 1, 2, 4, 8 are supported, with other values reserved. When multiple registers are transferred, the lowest-numbered vector register is held in the lowest-numbered memory addresses and successive vector register numbers are placed contiguously in memory.

The instructions operate with an $evl = nf * VLEN / EEW$, regardless of current settings in `vtype` and `v1`. The usual property that no elements are written if `vstart` \geq `v1` does not apply to these instructions. Instead, no elements are written if `vstart` \geq `evl`.

The instructions operate similarly to unmasked unit-stride load and store instructions of elements, with the base address passed in the scalar `x` register specified by `rs1`.

Implementations are allowed to raise a misaligned address exception on whole register loads and stores if the base address is not naturally aligned to the larger of the size of the encoded EEW in bytes (EEW/8) or the implementation's smallest supported SEW size in bytes (SEW_{MIN}/8).

Allowing misaligned exceptions to be raised based on non-alignment to encoded EEW simplifies the implementation of these instructions. Some implementations might not support smaller SEW widths, so are allowed to report misaligned exceptions for the smallest supported SEW even if larger than encoded EEW. An extreme implementation might have SEW_{MIN} > XLEN for example. Software environments can mandate the minimum alignment requirements to support an ABI. The base V extension mandates support for SEW=8.

```

# Format of whole register move instructions.
vl1r.v v3, (a0)      # Pseudo instruction equal to vl1re8.v

vl1re8.v  v3, (a0) # Load v3 with VLEN/8 bytes held at address in a0
vl1re16.v v3, (a0) # Load v3 with VLEN/16 halfwords held at address in a0
vl1re32.v v3, (a0) # Load v3 with VLEN/32 words held at address in a0
vl1re64.v v3, (a0) # Load v3 with VLEN/64 doublewords held at address in a0
vl1re128.v v3, (a0)
vl1re256.v v3, (a0)
vl1re512.v v3, (a0)
vl1re1024.v v3, (a0)

vl2r.v v2, (a0)      # Pseudo instruction equal to vl2re8.v v2, (a0)

vl2re8.v  v2, (a0) # Load v2-v3 with 2*VLEN/8 bytes from address in a0
vl2re16.v v2, (a0) # Load v2-v3 with 2*VLEN/16 halfwords held at address in a0
vl2re32.v v2, (a0) # Load v2-v3 with 2*VLEN/32 words held at address in a0
vl2re64.v v2, (a0) # Load v2-v3 with 2*VLEN/64 doublewords held at address in a0
vl2re128.v v2, (a0)
vl2re256.v v2, (a0)
vl2re512.v v2, (a0)
vl2re1024.v v2, (a0)

vl4r.v v4, (a0)      # Pseudo instruction equal to vl4re8.v

vl4re8.v  v4, (a0) # Load v4-v7 with 4*VLEN/8 bytes from address in a0
vl4re16.v v4, (a0)
vl4re32.v v4, (a0)
vl4re64.v v4, (a0)
vl4re128.v v4, (a0)
vl4re256.v v4, (a0)
vl4re512.v v4, (a0)
vl4re1024.v v4, (a0)

vl8r.v v8, (a0)      # Pseudo instruction equal to vl8re8.v

vl8re8.v  v8, (a0) # Load v8-v15 with 8*VLEN/8 bytes from address in a0
vl8re16.v v8, (a0)
vl8re32.v v8, (a0)
vl8re64.v v8, (a0)
vl8re128.v v8, (a0)
vl8re256.v v8, (a0)
vl8re512.v v8, (a0)
vl8re1024.v v8, (a0)

vs1r.v v3, (a1)      # Store v3 to address in a1
vs2r.v v2, (a1)      # Store v2-v3 to address in a1
vs4r.v v4, (a1)      # Store v4-v7 to address in a1
vs8r.v v8, (a1)      # Store v8-v15 to address in a1

```

Implementations should raise illegal instruction exceptions on `v1<nf>r` instructions for EEW values that are not supported.

These instructions can be implemented as unit-stride loads/stores of vector register groups, where EEW is 8, `nf` encodes EMUL, and `v1` = VLMAX for EEW and EMUL.

8. Vector AMO Operations


This instruction subset is given the ISA string `Zvamo`.

This set of instructions is included in the base "V" extension used for the Unix profile.

If vector AMO instructions are supported, then the scalar `Zaamo` instructions (atomic operations from the standard A extension) must be present.

Vector AMO operations are encoded using the unused width encodings under the standard AMO major opcode. Each active element performs an atomic read-modify-write of a single memory location.

Format for Vector AMO Instructions under AMO major opcode



`vs2[4:0]` specifies v register holding address

`vs3/vd[4:0]` specifies v register holding source operand and destination

`vm` specifies vector mask

`width[2:0]` specifies size of index elements, and distinguishes from scalar AMO

`amoop[4:0]` specifies the AMO operation

`wd` specifies whether the original memory value is written to `vd` (1=yes, 0=no)

The `vs2` vector register supplies the byte offset of each element, while the `vs3` vector register supplies the source data for the atomic memory operation.

AMOs have the same index EEW scheme as indexed operations, except without the `mew` bit, which is required to be zero, so offsets can have EEW=8,16,32,64 only. A vector of byte offsets in register `vs2` is added to the scalar base register in `rs1` to give the addresses of the AMO operations.

The data register `vs3` used dynamic SEW and MUL setting.

If the `wd` bit is set, the `vd` register is written with the initial value of the memory element. If the `wd` bit is clear, the `vd` register is not written.

When `wd` is clear, the memory system does not need to return the original memory value, and the original values in `vd` will be preserved.

The AMOs were defined to overwrite source data partly to reduce total memory pipeline read port count for implementations with register renaming. Also to support the same addressing mode as vector indexed operations, and because vector AMOs are less likely to need results given that the primary use is parallel in-memory reductions.

Vector AMOs operate as if `aq` and `r1` bits were zero on each element with regard to ordering relative to other instructions in the same hart.

Vector AMOs provide no ordering guarantee between element operations in the same vector AMO instruction.

Table 14. Vector AMO width encoding

	Width [2:0]			Index EEW	Mem data bits	Reg data bits	Opcode
Standard scalar AMO	0	1	0	-	32	XLEN	AMO*.W
Standard scalar AMO	0	1	1	-	64	XLEN	AMO*.D
Standard scalar AMO	1	0	0	-	128	XLEN	AMO*.Q
Vector AMO	0	0	0	8	SEW	SEW	VAMO*EI8.V
Vector AMO	1	0	1	16	SEW	SEW	VAMO*EI16.V
Vector AMO	1	1	0	32	SEW	SEW	VAMO*EI32.V
Vector AMO	1	1	1	64	SEW	SEW	VAMO*EI64.V

Index bits is the EEW of the offsets.

Mem bits is the size of element accessed in memory

Reg bits is the size of element accessed in register

If index EEW is less than XLEN, then addresses in the vector vs2 are zero-extended to XLEN. If index EEW is greater than XLEN, the instruction encoding is reserved.

Vector AMO instructions are only supported for the memory data element widths (in SEW) supported by AMOs in the implementation's scalar architecture. Other element width encodings are reserved.

The vector amoop[4:0] field uses the same encoding as the scalar 5-bit AMO instruction field, except that LR and SC are not supported.

Table 15. amoop

amoop					opcode
0	0	0	0	1	vamoswap
0	0	0	0	0	vamoadd
0	0	1	0	0	vamoxor
0	1	1	0	0	vamoand
0	1	0	0	0	vamoor
1	0	0	0	0	vamomin
1	0	1	0	0	vamomax
1	1	0	0	0	vamominu
1	1	1	0	0	vamomaxu

The assembly syntax uses x0 in the destination register position to indicate the return value is not required (wd=0).

```

# Vector AMOs for index EEW=8
vamoswapei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoswapei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoaddei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoaddei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoxorei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoxorei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoandei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoandei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoorei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoorei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominuei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominuei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxuei8.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxuei8.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

# Vector AMOs for index EEW=16
vamoswapei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoswapei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoaddei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoaddei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoxorei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoxorei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoandei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoandei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoorei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoorei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominuei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominuei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxuei16.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxuei16.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

# Vector AMOs for index EEW=32
vamoswapei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoswapei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoaddei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoaddei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

```

```

vamoxorei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoxorei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoandei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoandei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoorei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoorei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominuei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominuei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxuei32.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxuei32.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

# Vector AMOs for index EEW=64
vamoswapei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoswapei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoaddei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoaddei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoxorei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoxorei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoandei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoandei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamoorei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamoorei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamominuei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamominuei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

vamomaxuei64.v vd, (rs1), vs2, vd, v0.t # Write original value to register, wd=1
vamomaxuei64.v x0, (rs1), vs2, vs3, v0.t # Do not write original value to register, wd=0

```

9. Vector Memory Alignment Constraints

If an element accessed by a vector memory instruction is not naturally aligned to the size of the element, either the element is transferred successfully or an address misaligned exception is raised on that element.

Support for misaligned vector memory accesses is independent of an implementation's support for misaligned scalar memory accesses.

An implementation may have neither, one, or both scalar and vector memory accesses support some or all misaligned accesses in hardware. A separate PMA should be defined to determine if vector misaligned accesses are supported in the associated address range.

Vector misaligned memory accesses follow the same rules for atomicity as scalar misaligned memory accesses.

10. Vector Memory Consistency Model

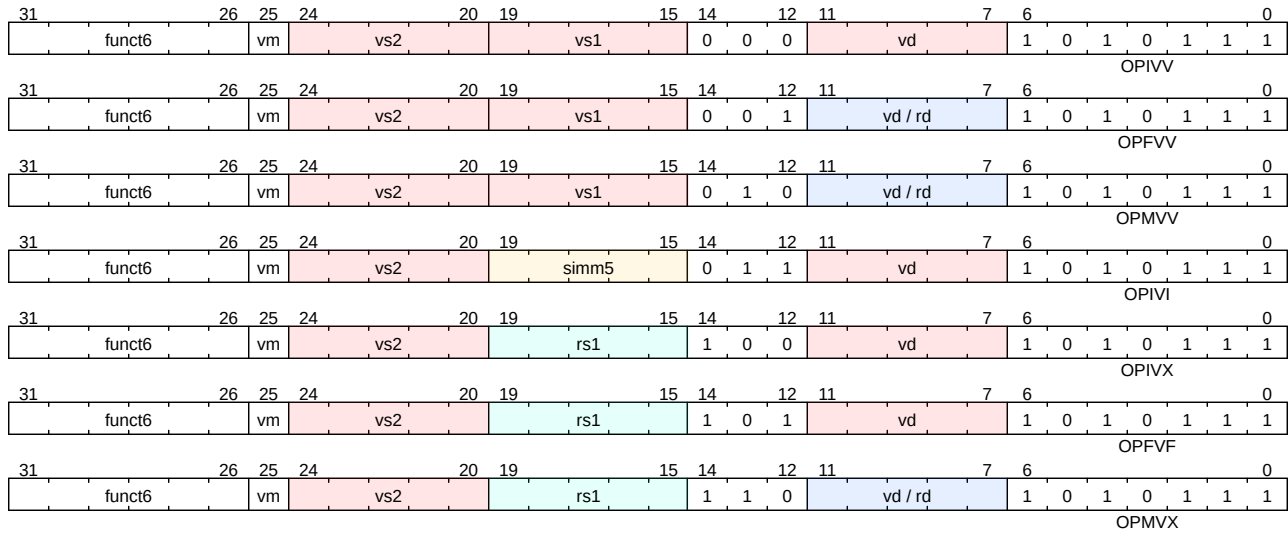
Vector memory instructions appear to execute in program order on the local hart. Vector memory instructions follow RVWMO at the instruction level, and element operations are ordered within the instruction as if performed by an element-ordered sequence of syntactically independent scalar instructions. Vector indexed-ordered stores write elements to memory in element order. Vector indexed-unordered stores do not preserve element order for writes within a single vector store instruction.

Need to flesh out details.

11. Vector Arithmetic Instruction Formats

The vector arithmetic instructions use a new major opcode (OP-V = 1010111₂) which neighbors OP-FP. The three-bit funct3 field is used to define sub-categories of vector instructions.

Formats for Vector Arithmetic Instructions under OP-V major opcode



11.1. Vector Arithmetic Instruction encoding

The funct3 field encodes the operand type and source locations.

Table 16. funct3

funct3[2:0]			Operands	Source of scalar(s)	
0	0	0	OPIVV	vector-vector	-
0	0	1	OPFVV	vector-vector	-
0	1	0	OPMVV	vector-vector	-
0	1	1	OPIVI	vector-immediate	imm[4:0]
1	0	0	OPIVX	vector-scalar	GPR x register rs1
1	0	1	OPFVF	vector-scalar	FP f register rs1
1	1	0	OPMVX	vector-scalar	GPR x register rs1
1	1	1	OPCFG	scalars-imms	GPR x register rs1 & rs2/imm

Integer operations are performed using unsigned or two's-complement signed integer arithmetic depending on the opcode.

In this discussion, fixed-point operations are considered to be integer operations.

All standard vector floating-point arithmetic operations follow the IEEE-754/2008 standard. All vector floating-point operations use the dynamic rounding mode in the frm register. Use of the frm field when it contains an invalid rounding mode by any vector floating-point instruction, even those that do not depend on the rounding mode, or when v1=0, or when vstart ≥ v1, is reserved.

All vector floating-point code will rely on a valid value in frm. Implementations can make all vector FP instructions report exceptions when the rounding mode is invalid to simplify control logic.

Vector-vector operations take two vectors of operands from vector register groups specified by vs2 and vs1 respectively.

Vector-scalar operations can have three possible forms, but in all cases take one vector of operands from a vector register group specified by vs2 and a second scalar source operand from one of three alternative sources.

1. For integer operations, the scalar can be a 5-bit immediate encoded in the rs1 field. The value is sign-extended to SEW bits, unless otherwise specified. . For integer operations, the scalar can be taken from the scalar x register specified by rs1. If XLEN>SEW, the least-significant SEW bits of the x register are used, unless otherwise specified. If XLEN<SEW, the value from the x register is sign-extended to SEW bits. . For floating-point operations, the scalar can be taken from a scalar f register. If FLEN > SEW, the value in the f registers is checked for a valid NaN-boxed value, in which case the least-significant SEW bits of the f register are used, else the canonical NaN value is used. Vector instructions where any floating-point vector operand's EEW is not a supported floating-point type width (which includes when FLEN < SEW) are reserved.

Some instructions *zero*-extend the 5-bit immediate, and denote this by naming the immediate `uimm` in the assembly syntax.

The proposed Zfinx variants will take the floating-point scalar argument from the x registers.

Vector arithmetic instructions are masked under control of the `vm` field.

Assembly syntax pattern for vector binary arithmetic instructions

Operations returning vector results, masked by `vm` (`v0.t`, `<nothing>`)

```
vop.vv vd, vs2, vs1, vm # integer vector-vector      vd[i] = vs2[i] op vs1[i]
vop.vx vd, vs2, rs1, vm # integer vector-scalar      vd[i] = vs2[i] op x[rs1]
vop.vi vd, vs2, imm, vm # integer vector-immediate   vd[i] = vs2[i] op imm
```

```
vfop.vv vd, vs2, vs1, vm # FP vector-vector operation vd[i] = vs2[i] fop vs1[i]
```

```
vfop.vf vd, vs2, rs1, vm # FP vector-scalar operation vd[i] = vs2[i] fop f[rs1]
```

In the encoding, `vs2` is the first operand, while `rs1/simm5` is the second operand. This is the opposite to the standard scalar ordering. This arrangement retains the existing encoding conventions that instructions that read only one scalar register, read it from `rs1`, and that 5-bit immediates are sourced from the `rs1` field.

Assembly syntax pattern for vector ternary arithmetic instructions (multiply-add)

Integer operations overwriting sum input

```
vop.vv vd, vs1, vs2, vm # vd[i] = vs1[i] * vs2[i] + vd[i]
vop.vx vd, rs1, vs2, vm # vd[i] = x[rs1] * vs2[i] + vd[i]
```

Integer operations overwriting product input

```
vop.vv vd, vs1, vs2, vm # vd[i] = vs1[i] * vd[i] + vs2[i]
vop.vx vd, rs1, vs2, vm # vd[i] = x[rs1] * vd[i] + vs2[i]
```

Floating-point operations overwriting sum input

```
vfop.vv vd, vs1, vs2, vm # vd[i] = vs1[i] * vs2[i] + vd[i]
vfop.vf vd, rs1, vs2, vm # vd[i] = f[rs1] * vs2[i] + vd[i]
```

Floating-point operations overwriting product input

```
vfop.vv vd, vs1, vs2, vm # vd[i] = vs1[i] * vd[i] + vs2[i]
vfop.vf vd, rs1, vs2, vm # vd[i] = f[rs1] * vd[i] + vs2[i]
```

For ternary multiply-add operations, the assembler syntax always places the destination vector register first, followed by either `rs1` or `vs1`, then `vs2`. This ordering provides a more natural reading of the assembler for these ternary operations, as the multiply operands are always next to each other.

11.2. Widening Vector Arithmetic Instructions

A few vector arithmetic instructions are defined to be *widening* operations where the destination elements have $EEW=2*SEW$ and $EMUL=2*LMUL$.

The first operand can be either single or double-width. These are generally written with a `vw*` prefix on the opcode or `vw*` for vector floating-point operations.

Assembly syntax pattern for vector widening arithmetic instructions

```
# Double-width result, two single-width sources: 2*SEW = SEW op SEW
vwop.vv  vd, vs2, vs1, vm # integer vector-vector      vd[i] = vs2[i] op vs1[i]
vwop.vx  vd, vs2, rs1, vm # integer vector-scalar      vd[i] = vs2[i] op x[rs1]

# Double-width result, first source double-width, second source single-width: 2*SEW = 2*SEW op SEW
vwop.wv  vd, vs2, vs1, vm # integer vector-vector      vd[i] = vs2[i] op vs1[i]
vwop.wx  vd, vs2, rs1, vm # integer vector-scalar      vd[i] = vs2[i] op x[rs1]
```

Originally, a *w* suffix was used on opcode, but this could be confused with the use of a *w* suffix to mean word-sized operations in doubleword integers, so the *w* was moved to prefix.

The floating-point widening operations were changed to *vfw** from *vwf** to be more consistent with any scalar widening floating-point operations that will be written as *fw**.

For integer multiply-add, another possible widening option increases the size of the accumulator to $EEW=4*SEW$ (i.e., $4*SEW \neq SEW*SEW$). These would be distinguished by a *vw4** prefix on the opcode. These are not included at this time, but are a possible addition to spec.

The destination vector register group results are arranged as if both SEW and LMUL were at twice their current settings (i.e., $EEW=2*SEW$, $EMUL=2*LMUL$).

For all widening instructions, the destination EEW and EMUL values must be a supported configuration, otherwise the instruction encoding is reserved.

The destination vector register group must be specified using a vector register number that is valid for the destination's EMUL, otherwise the instruction encoding is reserved.

This constraint is necessary to support restart with non-zero *vstart*.

For the *vw<op>.vv vd, vs2, vs1* format instructions, it is legal for *vd* to equal *vs2*.

11.3. Narrowing Vector Arithmetic Instructions

A few instructions are provided to convert double-width source vectors into single-width destination vectors. These instructions convert a vector register group with $EEW/EMUL=2*SEW/2*LMUL$ to a vector register group with the current LMUL/SEW vectors/elements.

If $EEW > ELEN$ or $EMUL > 8$, the instruction encoding is reserved.

An alternative design decision would have been to treat LMUL as defining the size of the source vector register group. The choice here is motivated by the belief the chosen approach will require fewer LMUL changes.

The source and destination vector register groups have to be specified with a vector register number that is legal for the source and destination EMUL values respectively, otherwise the instruction encoding is reserved.

Where there is a second source vector register group (specified by *vs1*), this has the same (narrower) width as the result (i.e., $EEW=SEW$).

It is safe to overwrite a second source vector register group with the same LMUL and element width as the result.

A *vn** prefix on the opcode is used to distinguish these instructions in the assembler, or a *vfn** prefix for narrowing floating-point opcodes. The double-width source vector register group is signified by a *w* in the source operand suffix (e.g., *vnsw.vv*)

Comparison operations that set a mask register are also implicitly a narrowing operation.

12. Vector Integer Arithmetic Instructions

A set of vector integer arithmetic instructions is provided.

12.1. Vector Single-Width Integer Add and Subtract

Vector integer add and subtract are provided. Reverse-subtract instructions are also provided for the vector-scalar forms.

```
# Integer adds.
vadd.vv vd, vs2, vs1, vm # Vector-vector
vadd.vx vd, vs2, rs1, vm # vector-scalar
vadd.vi vd, vs2, imm, vm # vector-immediate

# Integer subtract
vsub.vv vd, vs2, vs1, vm # Vector-vector
vsub.vx vd, vs2, rs1, vm # vector-scalar

# Integer reverse subtract
vrsb.vx vd, vs2, rs1, vm # vd[i] = rs1 - vs2[i]
vrsb.vi vd, vs2, imm, vm # vd[i] = imm - vs2[i]
```

A vector of integer values can be negated using a reverse-subtract instruction with a scalar operand of `x0`. Can define assembly pseudoinstruction `vneg.v vd, vs = vrsb.vx vd, vs, x0`.

12.2. Vector Widening Integer Add/Subtract

The widening add/subtract instructions are provided in both signed and unsigned variants, depending on whether the narrower source operands are first sign- or zero-extended before forming the double-width sum.

```
# Widening unsigned integer add/subtract, 2*SEW = SEW +/- SEW
vwaddu.vv vd, vs2, vs1, vm # vector-vector
vwaddu.vx vd, vs2, rs1, vm # vector-scalar
vwsbu.vv vd, vs2, vs1, vm # vector-vector
vwsbu.vx vd, vs2, rs1, vm # vector-scalar

# Widening signed integer add/subtract, 2*SEW = SEW +/- SEW
vwadd.vv vd, vs2, vs1, vm # vector-vector
vwadd.vx vd, vs2, rs1, vm # vector-scalar
vwsb.vv vd, vs2, vs1, vm # vector-vector
vwsb.vx vd, vs2, rs1, vm # vector-scalar

# Widening unsigned integer add/subtract, 2*SEW = 2*SEW +/- SEW
vwaddu.wv vd, vs2, vs1, vm # vector-vector
vwaddu.wx vd, vs2, rs1, vm # vector-scalar
vwsbu.wv vd, vs2, vs1, vm # vector-vector
vwsbu.wx vd, vs2, rs1, vm # vector-scalar

# Widening signed integer add/subtract, 2*SEW = 2*SEW +/- SEW
vwadd.wv vd, vs2, vs1, vm # vector-vector
vwadd.wx vd, vs2, rs1, vm # vector-scalar
vwsb.wv vd, vs2, vs1, vm # vector-vector
vwsb.wx vd, vs2, rs1, vm # vector-scalar
```

An integer value can be doubled in width using the widening add instructions with a scalar operand of `x0`. Can define assembly pseudoinstructions `vwcv.t.x.x.v vd, vs, vm = vwadd.vx vd, vs, x0, vm` and `vwcv.tu.x.x.v vd, vs, vm = vwaddu.vx vd, vs, x0, vm`.

12.3. Vector Integer Extension

The vector integer extension instructions zero- or sign-extend a source vector integer operand with EEW less than SEW to fill SEW-sized elements in the destination. The EEW of the source is 1/2, 1/4, or 1/8 of SEW, while EMUL of the source is (EEW/SEW)*LMUL. The destination has EEW equal to SEW and EMUL equal to LMUL.

```

vzext.vf2 vd, vs2, vm # Zero-extend SEW/2 source to SEW destination
vsxt.vf2 vd, vs2, vm # Sign-extend SEW/2 source to SEW destination
vzext.vf4 vd, vs2, vm # Zero-extend SEW/4 source to SEW destination
vsxt.vf4 vd, vs2, vm # Sign-extend SEW/4 source to SEW destination
vzext.vf8 vd, vs2, vm # Zero-extend SEW/8 source to SEW destination
vsxt.vf8 vd, vs2, vm # Sign-extend SEW/8 source to SEW destination

```

If the source EEW is not a supported width, or source EMUL would be below the minimum legal LMUL, the instruction encoding is reserved.

12.4. Vector Integer Add-with-Carry / Subtract-with-Borrow Instructions

To support multi-word integer arithmetic, instructions that operate on a carry bit are provided. For each operation (add or subtract), two instructions are provided: one to provide the result (SEW width), and the second to generate the carry output (single bit encoded as a mask boolean).

The carry inputs and outputs are represented using the mask register layout as described in Section [Mask Register Layout](#). Due to encoding constraints, the carry input must come from the implicit v0 register, but carry outputs can be written to any vector register that respects the source/destination overlap restrictions.

vadc and vsbc add or subtract the source operands and the carry-in or borrow-in, and write the result to vector register vd. These instructions are encoded as masked instructions (vm=0), but they operate on and write back all body elements. Encodings corresponding to the unmasked versions (vm=1) are reserved.

vmadc and vmsbc add or subtract the source operands, optionally add the carry-in or subtract the borrow-in if masked (vm=0), and write the result back to mask register vd. If unmasked (vm=1), there is no carry-in or borrow-in. These instructions operate on and write back all body elements, even if masked.

```

# Produce sum with carry.

# vd[i] = vs2[i] + vs1[i] + v0.mask[i]
vadc.vvm    vd, vs2, vs1, v0 # Vector-vector

# vd[i] = vs2[i] + x[rs1] + v0.mask[i]
vadc.vxm    vd, vs2, rs1, v0 # Vector-scalar

# vd[i] = vs2[i] + imm + v0.mask[i]
vadc.vim    vd, vs2, imm, v0 # Vector-immediate

# Produce carry out in mask register format

# vd.mask[i] = carry_out(vs2[i] + vs1[i] + v0.mask[i])
vmadc.vvm   vd, vs2, vs1, v0 # Vector-vector

# vd.mask[i] = carry_out(vs2[i] + x[rs1] + v0.mask[i])
vmadc.vxm   vd, vs2, rs1, v0 # Vector-scalar

# vd.mask[i] = carry_out(vs2[i] + imm + v0.mask[i])
vmadc.vim   vd, vs2, imm, v0 # Vector-immediate

# vd.mask[i] = carry_out(vs2[i] + vs1[i])
vmadc.vv    vd, vs2, vs1      # Vector-vector, no carry-in

# vd.mask[i] = carry_out(vs2[i] + x[rs1])
vmadc.vx    vd, vs2, rs1      # Vector-scalar, no carry-in

# vd.mask[i] = carry_out(vs2[i] + imm)
vmadc.vi    vd, vs2, imm      # Vector-immediate, no carry-in

```

Because implementing a carry propagation requires executing two instructions with unchanged inputs, destructive accumulations will require an additional move to obtain correct results.

```

# Example multi-word arithmetic sequence, accumulating into v4
vmadc.vvm v1, v4, v8, v0 # Get carry into temp register v1
vadc.vvm v4, v4, v8, v0  # Calc new sum
vmcpy.m v0, v1           # Move temp carry into v0 for next word

```

The subtract with borrow instruction vsbc performs the equivalent function to support long word arithmetic for subtraction. There are no subtract with immediate instructions.

```

# Produce difference with borrow.

# vd[i] = vs2[i] - vs1[i] - v0.mask[i]
vsbc.vvm    vd, vs2, vs1, v0 # Vector-vector

# vd[i] = vs2[i] - x[rs1] - v0.mask[i]
vsbc.vxm    vd, vs2, rs1, v0 # Vector-scalar

# Produce borrow out in mask register format

# vd.mask[i] = borrow_out(vs2[i] - vs1[i] - v0.mask[i])
vmsbc.vvm   vd, vs2, vs1, v0 # Vector-vector

# vd.mask[i] = borrow_out(vs2[i] - x[rs1] - v0.mask[i])
vmsbc.vxm   vd, vs2, rs1, v0 # Vector-scalar

# vd.mask[i] = borrow_out(vs2[i] - vs1[i])
vmsbc.vv    vd, vs2, vs1      # Vector-vector, no borrow-in

# vd.mask[i] = borrow_out(vs2[i] - x[rs1])
vmsbc.vx    vd, vs2, rs1      # Vector-scalar, no borrow-in

```

For vmsbc, the borrow is defined to be 1 iff the difference, prior to truncation, is negative.

For vadc and vsbc, the instruction encoding is reserved if the destination vector register is v0.

This constraint corresponds to the constraint on masked vector operations that overwrite the mask register.

12.5. Vector Bitwise Logical Instructions

```

# Bitwise logical operations.
vand.vv vd, vs2, vs1, vm # Vector-vector
vand.vx vd, vs2, rs1, vm # vector-scalar
vand.vi vd, vs2, imm, vm # vector-immediate

vor.vv vd, vs2, vs1, vm # Vector-vector
vor.vx vd, vs2, rs1, vm # vector-scalar
vor.vi vd, vs2, imm, vm # vector-immediate

vxor.vv vd, vs2, vs1, vm # Vector-vector
vxor.vx vd, vs2, rs1, vm # vector-scalar
vxor.vi vd, vs2, imm, vm # vector-immediate

```

With an immediate of -1, scalar-immediate forms of the vxor instruction provide a bitwise NOT operation. This can be provided as an assembler pseudoinstruction `vnot.v`.

12.6. Vector Single-Width Bit Shift Instructions

A full complement of vector shift instructions are provided, including logical shift left, and logical (zero-extending) and arithmetic (sign-extending) shift right.

```

# Bit shift operations
vsll.vv vd, vs2, vs1, vm    # Vector-vector
vsll.vx vd, vs2, rs1, vm    # vector-scalar
vsll.vi vd, vs2, uimm, vm   # vector-immediate

vsrl.vv vd, vs2, vs1, vm    # Vector-vector
vsrl.vx vd, vs2, rs1, vm    # vector-scalar
vsrl.vi vd, vs2, uimm, vm   # vector-immediate

vsra.vv vd, vs2, vs1, vm    # Vector-vector
vsra.vx vd, vs2, rs1, vm    # vector-scalar
vsra.vi vd, vs2, uimm, vm   # vector-immediate

```

The low lg2(SEW) bits of the vector or scalar shift amount value are used; immediates are zero-extended.

12.7. Vector Narrowing Integer Right Shift Instructions

The narrowing right shifts extract a smaller field from a wider operand and have both zero-extending (srl) and sign-extending (sra) forms. The shift amount can come from a vector or a scalar x register or a 5-bit immediate. The low lg2(2*SEW) bits of the vector or scalar shift amount value are used (e.g., the low 6 bits for a SEW=64-bit to SEW=32-bit narrowing operation). The immediate forms zero-extend their immediate operand.

```

# Narrowing shift right logical, SEW = (2*SEW) >> SEW
vnsrl.wv vd, vs2, vs1, vm    # vector-vector
vnsrl.wx vd, vs2, rs1, vm    # vector-scalar
vnsrl.wi vd, vs2, uimm, vm   # vector-immediate

# Narrowing shift right arithmetic, SEW = (2*SEW) >> SEW
vnsra.wv vd, vs2, vs1, vm    # vector-vector
vnsra.wx vd, vs2, rs1, vm    # vector-scalar
vnsra.wi vd, vs2, uimm, vm   # vector-immediate

```

It could be useful to add support for n4 variants, where the destination is 1/4 width of source.

An integer value can be halved in width using the narrowing integer shift instructions with a scalar operand of x0. Can define assembly pseudoinstructions `vncvt.x.x.v vd,vs,vm = vnsrl.wx vd,vs,x0,vm`.

12.8. Vector Integer Comparison Instructions

The following integer compare instructions write 1 to the destination mask register element if the comparison evaluates to true, and 0 otherwise. The destination mask vector is always held in a single vector register, with a layout of elements as described in Section [Mask Register Layout](#). The destination mask vector register may be the same as the source vector mask register (v0).

```

# Set if equal
vmsseq.vv vd, vs2, vs1, vm # Vector-vector
vmsseq.vx vd, vs2, rs1, vm # vector-scalar
vmsseq.vi vd, vs2, imm, vm # vector-immediate

# Set if not equal
vmsne.vv vd, vs2, vs1, vm # Vector-vector
vmsne.vx vd, vs2, rs1, vm # vector-scalar
vmsne.vi vd, vs2, imm, vm # vector-immediate

# Set if less than, unsigned
vmsltu.vv vd, vs2, vs1, vm # Vector-vector
vmsltu.vx vd, vs2, rs1, vm # Vector-scalar

# Set if less than, signed
vmslt.vv vd, vs2, vs1, vm # Vector-vector
vmslt.vx vd, vs2, rs1, vm # vector-scalar

# Set if less than or equal, unsigned
vmsleu.vv vd, vs2, vs1, vm # Vector-vector
vmsleu.vx vd, vs2, rs1, vm # vector-scalar
vmsleu.vi vd, vs2, imm, vm # Vector-immediate

# Set if less than or equal, signed
vmsle.vv vd, vs2, vs1, vm # Vector-vector
vmsle.vx vd, vs2, rs1, vm # vector-scalar
vmsle.vi vd, vs2, imm, vm # vector-immediate

# Set if greater than, unsigned
vmsgtu.vx vd, vs2, rs1, vm # Vector-scalar
vmsgtu.vi vd, vs2, imm, vm # Vector-immediate

# Set if greater than, signed
vmsgt.vx vd, vs2, rs1, vm # Vector-scalar
vmsgt.vi vd, vs2, imm, vm # Vector-immediate

# Following two instructions are not provided directly
# Set if greater than or equal, unsigned
# vmsgeu.vx vd, vs2, rs1, vm # Vector-scalar
# Set if greater than or equal, signed
# vmsge.vx vd, vs2, rs1, vm # Vector-scalar

```

The following table indicates how all comparisons are implemented in native machine code.

Comparison	Assembler Mapping	Assembler Pseudoinstruction
<code>va < vb</code>	<code>vmslt{u}.vv vd, va, vb, vm</code>	
<code>va <= vb</code>	<code>vmsle{u}.vv vd, va, vb, vm</code>	
<code>va > vb</code>	<code>vmslt{u}.vv vd, vb, va, vm</code>	<code>vmsgt{u}.vv vd, va, vb, vm</code>
<code>va >= vb</code>	<code>vmsle{u}.vv vd, vb, va, vm</code>	<code>vmsge{u}.vv vd, va, vb, vm</code>
<code>va < x</code>	<code>vmslt{u}.vx vd, va, x, vm</code>	
<code>va <= x</code>	<code>vmsle{u}.vx vd, va, x, vm</code>	
<code>va > x</code>	<code>vmsgt{u}.vx vd, va, x, vm</code>	
<code>va >= x</code>	see below	
<code>va < i</code>	<code>vmsle{u}.vi vd, va, i-1, vm</code>	<code>vmslt{u}.vi vd, va, i, vm</code>
<code>va <= i</code>	<code>vmsle{u}.vi vd, va, i, vm</code>	
<code>va > i</code>	<code>vmsgt{u}.vi vd, va, i, vm</code>	
<code>va >= i</code>	<code>vmsgt{u}.vi vd, va, i-1, vm</code>	<code>vmsge{u}.vi vd, va, i, vm</code>
va, vb vector register groups		
x scalar integer register		
i immediate		

The immediate forms of `vmslt{u}.vi` are not provided as the immediate value can be decreased by 1 and the `vmsle{u}.vi` variants used instead. The `vmsle.vi` range is -16 to 15, resulting in an effective `vmslt.vi` range of -15 to 16. The `vmsleu.vi` range is 0 to 15 giving an effective `vmsltu.vi` range of 1 to 16 (Note, `vmsltu.vi` with immediate 0 is not useful as it is always false). Because the 5-bit vector immediates are always sign-extended, `vmsleu.vi` also supports unsigned immediate values in the range $2^{SEW}-16$ to $2^{SEW}-1$, allowing corresponding `vmsltu.vi` comparisons against unsigned immediates in the range $2^{SEW}-15$ to 2^{SEW} . Note that `vlsltu.vi` with immediate 2^{SEW} is not useful as it is always true.

Similarly, `vmsge{u}.vi` is not provided and the comparison is implemented using `vmsgt{u}.vi` with the immediate decremented by one. The resulting effective `vmsge.vi` range is -15 to 16, and the resulting effective `vmsgeu.vi` range is 1 to 16 (Note, `vmsgeu.vi` with immediate 0 is not useful as it is always true).

The `vmsgt` forms for register scalar and immediates are provided to allow a single comparison instruction to provide the correct polarity of mask value without using additional mask logical instructions.

To reduce encoding space, the `vmsge{u}.vx` form is not directly provided, and so the `va ≥ x` case requires special treatment.

The `vmsge{u}.vx` could potentially be encoded in a non-orthogonal way under the unused OPIVI variant of `vmslt{u}`. These would be the only instructions in OPIVI that use a scalar ‘x’ register however. Alternatively, a further two `funct6` encodings could be used, but these would have a different operand format (writes to mask register) than others in the same group of 8 `funct6` encodings. The current PoR is to omit these instructions and to synthesize where needed as described below.

The `vmsge{u}.vx` operation can be synthesized by reducing the value of x by 1 and using the `vmsgt{u}.vx` instruction, when it is known that this will not underflow the representation in x.

Sequences to synthesize ‘`vmsge{u}.vx`’ instruction

`va >= x, x > minimum`

```
addi t0, x, -1; vmsgt{u}.vx vd, va, t0, vm
```

The above sequence will usually be the most efficient implementation, but assembler pseudoinstructions can be provided for cases where the range of x is unknown.

unmasked va >= x

```
pseudoinstruction: vmsge{u}.vx vd, va, x
expansion: vmslt{u}.vx vd, va, x; vmnand.mm vd, vd, vd
```

masked va >= x, vd != v0

```
pseudoinstruction: vmsge{u}.vx vd, va, x, v0.t
expansion: vmslt{u}.vx vd, va, x, v0.t; vmxor.mm vd, vd, v0
```

masked va >= x, vd == v0

```
pseudoinstruction: vmsge{u}.vx vd, va, x, v0.t, vt
expansion: vmslt{u}.vx vt, va, x; vmmandnot.mm vd, vd, vt
```

masked va >= x, any vd

```
pseudoinstruction: vmsge{u}.vx vd, va, x, v0.t, vt
expansion: vmslt{u}.vx vt, va, x; vmmandnot.mm vt, v0, vt; vmmandnot.mm vd, vd, v0; vmor.mm vd, vt,
```

The vt argument to the pseudoinstruction must name a temporary vector register that is not same as vd and which will be clobbered by the pseudoinstruction

Comparisons effectively AND in the mask, e.g,

```
# (a < b) && (b < c) in two instructions
vmslt.vv    v0, va, vb          # All body elements written
vmslt.vv    v0, vb, vc, v0.t    # Only update at set mask
```

12.9. Vector Integer Min/Max Instructions

Signed and unsigned integer minimum and maximum instructions are supported.

```
# Unsigned minimum
vminu.vv vd, vs2, vs1, vm    # Vector-vector
vminu.vx vd, vs2, rs1, vm    # vector-scalar

# Signed minimum
vmin.vv vd, vs2, vs1, vm     # Vector-vector
vmin.vx vd, vs2, rs1, vm     # vector-scalar

# Unsigned maximum
vmaxu.vv vd, vs2, vs1, vm    # Vector-vector
vmaxu.vx vd, vs2, rs1, vm    # vector-scalar

# Signed maximum
vmax.vv vd, vs2, vs1, vm     # Vector-vector
vmax.vx vd, vs2, rs1, vm     # vector-scalar
```

12.10. Vector Single-Width Integer Multiply Instructions

The single-width multiply instructions perform a SEW-bit*SEW-bit multiply and return an SEW-bit-wide result. The **mulh** versions write the high word of the product to the destination register.

```

# Signed multiply, returning low bits of product
vmul.vv vd, vs2, vs1, vm  # Vector-vector
vmul.vx vd, vs2, rs1, vm  # vector-scalar

# Signed multiply, returning high bits of product
vmulh.vv vd, vs2, vs1, vm  # Vector-vector
vmulh.vx vd, vs2, rs1, vm  # vector-scalar

# Unsigned multiply, returning high bits of product
vmulhu.vv vd, vs2, vs1, vm  # Vector-vector
vmulhu.vx vd, vs2, rs1, vm  # vector-scalar

# Signed(vs2)-Unsigned multiply, returning high bits of product
vmulhsu.vv vd, vs2, vs1, vm  # Vector-vector
vmulhsu.vx vd, vs2, rs1, vm  # vector-scalar

```

There is no vmulhus opcode to return high half of unsigned-vector * signed-scalar product.

The current vmulh* opcodes perform simple fractional multiplies, but with no option to scale, round, and/or saturate the result. Can consider changing definition of vmulh, vmulhu, vmulhsu to use vxrm rounding mode when discarding low half of product. There is no possibility of overflow in this case.

12.11. Vector Integer Divide Instructions

The divide and remainder instructions are equivalent to the RISC-V standard scalar integer multiply/divides, with the same results for extreme inputs.

```

# Unsigned divide.
vdivu.vv vd, vs2, vs1, vm  # Vector-vector
vdivu.vx vd, vs2, rs1, vm  # vector-scalar

# Signed divide
vdiv.vv vd, vs2, vs1, vm  # Vector-vector
vdiv.vx vd, vs2, rs1, vm  # vector-scalar

# Unsigned remainder
vremu.vv vd, vs2, vs1, vm  # Vector-vector
vremu.vx vd, vs2, rs1, vm  # vector-scalar

# Signed remainder
vrem.vv vd, vs2, vs1, vm  # Vector-vector
vrem.vx vd, vs2, rs1, vm  # vector-scalar

```

The decision to include integer divide and remainder was contentious. The argument in favor is that without a standard instruction, software would have to pick some algorithm to perform the operation, which would likely perform poorly on some microarchitectures versus others.

There is no instruction to perform a "scalar divide by vector" operation.

12.12. Vector Widening Integer Multiply Instructions

The widening integer multiply instructions return the full 2*SEW-bit product from an SEW-bit*SEW-bit multiply.

```
# Widening signed-integer multiply
vwmul.vv vd, vs2, vs1, vm # vector-vector
vwmul.vx vd, vs2, rs1, vm # vector-scalar
```

```
# Widening unsigned-integer multiply
vwmulu.vv vd, vs2, vs1, vm # vector-vector
vwmulu.vx vd, vs2, rs1, vm # vector-scalar
```

```
# Widening signed-unsigned integer multiply
vwmulsu.vv vd, vs2, vs1, vm # vector-vector
vwmulsu.vx vd, vs2, rs1, vm # vector-scalar
```

12.13. Vector Single-Width Integer Multiply-Add Instructions

The integer multiply-add instructions are destructive and are provided in two forms, one that overwrites the addend or minuend (vmacc, vnmsac) and one that overwrites the first multiplicand (vmadd, vnmsub).

The low half of the product is added or subtracted from the third operand.

"sac" is intended to be read as "subtract from accumulator". The opcode is "vnmsac" to match the (unfortunately counterintuitive) floating-point fnmsub instruction definition. Similarly for the "vnmsub" opcode.

```
# Integer multiply-add, overwrite addend
vmacc.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vmacc.vx vd, rs1, vs2, vm # vd[i] = +(x[rs1] * vs2[i]) + vd[i]
```

```
# Integer multiply-sub, overwrite minuend
vnmsac.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vs2[i]) + vd[i]
vnmsac.vx vd, rs1, vs2, vm # vd[i] = -(x[rs1] * vs2[i]) + vd[i]
```

```
# Integer multiply-add, overwrite multiplicand
vmadd.vv vd, vs1, vs2, vm # vd[i] = (vs1[i] * vd[i]) + vs2[i]
vmadd.vx vd, rs1, vs2, vm # vd[i] = (x[rs1] * vd[i]) + vs2[i]
```

```
# Integer multiply-sub, overwrite multiplicand
vnmsub.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vd[i]) + vs2[i]
vnmsub.vx vd, rs1, vs2, vm # vd[i] = -(x[rs1] * vd[i]) + vs2[i]
```

12.14. Vector Widening Integer Multiply-Add Instructions

The widening integer multiply-add instructions add a SEW-bit*SEW-bit multiply result to (from) a 2*SEW-bit value and produce a 2*SEW-bit result. All combinations of signed and unsigned multiply operands are supported.

```
# Widening unsigned-integer multiply-add, overwrite addend
vwmaccu.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vwmaccu.vx vd, rs1, vs2, vm # vd[i] = +(x[rs1] * vs2[i]) + vd[i]
```

```
# Widening signed-integer multiply-add, overwrite addend
vwmacc.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vwmacc.vx vd, rs1, vs2, vm # vd[i] = +(x[rs1] * vs2[i]) + vd[i]
```

```
# Widening signed-unsigned-integer multiply-add, overwrite addend
vwmaccsu.vv vd, vs1, vs2, vm # vd[i] = +(signed(vs1[i]) * unsigned(vs2[i])) + vd[i]
vwmaccsu.vx vd, rs1, vs2, vm # vd[i] = +(signed(x[rs1]) * unsigned(vs2[i])) + vd[i]
```

```
# Widening unsigned-signed-integer multiply-add, overwrite addend
vwmaccus.vx vd, rs1, vs2, vm # vd[i] = +(unsigned(x[rs1]) * signed(vs2[i])) + vd[i]
```

12.15. Vector Integer Merge Instructions

The vector integer merge instructions combine two source operands based on a mask. Unlike regular arithmetic instructions, the merge operates on all body elements (i.e., the set of elements from `vstart` up to the current vector length in `v1`).

The `vmerge` instructions are always masked (`vm=0`). The instructions combine two sources as follows. At elements where the mask value is zero, the first operand is copied to the destination element, otherwise the second operand is copied to the destination element. The first operand is always a vector register group specified by `vs2`. The second operand is a vector register group specified by `vs1` or a scalar x register specified by `rs1` or a 5-bit sign-extended immediate.

```
vmerge.vvm vd, vs2, vs1, v0 # vd[i] = v0.mask[i] ? vs1[i] : vs2[i]
vmerge.vxm vd, vs2, rs1, v0 # vd[i] = v0.mask[i] ? x[rs1] : vs2[i]
vmerge.vim vd, vs2, imm, v0 # vd[i] = v0.mask[i] ? imm : vs2[i]
```

12.16. Vector Integer Move Instructions

The vector integer move instructions copy a source operand to a vector register group. The `vmv.v.v` variant copies a vector register group, whereas the `vmv.v.x` and `vmv.v.i` variants *splat* a scalar register or immediate to all active elements of the destination vector register group. These instructions are always unmasked (`vm=1`). The first operand specifier (`vs2`) must contain `v0`, and any other vector register number in `vs2` is *reserved*.

```
vmv.v.v vd, vs1 # vd[i] = vs1[i]
vmv.v.x vd, rs1 # vd[i] = rs1
vmv.v.i vd, imm # vd[i] = imm
```

Mask values can be widened into SEW-width elements using a sequence `vmv.v.i vd, 0; vmerge.vim vd, vd, 1, v0`.

The vector integer move instructions share the encoding with the vector merge instructions, but with `vm=1` and `vs2=v0`.

The form `vmv.v.v vd, vd`, which leaves body elements unchanged, is used as a hint to indicate that the register will be used with an EEW equal to SEW. Implementations that internally reorganize data according to EEW can shuffle the internal representation according to SEW. Implementations that do not internally reorganize data can dynamically elide this instruction, and treat as a NOP.

13. Vector Fixed-Point Arithmetic Instructions

The preceding set of integer arithmetic instructions is extended to support fixed-point arithmetic.

A fixed-point number is a two's-complement signed or unsigned integer interpreted as the numerator in a fraction with an implicit denominator. The fixed-point instructions are intended to be applied to the numerators; it is the responsibility of software to manage the denominators. An N-bit element can hold two's-complement signed integers in the range $-2^{N-1} \dots +2^{N-1}-1$, and unsigned integers in the range $0 \dots +2^N-1$. The fixed-point instructions help preserve precision in narrow operands by supporting scaling and rounding, and can handle overflow by saturating results into the destination format range.

The widening integer operations described above can also be used to remove the possibility of overflow.

13.1. Vector Single-Width Saturating Add and Subtract

Saturating forms of integer add and subtract are provided, for both signed and unsigned integers. If the result would overflow the destination, the result is replaced with the closest representable value, and the `vxsat` bit is set.

```
# Saturating adds of unsigned integers.
vsaddu.vv vd, vs2, vs1, vm # Vector-vector
vsaddu.vx vd, vs2, rs1, vm # vector-scalar
vsaddu.vi vd, vs2, imm, vm # vector-immediate

# Saturating adds of signed integers.
vsadd.vv vd, vs2, vs1, vm # Vector-vector
vsadd.vx vd, vs2, rs1, vm # vector-scalar
vsadd.vi vd, vs2, imm, vm # vector-immediate

# Saturating subtract of unsigned integers.
vssubu.vv vd, vs2, vs1, vm # Vector-vector
vssubu.vx vd, vs2, rs1, vm # vector-scalar

# Saturating subtract of signed integers.
vssub.vv vd, vs2, vs1, vm # Vector-vector
vssub.vx vd, vs2, rs1, vm # vector-scalar
```

13.2. Vector Single-Width Averaging Add and Subtract

The averaging add and subtract instructions right shift the result by one bit and round off the result according to the setting in `vxrm`. Both unsigned and signed versions are provided. For `vaaddu`, `vaadd`, and `vasub`, there can be no overflow in the result. For `vasubu`, overflow is ignored.

```

# Averaging add

# Averaging adds of unsigned integers.
vaaddu.vv vd, vs2, vs1, vm # roundoff_unsigned(vs2[i] + vs1[i], 1)
vaaddu.vx vd, vs2, rs1, vm # roundoff_unsigned(vs2[i] + x[rs1], 1)

# Averaging adds of signed integers.
vaadd.vv vd, vs2, vs1, vm # roundoff_signed(vs2[i] + vs1[i], 1)
vaadd.vx vd, vs2, rs1, vm # roundoff_signed(vs2[i] + x[rs1], 1)

# Averaging subtract

# Averaging subtract of unsigned integers.
vasubu.vv vd, vs2, vs1, vm # roundoff_unsigned(vs2[i] - vs1[i], 1)
vasubu.vx vd, vs2, rs1, vm # roundoff_unsigned(vs2[i] - x[rs1], 1)

# Averaging subtract of signed integers.
vasub.vv vd, vs2, vs1, vm # roundoff_signed(vs2[i] - vs1[i], 1)
vasub.vx vd, vs2, rs1, vm # roundoff_signed(vs2[i] - x[rs1], 1)

```

13.3. Vector Single-Width Fractional Multiply with Rounding and Saturation

The signed fractional multiply instruction produces a $2 \times \text{SEW}$ product of the two SEW inputs, then shifts the result right by SEW-1 bits, rounding these bits according to vxrm, then saturates the result to fit into SEW bits. If the result causes saturation, the vxsat bit is set.

```

# Signed saturating and rounding fractional multiply
# See vxrm description for rounding calculation
vsmul.vv vd, vs2, vs1, vm # vd[i] = clip(roundoff_signed(vs2[i]*vs1[i], SEW-1))
vsmul.vx vd, vs2, rs1, vm # vd[i] = clip(roundoff_signed(vs2[i]*x[rs1], SEW-1))

```

When multiplying two N-bit signed numbers, the largest magnitude is obtained for $-2^{N-1} * -2^{N-1}$ producing a result $+2^{2N-2}$, which has a single (zero) sign bit when held in 2N bits. All other products have two sign bits in 2N bits. To retain greater precision in N result bits, the product is shifted right by one bit less than N, saturating the largest magnitude result but increasing result precision by one bit for all other products.

13.4. Vector Single-Width Scaling Shift Instructions

These instructions shift the input value right, and round off the shifted out bits according to vxrm. The scaling right shifts have both zero-extending (vssrl) and sign-extending (vssra) forms. The low $\lg_2(\text{SEW})$ bits of the vector or scalar shift amount value are used; immediates are zero-extended.

```

# Scaling shift right logical
vssrl.vv vd, vs2, vs1, vm # vd[i] = roundoff_unsigned(vs2[i], vs1[i])
vssrl.vx vd, vs2, rs1, vm # vd[i] = roundoff_unsigned(vs2[i], x[rs1])
vssrl.vi vd, vs2, uimm, vm # vd[i] = roundoff_unsigned(vs2[i], uimm)

# Scaling shift right arithmetic
vssra.vv vd, vs2, vs1, vm # vd[i] = roundoff_signed(vs2[i], vs1[i])
vssra.vx vd, vs2, rs1, vm # vd[i] = roundoff_signed(vs2[i], x[rs1])
vssra.vi vd, vs2, uimm, vm # vd[i] = roundoff_signed(vs2[i], uimm)

```

13.5. Vector Narrowing Fixed-Point Clip Instructions

The vnclip instructions are used to pack a fixed-point value into a narrower destination. The instructions support rounding, scaling, and saturation into the final destination format.

The second argument (vector element, scalar value, immediate value) gives the amount to right shift the source as in the narrowing shift instructions, which provides the scaling. The low $\lg_2(2 \times \text{SEW})$ bits of the vector or scalar shift amount value

are used (e.g., the low 6 bits for a SEW=64-bit to SEW=32-bit narrowing operation). The immediate forms zero-extend their immediate operand.

```
# Narrowing unsigned clip
#
#          SEW          2*SEW  SEW
vnclipu.wv vd, vs2, vs1, vm # vd[i] = clip(roundoff_unsigned(vs2[i], vs1[i]))
vnclipu.wx vd, vs2, rs1, vm # vd[i] = clip(roundoff_unsigned(vs2[i], x[rs1]))
vnclipu.wi vd, vs2, uimm, vm # vd[i] = clip(roundoff_unsigned(vs2[i], uimm5))

# Narrowing signed clip
vnclip.wv vd, vs2, vs1, vm # vd[i] = clip(roundoff_signed(vs2[i], vs1[i]))
vnclip.wx vd, vs2, rs1, vm # vd[i] = clip(roundoff_signed(vs2[i], x[rs1]))
vnclip.wi vd, vs2, uimm, vm # vd[i] = clip(roundoff_signed(vs2[i], uimm5))
```

For vnclipu/vnclip, the rounding mode is specified in the vxrm CSR. Rounding occurs around the least-significant bit of the destination and before saturation.

For vnclipu, the shifted rounded source value is treated as an unsigned integer and saturates if the result would overflow the destination viewed as an unsigned integer.

For vnclip, the shifted rounded source value is treated as a signed integer and saturates if the result would overflow the destination viewed as a signed integer.

If any destination element is saturated, the vxsat bit is set in the vxsat register.

14. Vector Floating-Point Instructions

The standard vector floating-point instructions treat 16-bit, 32-bit, 64-bit, and 128-bit elements as IEEE-754/2008-compatible values. If the EEW of a vector floating-point operand does not correspond to a supported IEEE floating-point type, the instruction encoding is reserved.

The floating-point element widths that are supported depend on the profile.

Vector floating-point instructions require the presence of base scalar floating-point extensions corresponding to the supported vector floating-point element widths.

Profiles supporting 16-bit half-precision floating-point values will also have to implement scalar half-precision floating-point support in the f registers.

If the floating-point unit status field `mstatus.FS` is `Off` then any attempt to execute a vector floating-point instruction will raise an illegal instruction exception. Any vector floating-point instruction that modifies any floating-point extension state (i.e., floating-point CSRs or f registers) must set `mstatus.FS` to `Dirty`.

The vector floating-point instructions have the same behavior as the scalar floating-point instructions with regard to NaNs.

Scalar values for vector-scalar operations can be sourced from the standard scalar f registers, as described in Section [Vector Arithmetic Instruction encoding](#).

14.1. Vector Floating-Point Exception Flags

A vector floating-point exception at any active floating-point element sets the standard FP exception flags in the `fflags` register. Inactive elements do not set FP exception flags.

14.2. Vector Single-Width Floating-Point Add/Subtract Instructions

```
# Floating-point add
vfadd.vv vd, vs2, vs1, vm # Vector-vector
vfadd.vf vd, vs2, rs1, vm # vector-scalar

# Floating-point subtract
vfsub.vv vd, vs2, vs1, vm # Vector-vector
vfsub.vf vd, vs2, rs1, vm # Vector-scalar vd[i] = vs2[i] - f[rs1]
vfsub.vf vd, vs2, rs1, vm # Scalar-vector vd[i] = f[rs1] - vs2[i]
```

14.3. Vector Widening Floating-Point Add/Subtract Instructions

```
# Widening FP add/subtract, 2*SEW = SEW +/- SEW
vfwadd.vv vd, vs2, vs1, vm # vector-vector
vfwadd.vf vd, vs2, rs1, vm # vector-scalar
vfwsb.vv vd, vs2, vs1, vm # vector-vector
vfwsb.vf vd, vs2, rs1, vm # vector-scalar

# Widening FP add/subtract, 2*SEW = 2*SEW +/- SEW
vfwadd.wv vd, vs2, vs1, vm # vector-vector
vfwadd.wf vd, vs2, rs1, vm # vector-scalar
vfwsb.wv vd, vs2, vs1, vm # vector-vector
vfwsb.wf vd, vs2, rs1, vm # vector-scalar
```

14.4. Vector Single-Width Floating-Point Multiply/Divide Instructions

```

# Floating-point multiply
vfmul.vv vd, vs2, vs1, vm # Vector-vector
vfmul.vf vd, vs2, rs1, vm # vector-scalar

# Floating-point divide
vfdiv.vv vd, vs2, vs1, vm # Vector-vector
vfdiv.vf vd, vs2, rs1, vm # vector-scalar

# Reverse floating-point divide vector = scalar / vector
vfrdiv.vf vd, vs2, rs1, vm # scalar-vector, vd[i] = f[rs1]/vs2[i]

```

14.5. Vector Widening Floating-Point Multiply

```

# Widening floating-point multiply
vfwmul.vv vd, vs2, vs1, vm # vector-vector
vfwmul.vf vd, vs2, rs1, vm # vector-scalar

```

14.6. Vector Single-Width Floating-Point Fused Multiply-Add Instructions

All four varieties of fused multiply-add are provided, and in two destructive forms that overwrite one of the operands, either the addend or the first multiplicand.

```

# FP multiply-accumulate, overwrites addend
vfmac.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vfmac.vf vd, rs1, vs2, vm # vd[i] = +(f[rs1] * vs2[i]) + vd[i]

# FP negate-(multiply-accumulate), overwrites subtrahend
vfnmac.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vs2[i]) - vd[i]
vfnmac.vf vd, rs1, vs2, vm # vd[i] = -(f[rs1] * vs2[i]) - vd[i]

# FP multiply-subtract-accumulator, overwrites subtrahend
vfmsac.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vs2[i]) - vd[i]
vfmsac.vf vd, rs1, vs2, vm # vd[i] = +(f[rs1] * vs2[i]) - vd[i]

# FP negate-(multiply-subtract-accumulator), overwrites minuend
vfnmsac.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vs2[i]) + vd[i]
vfnmsac.vf vd, rs1, vs2, vm # vd[i] = -(f[rs1] * vs2[i]) + vd[i]

# FP multiply-add, overwrites multiplicand
vfmad.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vd[i]) + vs2[i]
vfmad.vf vd, rs1, vs2, vm # vd[i] = +(f[rs1] * vd[i]) + vs2[i]

# FP negate-(multiply-add), overwrites multiplicand
vfnmad.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vd[i]) - vs2[i]
vfnmad.vf vd, rs1, vs2, vm # vd[i] = -(f[rs1] * vd[i]) - vs2[i]

# FP multiply-sub, overwrites multiplicand
vfmsub.vv vd, vs1, vs2, vm # vd[i] = +(vs1[i] * vd[i]) - vs2[i]
vfmsub.vf vd, rs1, vs2, vm # vd[i] = +(f[rs1] * vd[i]) - vs2[i]

# FP negate-(multiply-sub), overwrites multiplicand
vfnmsub.vv vd, vs1, vs2, vm # vd[i] = -(vs1[i] * vd[i]) + vs2[i]
vfnmsub.vf vd, rs1, vs2, vm # vd[i] = -(f[rs1] * vd[i]) + vs2[i]

```

It would be possible to use the two unused rounding modes in the scalar FP FMA encoding to provide a few non-destructive FMAs. However, this would be the only maskable operation with three inputs and separate output.

14.7. Vector Widening Floating-Point Fused Multiply-Add Instructions

The widening floating-point fused multiply-add instructions all overwrite the wide addend with the result. The multiplier inputs are all SEW wide, while the addend and destination is 2*SEW bits wide.

```
# FP widening multiply-accumulate, overwrites addend
vfwmac.vv vd, vs1, vs2, vm    # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vfwmac.vf vd, rs1, vs2, vm    # vd[i] = +(f[rs1] * vs2[i]) + vd[i]

# FP widening negate-(multiply-accumulate), overwrites addend
vfwnmacc.vv vd, vs1, vs2, vm  # vd[i] = -(vs1[i] * vs2[i]) - vd[i]
vfwnmacc.vf vd, rs1, vs2, vm  # vd[i] = -(f[rs1] * vs2[i]) - vd[i]

# FP widening multiply-subtract-accumulator, overwrites addend
vfwmsac.vv vd, vs1, vs2, vm   # vd[i] = +(vs1[i] * vs2[i]) - vd[i]
vfwmsac.vf vd, rs1, vs2, vm   # vd[i] = +(f[rs1] * vs2[i]) - vd[i]

# FP widening negate-(multiply-subtract-accumulator), overwrites addend
vfwnmsac.vv vd, vs1, vs2, vm  # vd[i] = -(vs1[i] * vs2[i]) + vd[i]
vfwnmsac.vf vd, rs1, vs2, vm  # vd[i] = -(f[rs1] * vs2[i]) + vd[i]
```

14.8. Vector Floating-Point Square-Root Instruction

This is a unary vector-vector instruction.

```
# Floating-point square root
vfsqrt.v vd, vs2, vm    # Vector-vector square root
```

14.9. Vector Floating-Point Reciprocal Square-Root Estimate Instruction

```
# Floating-point reciprocal square-root estimate to 7 bits.
vfrsqste7.v vd, vs2, vm
```

This is a unary vector-vector instruction that returns an estimate of $1/\sqrt{x}$ accurate to 7 bits.

The following table describes the instruction's behavior for all classes of floating-point inputs:

Input	Output	Exceptions raised
$-\infty \leq x < -0.0$	canonical NaN	NV
-0.0	$-\infty$	DZ
+0.0	$+\infty$	DZ
$+0.0 < x < +\infty$	<i>estimate of $1/\sqrt{x}$</i>	
$+\infty$	+0.0	
qNaN	canonical NaN	
sNaN	canonical NaN	NV

All positive normal and subnormal inputs produce normal outputs.

The output value is independent of the dynamic rounding mode.

For the non-exceptional cases, the low bit of the exponent and the six high bits of significand (after the leading one) are concatenated and used to address the following table. The output of the table becomes the seven high bits of the result significand (after the leading one); the remainder of the result significand is zero. Subnormal inputs are normalized and the exponent adjusted appropriately before the lookup. The output exponent is chosen to make the result approximate the reciprocal of the square root of the argument.

More precisely, the result is computed as follows. Let the normalized input exponent be equal to the input exponent if the input is normal, or 0 minus the number of leading zeros in the significand otherwise. If the input is subnormal, the normalized input significand is given by shifting the input significand left by 1 minus the normalized input exponent,

discarding the leading 1 bit. The output exponent equals $\text{floor}((3*B - 1 - \text{the normalized input exponent}) / 2)$. The output sign equals the input sign.

The following table gives the seven MSBs of the output significand as a function of the LSB of the normalized input exponent and the six MSBs of the normalized input significand; the other bits of the output significand are zero.

Table 17. vfrsqtrte7.v common-case lookup table contents

exp[0]	sig[MSB -: 6]	sig_out[MSB -: 7]
0	0	52
	1	51
	2	50
	3	48
	4	47
	5	46
	6	44
	7	43
	8	42
	9	41
	10	40
	11	39
	12	38
	13	36
	14	35
	15	34
	16	33
	17	32
	18	31
	19	30
	20	30
	21	29
	22	28
	23	27
	24	26
	25	25
	26	24
	27	23
	28	23
	29	22
	30	21
	31	20
	32	19
	33	19
	34	18
	35	17
	36	16
	37	16
	38	15
	39	14
	40	14
	41	13
	42	12
	43	12
	44	11

1	45	10
	46	10
	47	9
	48	9
	49	8
	50	7
	51	7
	52	6
	53	6
	54	5
	55	4
	56	4
	57	3
	58	3
	59	2
	60	2
	61	1
	62	1
	63	0
	0	127
	1	125
	2	123
	3	121
	4	119
	5	118
	6	116
	7	114
	8	113
	9	111
	10	109
	11	108
	12	106
	13	105
	14	103
	15	102
	16	100
	17	99
	18	97
	19	96
	20	95
	21	93
	22	92
	23	91
	24	90
	25	88
	26	87
	27	86
	28	85

29	84
30	83
31	82
32	80
33	79
34	78
35	77
36	76
37	75
38	74
39	73
40	72
41	71
42	70
43	70
44	69
45	68
46	67
47	66
48	65
49	64
50	63
51	63
52	62
53	61
54	60
55	59
56	59
57	58
58	57
59	56
60	56
61	55
62	54
63	53

For example, when $SEW=32$, $vfrsqste7(0x00718abc \approx 1.043e-38) = 0x5f080000 \approx 9.800e18$, and $vfrsqste7(0x7f765432 \approx 3.274e38) = 0x1f820000 \approx 5.506e-20$.

The 7 bit accuracy was chosen as it requires 0,1,2,3 Newton-Raphson iterations to converge to close to bfloat16, FP16, FP32, FP64 accuracy respectively. Future instructions can be defined with greater estimate accuracy.

14.10. Vector Floating-Point Reciprocal Estimate Instruction

```
# Floating-point reciprocal estimate to 7 bits.
vfrece7.v vd, vs2, vm
```

This is a unary vector-vector instruction that returns an estimate of $1/x$ accurate to 7 bits.

The following table describes the instruction's behavior for all classes of floating-point inputs, where B is the exponent bias:

Input (x)	Rounding Mode	Output ($y \approx 1/x$)	Exceptions raised
$-\infty$	<i>any</i>	-0.0	
$-2^{B+1} < x \leq -2^B$ (normal)	<i>any</i>	$-2^{-(B+1)} \geq y > -2^{-B}$ (subnormal, sig=01...)	
$-2^B < x \leq -2^{B-1}$ (normal)	<i>any</i>	$-2^{-B} \geq y > -2^{-B+1}$ (subnormal, sig=1...)	
$-2^{B-1} < x \leq -2^{-B+1}$ (normal)	<i>any</i>	$-2^{-B+1} \geq y > -2^{B-1}$ (normal)	
$-2^{-B+1} < x \leq -2^{-B}$ (subnormal, sig=1...)	<i>any</i>	$-2^{B-1} \geq y > -2^B$ (normal)	
$-2^{-B} < x \leq -2^{-(B+1)}$ (subnormal, sig=01...)	<i>any</i>	$-2^B \geq y > -2^{B+1}$ (normal)	
$-2^{-(B+1)} < x < -0.0$ (subnormal, sig=00...)	RUP, RTZ	greatest-mag. negative finite value	NX, OF
$-2^{-(B+1)} < x < -0.0$ (subnormal, sig=00...)	RDN, RNE, RMM	$-\infty$	NX, OF
-0.0	<i>any</i>	$-\infty$	DZ
+0.0	<i>any</i>	$+\infty$	DZ
$+0.0 < x < 2^{-(B+1)}$ (subnormal, sig=00...)	RUP, RNE, RMM	$+\infty$	NX, OF
$+0.0 < x < 2^{-(B+1)}$ (subnormal, sig=00...)	RDN, RTZ	greatest finite value	NX, OF
$2^{-(B+1)} \leq x < 2^{-B}$ (subnormal, sig=01...)	<i>any</i>	$2^{B+1} > y \geq 2^B$ (normal)	
$2^{-B} \leq x < 2^{-B+1}$ (subnormal, sig=1...)	<i>any</i>	$2^B > y \geq 2^{B-1}$ (normal)	
$2^{-B+1} \leq x < 2^{B-1}$ (normal)	<i>any</i>	$2^{B-1} > y \geq 2^{-B+1}$ (normal)	
$2^{B-1} \leq x < 2^B$ (normal)	<i>any</i>	$2^{-B+1} > y \geq 2^{-B}$ (subnormal, sig=1...)	
$2^B \leq x < 2^{B+1}$ (normal)	<i>any</i>	$2^{-B} > y \geq 2^{-(B+1)}$ (subnormal, sig=01...)	
$+\infty$	<i>any</i>	+0.0	
qNaN	<i>any</i>	canonical NaN	
sNaN	<i>any</i>	canonical NaN	NV

Subnormal inputs with magnitude at least $2^{-(B+1)}$ produce normal outputs; other subnormal inputs produce infinite outputs. Normal inputs with magnitude at least 2^{B-1} produce subnormal outputs; other normal inputs produce normal outputs.

The output value depends on the dynamic rounding mode when the overflow exception is raised.

For the non-exceptional cases, the seven high bits of significand (after the leading one) are used to address the following table. The output of the table becomes the seven high bits of the result significand (after the leading one); the remainder of the result significand is zero. Subnormal inputs are normalized and the exponent adjusted appropriately before the lookup. The output exponent is chosen to make the result approximate the reciprocal of the argument, and subnormal outputs are denormalized accordingly.

More precisely, the result is computed as follows. Let the normalized input exponent be equal to the input exponent if the input is normal, or 0 minus the number of leading zeros in the significand otherwise. The normalized output exponent equals $(2*B - 1 - \text{the normalized input exponent})$. If the normalized output exponent is outside the range $[-1, 2*B]$, the result corresponds to one of the exceptional cases in the table above.

If the input is subnormal, the normalized input significand is given by shifting the input significand left by 1 minus the normalized input exponent, discarding the leading 1 bit. Otherwise, the normalized input significand equals the input significand. The following table gives the seven MSBs of the normalized output significand as a function of the seven MSBs of the normalized input significand; the other bits of the normalized output significand are zero.

Table 18. vfrece7.v common-case lookup table contents

sig[MSB -: 7]	sig_out[MSB -: 7]
0	127
1	125
2	123
3	121
4	119
5	117
6	116
7	114
8	112
9	110
10	109
11	107
12	105
13	104
14	102
15	100
16	99
17	97
18	96
19	94
20	93
21	91
22	90
23	88
24	87
25	85
26	84
27	83
28	81
29	80
30	79
31	77
32	76
33	75
34	74
35	72
36	71
37	70
38	69
39	68
40	66
41	65
42	64
43	63
44	62

45	61
46	60
47	59
48	58
49	57
50	56
51	55
52	54
53	53
54	52
55	51
56	50
57	49
58	48
59	47
60	46
61	45
62	44
63	43
64	42
65	41
66	40
67	40
68	39
69	38
70	37
71	36
72	35
73	35
74	34
75	33
76	32
77	31
78	31
79	30
80	29
81	28
82	28
83	27
84	26
85	25
86	25
87	24
88	23
89	23
90	22
91	21
92	21

93	20
94	19
95	19
96	18
97	17
98	17
99	16
100	15
101	15
102	14
103	14
104	13
105	12
106	12
107	11
108	11
109	10
110	9
111	9
112	8
113	8
114	7
115	7
116	6
117	5
118	5
119	4
120	4
121	3
122	3
123	2
124	2
125	1
126	1
127	0

If the normalized output exponent is 0 or -1, the result is subnormal: the output exponent is 0, and the output significand is given by concatenating a 1 bit to the left of the normalized output significand, then shifting that quantity right by 1 minus the normalized output exponent. Otherwise, the output exponent equals the normalized output exponent, and the output significand equals the normalized output significand. The output sign equals the input sign.

For example, when SEW=32, $\text{vfrece7}(0x00718abc \approx 1.043e-38) = 0x7e900000 \approx 9.570e37$, and $\text{vfrece7}(0x7f765432 \approx 3.274e38) = 0x00214000 \approx 3.053e-39$.

The 7 bit accuracy was chosen as it requires 0,1,2,3 Newton-Raphson iterations to converge to close to bfloat16, FP16, FP32, FP64 accuracy respectively. Future instructions can be defined with greater estimate accuracy.

14.11. Vector Floating-Point MIN/MAX Instructions

The vector floating-point `vfmin` and `vfmax` instructions have the same behavior as the corresponding scalar floating-point instructions in version 2.2 of the RISC-V F/D/Q extension.

```

# Floating-point minimum
vfmmin.vv vd, vs2, vs1, vm # Vector-vector
vfmmin.vf vd, vs2, rs1, vm # vector-scalar

# Floating-point maximum
vfmax.vv vd, vs2, vs1, vm # Vector-vector
vfmax.vf vd, vs2, rs1, vm # vector-scalar

```

14.12. Vector Floating-Point Sign-Injection Instructions

Vector versions of the scalar sign-injection instructions. The result takes all bits except the sign bit from the vector `vs2` operands.

```

vfsgnj.vv vd, vs2, vs1, vm # Vector-vector
vfsgnj.vf vd, vs2, rs1, vm # vector-scalar

vfsgnjn.vv vd, vs2, vs1, vm # Vector-vector
vfsgnjn.vf vd, vs2, rs1, vm # vector-scalar

vfsgnjx.vv vd, vs2, vs1, vm # Vector-vector
vfsgnjx.vf vd, vs2, rs1, vm # vector-scalar

```

A vector of floating-point values can be negated using a sign-injection instruction with both source operands set to the same vector operand. Can define assembly pseudoinstruction `vneg.v vd, vs = vfsgnjn.vv vd, vs, vs`.

14.13. Vector Floating-Point Compare Instructions

These vector FP compare instructions compare two source operands and write the comparison result to a mask register. The destination mask vector is always held in a single vector register, with a layout of elements as described in Section [Mask Register Layout](#). The destination mask vector register may be the same as the source vector mask register (`v0`).

The compare instructions follow the semantics of the scalar floating-point compare instructions. `vmfeq` and `vmfne` raise the invalid operation exception only on signaling NaN inputs. `vmflt`, `vmfle`, `vmfgt`, and `vmfge` raise the invalid operation exception on both signaling and quiet NaN inputs. `vmfne` writes 1 to the destination element when either operand is NaN, whereas the other comparisons write 0 when either operand is NaN.

```

# Compare equal
vmfeq.vv vd, vs2, vs1, vm # Vector-vector
vmfeq.vf vd, vs2, rs1, vm # vector-scalar

# Compare not equal
vmfne.vv vd, vs2, vs1, vm # Vector-vector
vmfne.vf vd, vs2, rs1, vm # vector-scalar

# Compare less than
vmflt.vv vd, vs2, vs1, vm # Vector-vector
vmflt.vf vd, vs2, rs1, vm # vector-scalar

# Compare less than or equal
vmfle.vv vd, vs2, vs1, vm # Vector-vector
vmfle.vf vd, vs2, rs1, vm # vector-scalar

# Compare greater than
vmfgt.vf vd, vs2, rs1, vm # vector-scalar

# Compare greater than or equal
vmfge.vf vd, vs2, rs1, vm # vector-scalar

```

Comparison	Assembler Mapping	Assembler pseudoinstruction
<code>va < vb</code>	<code>vmflt.vv vd, va, vb, vm</code>	
<code>va <= vb</code>	<code>vmfle.vv vd, va, vb, vm</code>	
<code>va > vb</code>	<code>vmflt.vv vd, vb, va, vm</code>	<code>vmfgt.vv vd, va, vb, vm</code>
<code>va >= vb</code>	<code>vmfle.vv vd, vb, va, vm</code>	<code>vmfge.vv vd, va, vb, vm</code>
<code>va < f</code>	<code>vmflt.vf vd, va, f, vm</code>	
<code>va <= f</code>	<code>vmfle.vf vd, va, f, vm</code>	
<code>va > f</code>	<code>vmfgt.vf vd, va, f, vm</code>	
<code>va >= f</code>	<code>vmfge.vf vd, va, f, vm</code>	
<code>va, vb</code>	vector register groups	
<code>f</code>	scalar floating-point register	

Providing all forms is necessary to correctly handle unordered comparisons for NaNs.

C99 floating-point quiet comparisons can be implemented by masking the signaling comparisons when either input is NaN, as follows. When the comparand is a non-NaN constant, the middle two instructions can be omitted.

```
# Example of implementing isgreater()
vmfeq.vv v0, va, va      # Only set where A is not NaN.
vmfeq.vv v1, vb, vb      # Only set where B is not NaN.
vmand.mm v0, v0, v1      # Only set where A and B are ordered,
vmfgt.vv v0, va, vb, v0.t # so only set flags on ordered values.
```

In the above sequence, it is tempting to mask the second `vmfeq` instruction and remove the `vmand` instruction, but this more efficient sequence incorrectly fails to raise the invalid exception when an element of `va` contains a quiet NaN and the corresponding element in `vb` contains a signaling NaN.

14.14. Vector Floating-Point Classify Instruction

This is a unary vector-vector instruction that operates in the same way as the scalar classify instruction.

```
vfclass.v vd, vs2, vm    # Vector-vector
```

The 10-bit mask produced by this instruction is placed in the least-significant bits of the result elements. The upper (SEW-10) bits of the result are filled with zeros. The instruction is only defined for SEW=16b and above, so the result will always fit in the destination elements.

14.15. Vector Floating-Point Merge Instruction

A vector-scalar floating-point merge instruction is provided, which operates on all body elements, from `vstart` up to the current vector length in `v1` regardless of mask value.

The `vfmerge.vfm` instruction is always masked (`vm=0`). At elements where the mask value is zero, the first vector operand is copied to the destination element, otherwise a scalar floating-point register value is copied to the destination element.

```
vfmerge.vfm vd, vs2, rs1, v0 # vd[i] = v0.mask[i] ? f[rs1] : vs2[i]
```

14.16. Vector Floating-Point Move Instruction

The vector floating-point move instruction *splats* a floating-point scalar operand to a vector register group. The instruction copies a scalar `f` register value to all active elements of a vector register group. This instruction is always unmasked (`vm=1`). The instruction must have the `vs2` field set to `v0`, with all other values for `vs2` reserved.

```
vfmv.v.f vd, rs1 # vd[i] = f[rs1]
```

The `vfmv.v.f` instruction shares the encoding with the `vfmerge.vfm` instruction, but with `vm=1` and `vs2=v0`.

14.17. Single-Width Floating-Point/Integer Type-Convert Instructions

Conversion operations are provided to convert to and from floating-point values and unsigned and signed integers, where both source and destination are SEW wide.

<code>vfcvt.xu.f.v vd, vs2, vm</code>	# Convert float to unsigned integer.
<code>vfcvt.x.f.v vd, vs2, vm</code>	# Convert float to signed integer.
<code>vfcvt.rtz.xu.f.v vd, vs2, vm</code>	# Convert float to unsigned integer, truncating.
<code>vfcvt.rtz.x.f.v vd, vs2, vm</code>	# Convert float to signed integer, truncating.
<code>vfcvt.f.xu.v vd, vs2, vm</code>	# Convert unsigned integer to float.
<code>vfcvt.f.x.v vd, vs2, vm</code>	# Convert signed integer to float.

The conversions follow the same rules on exceptional conditions as the scalar conversion instructions. The conversions use the dynamic rounding mode in `frm`, except for the `rtz` variants, which round towards zero.

The `rtz` variants are provided to accelerate truncating conversions from floating-point to integer, as is common in languages like C and Java.

14.18. Widening Floating-Point/Integer Type-Convert Instructions

A set of conversion instructions is provided to convert between narrower integer and floating-point datatypes to a type of twice the width.

<code>vfwcvt.xu.f.v vd, vs2, vm</code>	# Convert float to double-width unsigned integer.
<code>vfwcvt.x.f.v vd, vs2, vm</code>	# Convert float to double-width signed integer.
<code>vfwcvt.rtz.xu.f.v vd, vs2, vm</code>	# Convert float to double-width unsigned integer, truncating.
<code>vfwcvt.rtz.x.f.v vd, vs2, vm</code>	# Convert float to double-width signed integer, truncating.
<code>vfwcvt.f.xu.v vd, vs2, vm</code>	# Convert unsigned integer to double-width float.
<code>vfwcvt.f.x.v vd, vs2, vm</code>	# Convert signed integer to double-width float.
<code>vfwcvt.f.f.v vd, vs2, vm</code>	# Convert single-width float to double-width float.

These instructions have the same constraints on vector register overlap as other widening instructions (see [Widening Vector Arithmetic Instructions](#)).

A double-width IEEE floating-point value can always represent a single-width integer exactly.

A double-width IEEE floating-point value can always represent a single-width IEEE floating-point value exactly.

A full set of floating-point widening conversions is not supported as single instructions, but any widening conversion can be implemented as several doubling steps with equivalent results and no additional exception flags raised.

14.19. Narrowing Floating-Point/Integer Type-Convert Instructions

A set of conversion instructions is provided to convert wider integer and floating-point datatypes to a type of half the width.

<code>vfncvt.xu.f.w vd, vs2, vm</code>	# Convert double-width float to unsigned integer.
<code>vfncvt.x.f.w vd, vs2, vm</code>	# Convert double-width float to signed integer.
<code>vfncvt.rtz.xu.f.w vd, vs2, vm</code>	# Convert double-width float to unsigned integer, truncating.
<code>vfncvt.rtz.x.f.w vd, vs2, vm</code>	# Convert double-width float to signed integer, truncating.
<code>vfncvt.f.xu.w vd, vs2, vm</code>	# Convert double-width unsigned integer to float.
<code>vfncvt.f.x.w vd, vs2, vm</code>	# Convert double-width signed integer to float.
<code>vfncvt.f.f.w vd, vs2, vm</code>	# Convert double-width float to single-width float.
<code>vfncvt.rod.f.f.w vd, vs2, vm</code>	# Convert double-width float to single-width float, # rounding towards odd.

These instructions have the same constraints on vector register overlap as other narrowing instructions (see [Narrowing Vector Arithmetic Instructions](#)).

A full set of floating-point widening conversions is not supported as single instructions. Conversions can be implemented in a sequence of halving steps. Results are equivalently rounded and the same exception flags are raised if all but the last halving step use round-towards-odd (`vfncvt.rod.f.f.w`). Only the final step should use the desired rounding mode.

An integer value can be halved in width using the narrowing integer shift instructions with a shift amount of 0.

15. Vector Reduction Operations

Vector reduction operations take a vector register group of elements and a scalar held in element 0 of a vector register, and perform a reduction using some binary operator, to produce a scalar result in element 0 of a vector register. The scalar input and output operands are held in element 0 of a single vector register, not a vector register group, so any vector register can be the scalar source or destination of a vector reduction regardless of LMUL setting.

The destination vector register can overlap the source operands, including the mask register.

Reductions read and write the scalar operand and result into element 0 of a vector register to avoid a loss of decoupling with the scalar processor, and to support future polymorphic use with future types not supported in the scalar unit.

Inactive elements from the source vector register group are excluded from the reduction, but the scalar operand is always included regardless of the mask values.

The other elements in the destination vector register ($0 < \text{index} < \text{VLEN}/\text{SEW}$) are left unchanged.

If $v1=0$, no operation is performed and the destination register is not updated.

Traps on vector reduction instructions are always reported with a `vstart` of 0. Vector reduction operations raise an illegal instruction exception if `vstart` is non-zero.

The assembler syntax for a reduction operation is `vredop.vs`, where the `.vs` suffix denotes the first operand is a vector register group and the second operand is a scalar stored in element 0 of a vector register.

15.1. Vector Single-Width Integer Reduction Instructions

All operands and results of single-width reduction instructions have the same SEW width. Overflows wrap around on arithmetic sums.

```
# Simple reductions, where [*] denotes all active elements:
vredsum.vs  vd, vs2, vs1, vm  # vd[0] = sum( vs1[0] , vs2[*] )
vredmaxu.vs vd, vs2, vs1, vm  # vd[0] = maxu( vs1[0] , vs2[*] )
vredmax.vs  vd, vs2, vs1, vm  # vd[0] = max( vs1[0] , vs2[*] )
vredminu.vs vd, vs2, vs1, vm  # vd[0] = minu( vs1[0] , vs2[*] )
vredmin.vs  vd, vs2, vs1, vm  # vd[0] = min( vs1[0] , vs2[*] )
vredand.vs  vd, vs2, vs1, vm  # vd[0] = and( vs1[0] , vs2[*] )
vredor.vs   vd, vs2, vs1, vm  # vd[0] = or( vs1[0] , vs2[*] )
vredxor.vs  vd, vs2, vs1, vm  # vd[0] = xor( vs1[0] , vs2[*] )
```

15.2. Vector Widening Integer Reduction Instructions

The unsigned `vwredsumu.vs` instruction zero-extends the SEW-wide vector elements before summing them, then adds the $2*\text{SEW}$ -width scalar element, and stores the result in a $2*\text{SEW}$ -width scalar element.

The `vwredsum.vs` instruction sign-extends the SEW-wide vector elements before summing them.

```
# Unsigned sum reduction into double-width accumulator
vwredsumu.vs vd, vs2, vs1, vm  #  $2*\text{SEW} = 2*\text{SEW} + \text{sum}(\text{zero-extend}(\text{SEW}))$ 

# Signed sum reduction into double-width accumulator
vwredsum.vs  vd, vs2, vs1, vm  #  $2*\text{SEW} = 2*\text{SEW} + \text{sum}(\text{sign-extend}(\text{SEW}))$ 
```

15.3. Vector Single-Width Floating-Point Reduction Instructions

```
# Simple reductions.
vfredosum.vs vd, vs2, vs1, vm # Ordered sum
vfredsum.vs  vd, vs2, vs1, vm # Unordered sum
vfredmax.vs  vd, vs2, vs1, vm # Maximum value
vfredmin.vs  vd, vs2, vs1, vm # Minimum value
```

15.3.1. Vector Ordered Single-Width Floating-Point Sum Reduction

The `vfredosum` instruction must sum the floating-point values in element order, starting with the scalar in `vs1[0]`--that is, it performs the computation: $((vs1[0] + vs2[0]) + vs2[1]) + \dots + vs2[v1-1]$, where each addition operates identically to the scalar floating-point instructions in terms of raising exception flags and generating or propagating special values.

The ordered reduction supports compiler autovectorization, while the unordered FP sum allows for faster implementations.

When the operation is masked (`vm=0`), the masked-off elements do not affect the result or the exception flags.

If no elements are active, no additions are performed, so the scalar in `vs1[0]` is simply copied to the destination register, without canonicalizing NaN values and without setting any exception flags. This behavior preserves the handling of NaNs, exceptions, and rounding when autovectorizing a scalar summation loop.

15.3.2. Vector Unordered Single-Width Floating-Point Sum Reduction

The unordered sum reduction instruction, `vfredsum`, provides an implementation more freedom in performing the reduction.

The implementation can produce a result equivalent to a reduction tree composed of binary operator nodes, with the inputs being elements from the source vector register group (`vs2`) and the source scalar value (`vs1[0]`). Each operator in the tree accepts two inputs and produces one result. Each operator first computes an exact sum as a RISC-V scalar floating-point addition with infinite exponent range and precision, then converts this exact sum to a floating-point format with range and precision each at least as great as the element floating-point format indicated by SEW, rounding using the currently active floating-point dynamic rounding mode. A different floating-point range and precision may be chosen for the result of each operator. A node where one input is derived only from elements masked-off or beyond the active vector length may either treat that input as the additive identity of the appropriate EEW or simply copy the other input to its output. The rounded result from the root node in the tree is converted (rounded again, using the dynamic rounding mode) to the standard floating-point format indicated by SEW. An implementation is allowed to add an additional additive identity to the final result.

The additive identity is $+0.0$ when rounding down (towards $-\infty$) or -0.0 for all other rounding modes.

The reduction tree structure must be deterministic for a given value in `vtype` and `v1`.

As a consequence of this definition, implementations need not propagate NaN payloads through the reduction tree when no elements are active. In particular, if no elements are active and the scalar input is NaN, implementations are permitted to canonicalize the NaN and, if the NaN is signaling, set the invalid exception flag. Implementations are alternatively permitted to pass through the original NaN and set no exception flags, as with `vfredosum`.

The `vfredosum` instruction is a valid implementation of the `vfredsum` instruction.

15.3.3. Vector Single-Width Floating Max and Min Reductions

Floating-point max and min reductions should return the same final value and raise the same exception flags regardless of operation order.

If no elements are active, the scalar in `vs1[0]` is simply copied to the destination register, without canonicalizing NaN values and without setting any exception flags.

15.4. Vector Widening Floating-Point Reduction Instructions

Widening forms of the sum reductions are provided that read and write a double-width reduction result.

```
# Simple reductions.  
vfwredosum.vs vd, vs2, vs1, vm # Ordered sum  
vfwredsum.vs vd, vs2, vs1, vm # Unordered sum
```

The reduction of the SEW-width elements is performed as in the single-width reduction case, with the elements in vs2 promoted to 2*SEW bits before adding to the 2*SEW-bit accumulator.

vfwredosum.vs handles inactive elements and NaN payloads analogously to vfredosum.vs; vfwredsum.vs does so analogously to vfredsum.vs.

16. Vector Mask Instructions

Several instructions are provided to help operate on mask values held in a vector register.

16.1. Vector Mask-Register Logical Instructions

Vector mask-register logical operations operate on mask registers. Each element in a mask register is a single bit, so these instructions all operate on single vector registers regardless of the setting of the `vlmul` field in `vtype`. They do not change the value of `vlmul`. The destination vector register may be the same as either source vector register.

As with other vector instructions, the elements with indices less than `vstart` are unchanged, and `vstart` is reset to zero after execution. Vector mask logical instructions are always unmasked so there are no inactive elements. Mask elements past `vl`, the tail elements, are handled according to the setting of `vta` in `vtype` (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)).

```
vmmand.mm vd, vs2, vs1    # vd.mask[i] = vs2.mask[i] && vs1.mask[i]
vmnand.mm vd, vs2, vs1    # vd.mask[i] = !(vs2.mask[i] && vs1.mask[i])
vmandnot.mm vd, vs2, vs1  # vd.mask[i] = vs2.mask[i] && !vs1.mask[i]
vmxor.mm  vd, vs2, vs1    # vd.mask[i] = vs2.mask[i] ^^ vs1.mask[i]
vmor.mm   vd, vs2, vs1    # vd.mask[i] = vs2.mask[i] || vs1.mask[i]
vmnor.mm  vd, vs2, vs1    # vd.mask[i] = !(vs2.mask[i] || vs1.mask[i])
vmornot.mm vd, vs2, vs1  # vd.mask[i] = vs2.mask[i] || !vs1.mask[i]
vmxnor.mm vd, vs2, vs1    # vd.mask[i] = !(vs2.mask[i] ^^ vs1.mask[i])
```

Several assembler pseudoinstructions are defined as shorthand for common uses of mask logical operations:

```
vmmv.m vd, vs => vmmand.mm vd, vs, vs # Copy mask register
vmclr.m vd    => vmxor.mm  vd, vd, vd  # Clear mask register
vmset.m vd    => vmxnor.mm vd, vd, vd  # Set mask register
vmnot.m vd, vs=> vmnand.mm vd, vs, vs  # Invert bits
```

The `vmmv.m` instruction was previously called `vmcpy.m`, but with new layout it is more consistent to name as a "mv" because bits are copied without interpretation. The `vmcpy.m` assembler pseudo-instruction can be retained for compatibility.

The set of eight mask logical instructions can generate any of the 16 possibly binary logical functions of the two input masks:

inputs				
0	0	1	1	src1
0	1	0	1	src2

output				instruction	pseudoinstruction
0	0	0	0	vmxor.mm vd, vd, vd	vmclr.m vd
1	0	0	0	vmnor.mm vd, src1, src2	
0	1	0	0	vmandnot.mm vd, src2, src1	
1	1	0	0	vmnand.mm vd, src1, src1	vmnot.m vd, src1
0	0	1	0	vmandnot.mm vd, src1, src2	
1	0	1	0	vmnand.mm vd, src2, src2	vmnot.m vd, src2
0	1	1	0	vmxor.mm vd, src1, src2	
1	1	1	0	vmnand.mm vd, src1, src2	
0	0	0	1	vmand.mm vd, src1, src2	
1	0	0	1	vmxnor.mm vd, src1, src2	
0	1	0	1	vmand.mm vd, src2, src2	vmcpy.m vd, src2
1	1	0	1	vmornot.mm vd, src2, src1	
0	0	1	1	vmand.mm vd, src1, src1	vmcpy.m vd, src1
1	0	1	1	vmornot.mm vd, src1, src2	
1	1	1	1	vmxnor.mm vd, vd, vd	vmset.m vd

The vector mask logical instructions are designed to be easily fused with a following masked vector operation to effectively expand the number of predicate registers by moving values into v0 before use.

16.2. Vector mask population count vpopc

```
vpopc.m rd, vs2, vm
```

The source operand is a single vector register holding mask register values as described in Section [Mask Register Layout](#).

The vpopc.m instruction counts the number of mask elements of the active elements of the vector source mask register that have the value 1 and writes the result to a scalar x register.

The operation can be performed under a mask, in which case only the masked elements are counted.

```
vpopc.m rd, vs2, v0.t # x[rd] = sum_i ( vs2.mask[i] && v0.mask[i] )
```

Traps on vpopc.m are always reported with a vstart of 0. The vpopc instruction will raise an illegal instruction exception if vstart is non-zero.

16.3. vfirst find-first-set mask bit

```
vfirst.m rd, vs2, vm
```

The vfirst instruction finds the lowest-numbered active element of the source mask vector that has the value 1 and writes that element's index to a GPR. If no active element has the value 1, -1 is written to the GPR.

Software can assume that any negative value (highest bit set) corresponds to no element found, as vector lengths will never exceed $2^{(XLEN-1)}$ on any implementation.

Traps on vfirst are always reported with a vstart of 0. The vfirst instruction will raise an illegal instruction exception if vstart is non-zero.

16.4. vmsbf.m set-before-first mask bit

vmsbf.m vd, vs2, vm

Example

7	6	5	4	3	2	1	0	Element number
1	0	0	1	0	1	0	0	v3 contents
vmsbf.m v2, v3								
0	0	0	0	0	0	1	1	v2 contents
1	0	0	1	0	1	0	1	v3 contents
vmsbf.m v2, v3								
0	0	0	0	0	0	0	0	v2
0	0	0	0	0	0	0	0	v3 contents
vmsbf.m v2, v3								
1	1	1	1	1	1	1	1	v2
1	1	0	0	0	0	1	1	v0 vcontents
1	0	0	1	0	1	0	0	v3 contents
vmsbf.m v2, v3, v0.t								
0	1	x	x	x	x	1	1	v2 contents

The `vmsbf.m` instruction takes a mask register as input and writes results to a mask register. The instruction writes a 1 to all active mask elements before the first source element that is a 1, then writes a 0 to that element and all following active elements. If there is no set bit in the source vector, then all active elements in the destination are written with a 1.

The tail elements in the destination mask register are handled according to the setting of the `vta` bit in `vtype` (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)).

Traps on `vmsbf.m` are always reported with a `vstart` of 0. The `vmsbf` instruction will raise an illegal instruction exception if `vstart` is non-zero.

The destination register cannot overlap the source register and, if masked, cannot overlap the mask register ('v0').

16.5. vmsif.m set-including-first mask bit

The vector mask set-including-first instruction is similar to set-before-first, except it also includes the element with a set bit.

vmsif.m vd, vs2, vm

Example

7	6	5	4	3	2	1	0	Element number
1	0	0	1	0	1	0	0	v3 contents
vmsif.m v2, v3								
0	0	0	0	0	1	1	1	v2 contents
1	0	0	1	0	1	0	1	v3 contents
vmsif.m v2, v3								
0	0	0	0	0	0	0	1	v2
1	1	0	0	0	0	1	1	v0 vcontents
1	0	0	1	0	1	0	0	v3 contents
vmsif.m v2, v3, v0.t								
1	1	x	x	x	x	1	1	v2 contents

The tail elements in the destination mask register are handled according to the setting of the `vta` bit in `vtype` (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)).

Traps on `vmsif.m` are always reported with a `vstart` of 0. The `vmsif` instruction will raise an illegal instruction exception if `vstart` is non-zero.

The destination register cannot overlap the source register and, if masked, cannot overlap the mask register ('v0').

16.6. `vmsof.m` set-only-first mask bit

The vector mask set-only-first instruction is similar to set-before-first, except it only sets the first element with a bit set, if any.

```
vmsof.m vd, vs2, vm
```

Example

7	6	5	4	3	2	1	0	Element number
1	0	0	1	0	1	0	0	v3 contents vmsof.m v2, v3
0	0	0	0	0	1	0	0	v2 contents
1	0	0	1	0	1	0	1	v3 contents vmsof.m v2, v3
0	0	0	0	0	0	0	1	v2
1	1	0	0	0	0	1	1	v0 vcontents
1	1	0	1	0	1	0	0	v3 contents vmsof.m v2, v3, v0.t
0	1	x	x	x	x	0	0	v2 contents

The tail elements in the destination mask register are handled according to the setting of the `vta` bit in `vtype` (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)).

Traps on `vmsof.m` are always reported with a `vstart` of 0. The `vmsof` instruction will raise an illegal instruction exception if `vstart` is non-zero.

The destination register cannot overlap the source register and, if masked, cannot overlap the mask register ('v0').

16.7. Example using vector mask instructions

The following is an example of vectorizing a data-dependent exit loop.

```
link:example/stncpy.s[]
```

```
link:example/strncpy.s[]
```

16.8. Vector Iota Instruction

The `viota.m` instruction reads a source vector mask register and writes to each element of the destination vector register group the sum of all the bits of elements in the mask register whose index is less than the element, e.g., a parallel prefix sum of the mask values.

This instruction can be masked, in which case only the enabled elements contribute to the sum and only the enabled elements are written.

`viota.m vd, vs2, vm`

Example

7	6	5	4	3	2	1	0	Element number
1	0	0	1	0	0	0	1	v2 contents
viota.m v4, v2 # Unmasked								
2	2	2	1	1	1	1	0	v4 result
1	1	1	0	1	0	1	1	v0 contents
1	0	0	1	0	0	0	1	v2 contents
2	3	4	5	6	7	8	9	v4 contents
viota.m v4, v2, v0.t # Masked								
1	1	1	5	1	7	1	0	v4 results

The result value is zero-extended to fill the destination element if SEW is wider than the result. If the result value would overflow the destination SEW, the least-significant SEW bits are retained.

Traps on `viota.m` are always reported with a `vstart` of 0, and execution is always restarted from the beginning when resuming after a trap handler. An illegal instruction exception is raised if `vstart` is non-zero.

The destination register group cannot overlap the source register and, if masked, cannot overlap the mask register ('v0').

These constraints exist for two reasons. First, to simplify avoidance of WAR hazards in implementations with temporally long vector registers and no vector register renaming. Second, to enable resuming execution after a trap simpler.

The `viota.m` instruction can be combined with memory scatter instructions (indexed stores) to perform vector compress functions.

```

# Compact non-zero elements from input memory array to output memory array
#
# size_t compact_non_zero(size_t n, const int* in, int* out)
# {
#     size_t i;
#     size_t count = 0;
#     int *p = out;
#
#     for (i=0; i<n; i++)
#     {
#         const int v = *in++;
#         if (v != 0)
#             *p++ = v;
#     }
#
#     return (size_t) (p - out);
# }
#
# a0 = n
# a1 = &in
# a2 = &out

compact_non_zero:
    li a6, 0                # Clear count of non-zero elements
loop:
    vsetvli a5, a0, e32,m8,ta,ma # 32-bit integers
    vle32.v v8, (a1)           # Load input vector
    sub a0, a0, a5             # Decrement number done
    slli a5, a5, 2             # Multiply by four bytes
    vmsne.vi v0, v8, 0         # Locate non-zero values
    add a1, a1, a5             # Bump input pointer
    vpopc.m a5, v0             # Count number of elements set in v0
    viota.m v16, v0            # Get destination offsets of active elements
    add a6, a6, a5             # Accumulate number of elements
    vsll.vi v16, v16, 2, v0.t   # Multiply offsets by four bytes
    slli a5, a5, 2             # Multiply number of non-zero elements by four bytes
    vsuxei32.v v8, (a2), v16, v0.t # Scatter using scaled viota results under mask
    add a2, a2, a5             # Bump output pointer
    bnez a0, loop              # Any more?

    mv a0, a6                 # Return count
    ret

```

16.9. Vector Element Index Instruction

The `vid.v` instruction writes each element's index to the destination vector register group, from 0 to `v1-1`.

```
vid.v vd, vm # Write element ID to destination.
```

The instruction can be masked.

The `vs2` field of the instruction must be set to `v0`, otherwise the encoding is *reserved*.

The result value is zero-extended to fill the destination element if SEW is wider than the result. If the result value would overflow the destination SEW, the least-significant SEW bits are retained.

Microarchitectures can implement `vid.v` instruction using the same datapath as `viota.m` but with an implicit set mask source.

17. Vector Permutation Instructions

A range of permutation instructions are provided to move elements around within the vector registers.

17.1. Integer Scalar Move Instructions

The integer scalar read/write instructions transfer a single value between a scalar *x* register and element 0 of a vector register. The instructions ignore LMUL and vector register groups.

```
vmv.x.s rd, vs2 # x[rd] = vs2[0] (rs1=0)
vmv.s.x vd, rs1 # vd[0] = x[rs1] (vs2=0)
```

The `vmv.x.s` instruction copies a single SEW-wide element from index 0 of the source vector register to a destination integer register. If $SEW > XLEN$, the least-significant $XLEN$ bits are transferred and the upper $SEW - XLEN$ bits are ignored. If $SEW < XLEN$, the value is sign-extended to $XLEN$ bits.

The `vmv.s.x` instruction copies the scalar integer register to element 0 of the destination vector register. If $SEW < XLEN$, the least-significant bits are copied and the upper $XLEN - SEW$ bits are ignored. If $SEW > XLEN$, the value is sign-extended to SEW bits. The other elements in the destination vector register ($0 < \text{index} < VLEN/SEW$) are unchanged. If $vstart \geq vl$, no operation is performed and the destination register is not updated.

As a consequence, when $vl=0$, no elements are updated in the destination vector register group, regardless of $vstart$.

The encodings corresponding to the masked versions ($vm=0$) of `vmv.x.s` and `vmv.s.x` are reserved.

17.2. Floating-Point Scalar Move Instructions

The floating-point scalar read/write instructions transfer a single value between a scalar *f* register and element 0 of a vector register. The instructions ignore LMUL and vector register groups.

```
vfmv.f.s rd, vs2 # f[rd] = vs2[0] (rs1=0)
vfmv.s.f vd, rs1 # vd[0] = f[rs1] (vs2=0)
```

The `vfmv.f.s` instruction copies a single SEW-wide element from index 0 of the source vector register to a destination scalar floating-point register.

The `vfmv.s.f` instruction copies the scalar floating-point register to element 0 of the destination vector register. The other elements in the destination vector register ($0 < \text{index} < VLEN/SEW$) are unchanged. If $vstart \geq vl$, no operation is performed and the destination register is not updated.

As a consequence, when $vl=0$, no elements are updated in the destination vector register group, regardless of $vstart$.

The encodings corresponding to the masked versions ($vm=0$) of `vfmv.f.s` and `vfmv.s.f` are reserved.

17.3. Vector Slide Instructions

The slide instructions move elements up and down a vector register group.

The slide operations can be implemented much more efficiently than using the arbitrary register gather instruction. Implementations may optimize certain OFFSET values for `vslideup` and `vslidedown`. In particular, power-of-2 offsets may operate substantially faster than other offsets.

For all of the `vslideup`, `vslidedown`, `v[f]slide1up`, and `v[f]slide1down` instructions, if $vstart \geq vl$, the instruction performs no operation and leaves the destination vector register unchanged.

As a consequence, when $vl=0$, no elements are updated in the destination vector register group, regardless of $vstart$.

The slide instructions may be masked, with mask element *i* controlling whether *destination* element *i* is written.

17.3.1. Vector Slideup Instructions

```
vslideup.vx vd, vs2, rs1, vm      # vd[i+rs1] = vs2[i]
vslideup.vi vd, vs2, uimm[4:0], vm # vd[i+uimm] = vs2[i]
```

For `vslideup`, the value in `v1` specifies the maximum number of destination elements that are written. The start index (*OFFSET*) for the destination can be either specified using an unsigned integer in the x register specified by `rs1`, or a 5-bit immediate, zero-extended to XLEN bits. If $XLEN > SEW$, *OFFSET* is *not* truncated to SEW bits. Destination elements *OFFSET* through `v1-1` are written if unmasked and if *OFFSET* < `v1`.

`vslideup` behavior for destination elements

OFFSET is amount to slideup, either from x register or a 5-bit immediate

$0 < i < \max(vstart, OFFSET)$	Unchanged
$\max(vstart, OFFSET) \leq i < v1$	$vd[i] = vs2[i-OFFSET]$ if <code>v0.mask[i]</code> enabled, unchanged if not
$v1 \leq i < VLMAX$	Tail elements, unchanged

The destination vector register group for `vslideup` cannot overlap the source vector register group, otherwise the instruction encoding is reserved.

The non-overlap constraint avoids WAR hazards on the input vectors during execution, and enables restart with non-zero `vstart`.

17.3.2. Vector Slidedown Instructions

```
vslidedown.vx vd, vs2, rs1, vm      # vd[i] = vs2[i+rs1]
vslidedown.vi vd, vs2, uimm[4:0], vm # vd[i] = vs2[i+uimm]
```

For `vslidedown`, the value in `v1` specifies the number of destination elements that are written.

The start index (*OFFSET*) for the source can be either specified using an unsigned integer in the x register specified by `rs1`, or a 5-bit immediate, zero-extended to XLEN bits. If $XLEN > SEW$, *OFFSET* is *not* truncated to SEW bits.

<code>vslidedown</code> behavior for source elements	for element <code>i</code> in slide
$0 \leq i+OFFSET < VLMAX$	Read <code>vs2[i+OFFSET]</code>
$VLMAX \leq i+OFFSET$	Read as 0

<code>vslidedown</code> behavior for destination element <code>i</code> in slide	
$0 < i < vstart$	Unchanged
$vstart \leq i < v1$	Updated if <code>v0.mask[i]</code> enabled, unchanged if not
$v1 \leq i < VLMAX$	Unchanged

17.3.3. Vector Slide1up

Variants of slide are provided that only move by one element but which also allow a scalar integer value to be inserted at the vacated element position.

```
vslide1up.vx vd, vs2, rs1, vm      # vd[0]=x[rs1], vd[i+1] = vs2[i]
vfslide1up.vf vd, vs2, rs1, vm      # vd[0]=f[rs1], vd[i+1] = vs2[i]
```

The `vslide1up` instruction places the x register argument at location 0 of the destination vector register group, provided that element 0 is active, otherwise the destination element is unchanged. If $XLEN < SEW$, the value is sign-extended to SEW bits. If $XLEN > SEW$, the least-significant bits are copied over and the high $SEW-XLEN$ bits are ignored.

The remaining active `v1-1` elements are copied over from index `i` in the source vector register group to index `i+1` in the destination vector register group.

The `v1` register specifies how many of the destination vector register elements are written with source values, and all tail elements are unchanged.

`vslide1up` behavior

$i < \text{vstart}$	unchanged
$0 = i = \text{vstart}$	$\text{vd}[i] = \text{x}[\text{rs1}]$ if $\text{v0.mask}[i]$ enabled, unchanged if not
$\text{max}(\text{vstart}, 1) \leq i < \text{v1}$	$\text{vd}[i] = \text{vs2}[i-1]$ if $\text{v0.mask}[i]$ enabled, unchanged if not
$\text{v1} \leq i < \text{VLMAX}$	unchanged

The `vslide1up` instruction requires that the destination vector register group does not overlap the source vector register group. Otherwise, the instruction encoding is reserved.

The `vslide1up` instruction is defined analogously, but sources its scalar argument from an `f` register.

17.3.4. Vector Slide1down Instruction

The `vslide1down` instruction copies the first `v1-1` active elements values from index `i+1` in the source vector register group to index `i` in the destination vector register group.

The `v1` register specifies how many of the destination vector register elements are written with source values, and all tail elements are unchanged.

<code>vslide1down.vx</code>	<code>vd, vs2, rs1, vm</code>	<code># vd[i] = vs2[i+1], vd[v1-1]=x[rs1]</code>
<code>vslide1down.vf</code>	<code>vd, vs2, rs1, vm</code>	<code># vd[i] = vs2[i+1], vd[v1-1]=f[rs1]</code>

The `vslide1down` instruction places the `x` register argument at location `v1-1` in the destination vector register, provided that element `v1-1` is active, otherwise the destination element is unchanged. If `XLEN < SEW`, the value is sign-extended to `SEW` bits. If `XLEN > SEW`, the least-significant bits are copied over and the high `SEW-XLEN` bits are ignored.

`vslide1down` behavior

$i < \text{vstart}$	unchanged
$\text{vstart} \leq i < \text{v1}-1$	$\text{vd}[i] = \text{vs2}[i+1]$ if $\text{v0.mask}[i]$ enabled, unchanged if not
$\text{vstart} \leq i = \text{v1}-1$	$\text{vd}[\text{v1}-1] = \text{x}[\text{rs1}]$ if $\text{v0.mask}[i]$ enabled, unchanged if not
$\text{v1} \leq i < \text{VLMAX}$	unchanged

The `vslide1down` instruction is defined analogously, but sources its scalar argument from an `f` register.

The `vslide1down` instruction can be used to load values into a vector register without using memory and without disturbing other vector registers. This provides a path for debuggers to modify the contents of a vector register, albeit slowly, with multiple repeated `vslide1down` invocations.

17.4. Vector Register Gather Instructions

The vector register gather instructions read elements from a first source vector register group at locations given by a second source vector register group. The index values in the second vector are treated as unsigned integers. The source vector can be read at any index `< VLMAX` regardless of `v1`. The number of elements to write to the destination register is given by `v1`, and the remaining elements past `v1` are handled according to the current tail policy (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)). The operation can be masked.

<code>vrgather.vv</code>	<code>vd, vs2, vs1, vm</code>	<code># vd[i] = (vs1[i] >= VLMAX) ? 0 : vs2[vs1[i]];</code>
<code>vrgatherei16.vv</code>	<code>vd, vs2, vs1, vm</code>	<code># vd[i] = (vs1[i] >= VLMAX) ? 0 : vs2[vs1[i]];</code>

The `vrgather.vv` form uses `SEW/LMUL` for both the data and indices. The `vrgatherei16.vv` form uses `SEW/LMUL` for the data in `vs2` but `EEW=16` and `EMUL = (16/SEW)*LMUL` for the indices in `vs1`.

When `SEW=8`, `vrgather.vv` can only reference vector elements 0-255. The `vrgatherei16` form can index 64K elements, and can also be used to reduce the register capacity needed to hold indices for wider `SEW` data.

If the element indices are out of range ($vs1[i] \geq VLMAX$) then zero is returned for the element value.

Vector-scalar and vector-immediate forms of the register gather are also provided. These read one element from the source vector at the given index, and write this value to the `v1` elements at the start of the destination vector register. The index value in the scalar register and the immediate, zero-extended to `XLEN` bits, are treated as unsigned integers. If `XLEN > SEW`, the index value is *not* truncated to `SEW` bits.

These forms allow any vector element to be "splatted" to an entire vector.

```
vrgather.vx vd, vs2, rs1, vm # vd[i] = (x[rs1] >= VLMAX) ? 0 : vs2[x[rs1]]
vrgather.vi vd, vs2, uimm, vm # vd[i] = (uimm >= VLMAX) ? 0 : vs2[uimm]
```

For any `vrgather` instruction, the destination vector register group cannot overlap with the source vector register groups, otherwise the instruction encoding is reserved.

17.5. Vector Compress Instruction

The vector compress instruction allows elements selected by a vector mask register from a source vector register group to be packed into contiguous elements at the start of the destination vector register group.

```
vcompress.vm vd, vs2, vs1 # Compress into vd elements of vs2 where vs1 is enabled
```

The vector mask register specified by `vs1` indicates which of the first `v1` elements of vector register group `vs2` should be extracted and packed into contiguous elements at the beginning of vector register `vd`. The remaining elements of `vd` are treated as tail elements according to the current tail policy (Section [Vector Tail Agnostic and Vector Mask Agnostic vta and vma](#)).

Example use of `vcompress` instruction

```
1 1 0 1 0 0 1 0 1   v0
8 7 6 5 4 3 2 1 0   v1
1 2 3 4 5 6 7 8 9   v2

                                vcompress.vm v2, v1, v0
1 2 3 4 8 7 5 2 0   v2
```

`vcompress` is encoded as an unmasked instruction (`vm=1`). The equivalent masked instruction (`vm=0`) is reserved.

The destination vector register group cannot overlap the source vector register group or the source mask register, otherwise the instruction encoding is reserved.

A trap on a `vcompress` instruction is always reported with a `vstart` of 0. Executing a `vcompress` instruction with a non-zero `vstart` raises an illegal instruction exception.

Although possible, `vcompress` is one of the more difficult instructions to restart with a non-zero `vstart`, so assumption is implementations will choose not to do that but will instead restart from element 0. This does mean elements in destination register after `vstart` will already have been updated.

17.5.1. Synthesizing `vdecompress`

There is no inverse `vdecompress` provided, as this operation can be readily synthesized using `iota` and a masked `vrgather`:

```
Desired functionality of 'vdecompress'
7 6 5 4 3 2 1 0   # vid

    e d c b a     # packed vector of 5 elements
1 0 0 1 1 1 0 1   # mask vector of 8 elements

e 0 0 d c b 0 a   # result of vdecompress
```

```

1 0 0 1 1 1 0 1    # mask vector
4 4 4 3 2 1 1 0    # viota.m
0 0 0 0 0 0 0 0    # zero result register
e 0 0 d c b 0 a    # vrgather using viota.m under mask

```

```

# v0 holds mask
# v1 holds packed data
# v11 holds decompressed data
viota.m v10, v0      # Calc iota from mask in v0
vmv.v.i v11, 0       # Clear destination
vrgather.vv v11, v1, v10, v0.t # Expand into destination

```

The destination zeroing can be omitted if the decompressed elements are overwriting elements in an existing vector.

17.6. Whole Vector Register Move

The `vmv<nr>r.v` instructions copy whole vector registers (i.e., all VLEN bits) and can copy whole vector register groups. The instructions operate as if `EEW=SEW`, `EMUL = nr`, effective length `evl= EMUL * VLEN/SEW`.

These instructions are intended to aid compilers to shuffle vector registers without needing to know or change `v1` or `vtype`.

The usual property that no elements are written if `vstart ≥ v1` does not apply to these instructions. Instead, no elements are written if `vstart ≥ evl`.

If `vd` is equal to `vs2` the instruction is an architectural NOP, but is treated as a HINT to implementations that rearrange data internally that the register group will next be accessed with an `EEW` equal to `SEW`.

The instruction is encoded as an OPIVI instruction. The number of vector registers to copy is encoded in the low three bits of the `simm` field using the same encoding as the `nf` field for memory instructions, i.e., `simm = nr-1`. The value of the `nr` field must be 1, 2, 4, or 8, with other values reserved.

A future extension may support other numbers of registers to be moved. Values of `simm` other than 0, 1, 3, and 7 are currently reserved.

The instruction uses the same `funct6` encoding as the `vsmul` instruction but with an immediate operand, and only the unmasked version (`vm=1`). This encoding is chosen as it is close to the related `vmerge` encoding, and it is unlikely the `vsmul` instruction would benefit from an immediate form.

```

vmv<nr>r.v vd, vs2  # General form

vmv1r.v v1, v2      # Copy v1=v2
vmv2r.v v10, v12    # Copy v10=v12; v11=v13
vmv4r.v v4, v8       # Copy v4=v8; v5=v9; v6=v10; v7=v11
vmv8r.v v0, v8       # Copy v0=v8; v1=v9; ...; v7=v15

```

The source and destination vector register numbers must be aligned appropriately for the vector register group size, and encodings with other vector register numbers are reserved.

A future extension may relax the vector register alignment restrictions.

18. Exception Handling

On a trap during a vector instruction (caused by either a synchronous exception or an asynchronous interrupt), the existing `*epc` CSR is written with a pointer to the errant vector instruction, while the `vstart` CSR contains the element index that caused the trap to be taken.

We chose to add a `vstart` CSR to allow resumption of a partially executed vector instruction to reduce interrupt latencies and to simplify forward-progress guarantees. This is similar to the scheme in the IBM 3090 vector facility. To ensure forward progress without the `vstart` CSR, implementations would have to guarantee an entire vector instruction can always complete atomically without generating a trap. This is particularly difficult to ensure in the presence of strided or scatter/gather operations and demand-paged virtual memory.

18.1. Precise vector traps

We assume most supervisor-mode environments with demand-paging will require precise vector traps.

Precise vector traps require that:

1. all instructions older than the trapping vector instruction have committed their results
2. no instructions newer than the trapping vector instruction have altered architectural state
3. any operations within the trapping vector instruction affecting result elements preceding the index in the `vstart` CSR have committed their results
4. no operations within the trapping vector instruction affecting elements at or following the `vstart` CSR have altered architectural state except if restarting and completing the affected vector instruction will recover the correct state.

We relax the last requirement to allow elements following `vstart` to have been updated at the time the trap is reported, provided that re-executing the instruction from the given `vstart` will correctly overwrite those elements.

In idempotent memory regions, vector store instructions may have updated elements in memory past the element causing a synchronous trap. Non-idempotent memory regions must not have been updated for indices equal to or greater than the element that caused a synchronous trap during a vector store instruction.

Except where noted above, vector instructions are allowed to overwrite their inputs, and so in most cases, the vector instruction restart must be from the `vstart` location. However, there are a number of cases where this overwrite is prohibited to enable execution of the vector instructions to be idempotent and hence restartable from any location.

18.2. Imprecise vector traps

Imprecise vector traps are traps that are not precise. In particular, instructions newer than `*epc` may have committed results, and instructions older than `*epc` may have not completed execution. Imprecise traps are primarily intended to be used in situations where reporting an error and terminating execution is the appropriate response.

A profile might specify that interrupts are precise while other traps are imprecise. We assume many embedded implementations will generate only imprecise traps for vector instructions on fatal errors, as they will not require resumable traps.

Imprecise traps shall report the faulting element in `vstart` for traps caused by synchronous vector exceptions.

18.3. Selectable precise/imprecise traps

Some profiles may choose to provide a privileged mode bit to select between precise and imprecise vector traps. Imprecise mode would run at high-performance but possibly make it difficult to discern error causes, while precise mode would run more slowly, but support debugging of errors albeit with a possibility of not experiencing the same errors as in imprecise mode.

18.4. Swappable traps

Another trap mode can support swappable state in the vector unit, where on a trap, special instructions can save and restore the vector unit microarchitectural state, to allow execution to continue correctly around imprecise traps.

This mechanism is not defined in the base vector ISA.

A future extension might define a standard way of saving and restoring opaque microarchitectural state from a vector unit implementation to support context switching with imprecise traps.

19. Vector Instruction Listing

Integer					Integer				FP			
funct3					funct3				funct3			
OPIVV	V				OPMVV	V			OPFVV	V		
OPIVX		X			OPMVX		X		OPFVF		F	
OPIVI			I									

funct6					funct6				funct6			
000000	V	X	I	vadd	000000	V		vredsum	000000	V	F	vfadd
000001					000001	V		vredand	000001	V		vfredsum
000010	V	X		vsub	000010	V		vredor	000010	V	F	vfsub
000011		X	I	vrsb	000011	V		vredxor	000011	V		vfredsum
000100	V	X		vminu	000100	V		vredminu	000100	V	F	vfmin
000101	V	X		vmin	000101	V		vredmin	000101	V		vfredmin
000110	V	X		vmaxu	000110	V		vredmaxu	000110	V	F	vfmax
000111	V	X		vmax	000111	V		vredmax	000111	V		vfredmax
001000					001000	V	X	vaaddu	001000	V	F	vfsgnj
001001	V	X	I	vand	001001	V	X	vaadd	001001	V	F	vfsgnjn
001010	V	X	I	vor	001010	V	X	vasubu	001010	V	F	vfsgnjx
001011	V	X	I	vxor	001011	V	X	vasub	001011			
001100	V	X	I	vrgather	001100				001100			
001101					001101				001101			
001110		X	I	vslideup	001110		X	vslide1up	001110		F	vfslide1up
001110	V			vrgatherei16								
001111		X	I	vslidedown	001111		X	vslide1down	001111		F	vfslide1down

funct6					funct6				funct6			
010000	V	X	I	vadc	010000	V		VWXUNARY0	010000	V		VWFUNARY0
					010000		X	VRXUNARY0	010000		F	VRFUNARY0
010001	V	X	I	vmadc	010001				010001			
010010	V	X		vsbc	010010	V		VXUNARY0	010010	V		VFUNARY0
010011	V	X		vmsbc	010011				010011	V		VFUNARY1
010100					010100	V		VMUNARY0	010100			
010101					010101				010101			
010110					010110				010110			
010111	V	X	I	vmerge/vmv	010111	V		vcompress	010111		F	vfmerge.vf/vfmv
011000	V	X	I	vmseq	011000	V		vmandnot	011000	V	F	vmfeq
011001	V	X	I	vmsne	011001	V		vmand	011001	V	F	vmfle
011010	V	X		vmsltu	011010	V		vmor	011010			
011011	V	X		vmslt	011011	V		vmxor	011011	V	F	vmflt
011100	V	X	I	vmsleu	011100	V		vmornot	011100	V	F	vmfne
011101	V	X	I	vmsle	011101	V		vmnand	011101		F	vmfgt
011110		X	I	vmsgtu	011110	V		vmnor	011110			
011111		X	I	vmsgt	011111	V		vmxnor	011111		F	vmfge

funct6					funct6					funct6				
100000	V	X	I	vsaddu	100000	V	X	vdivu		100000	V	F	vfddiv	
100001	V	X	I	vsadd	100001	V	X	vdiv		100001		F	vfrdiv	
100010	V	X		vssubu	100010	V	X	vremu		100010				
100011	V	X		vssub	100011	V	X	vrem		100011				
100100					100100	V	X	vmulhu	100100	V	F	vfmul		
100101	V	X	I	vsll	100101	V	X	vmul	100101					
100110					100110	V	X	vmulhsu	100110					
100111	V	X		vsmul	100111	V	X	vmulh	100111		F	vfrsub		
			I	vmv<nf>r										
101000	V	X	I	vsrl	101000				101000	V	F	vfmadd		
101001	V	X	I	vsra	101001	V	X	vmadd	101001	V	F	vfnmadd		
101010	V	X	I	vssrl	101010				101010	V	F	vfmsub		
101011	V	X	I	vssra	101011	V	X	vnmsub	101011	V	F	vfnmsub		
101100	V	X	I	vnsrl	101100				101100	V	F	vfmacc		
101101	V	X	I	vnsra	101101	V	X	vmacc	101101	V	F	vfnmacc		
101110	V	X	I	vnclipu	101110				101110	V	F	vfmsac		
101111	V	X	I	vnclip	101111	V	X	vnmsac	101111	V	F	vfnmsac		

funct6					funct6					funct6				
110000	V			vwredsumu	110000	V	X	vwaddu		110000	V	F	vfwadd	
110001	V			vwredsum	110001	V	X	vwadd		110001	V		vfwredsum	
110010					110010	V	X	vwsubu		110010	V	F	vfwsub	
110011					110011	V	X	vwsu		110011	V		vfwredosum	
110100					110100	V	X	vwaddu.w		110100	V	F	vfwadd.w	
110101					110101	V	X	vwadd.w		110101				
110110					110110	V	X	vwsu.w		110110	V	F	vfwsub.w	
110111					110111	V	X	vwsu.w		110111				
111000					111000	V	X	vwmulu		111000	V	F	vfwmul	
111001					111001					111001				
111010					111010	V	X	vwmulsu		111010				
111011					111011	V	X	vwmul		111011				
111100					111100	V	X	vwmaccu		111100	V	F	vfwmacc	
111101					111101	V	X	vwmacc		111101	V	F	vfwnmacc	
111110					111110		X	vwmaccus		111110	V	F	vfwmsac	
111111					111111	V	X	vwmaccsu		111111	V	F	vfwnmsac	

Table 19. VRXUNARY0 encoding space

vs2	
00000	vmv.s.x

Table 20. VWXUNARY0 encoding space

vs1	
00000	vmv.x.s
10000	vpopc
10001	vfirst

Table 21. VXUNARY0 encoding space

vs1	
00010	vzext.vf8
00011	vsext.vf8
00100	vzext.vf4
00101	vsext.vf4
00110	vzext.vf2
00111	vsext.vf2

Table 22. VRFUNARY0 encoding space

vs2	
00000	vfmv.s.f

Table 23. VWFUNARY0 encoding space

vs1	
00000	vfmv.f.s

Table 24. VFUNARY0 encoding space

vs1	name
single-width converts	
00000	vfcvt.xu.f.v
00001	vfcvt.x.f.v
00010	vfcvt.f.xu.v
00011	vfcvt.f.x.v
00110	vfcvt.rtz.xu.f.v
00111	vfcvt.rtz.x.f.v
widening converts	
01000	vfwcvt.xu.f.v
01001	vfwcvt.x.f.v
01010	vfwcvt.f.xu.v
01011	vfwcvt.f.x.v
01100	vfwcvt.f.f.v
01110	vfwcvt.rtz.xu.f.v
01111	vfwcvt.rtz.x.f.v
narrowing converts	
10000	vfncvt.xu.f.w
10001	vfncvt.x.f.w
10010	vfncvt.f.xu.w
10011	vfncvt.f.x.w
10100	vfncvt.f.f.w
10101	vfncvt.rod.f.f.w
10110	vfncvt.rtz.xu.f.w
10111	vfncvt.rtz.x.f.w

Table 25. VFUNARY1 encoding space

vs1	name
00000	vfsqrt.v
00100	vfrsqste7.v
00101	vfrc7.v
10000	vfclass.v

Table 26. VMUNARY0 encoding space

vs1	name
00001	vmsbf
00010	vmsof
00011	vmsif
10000	viota
10001	vid

Appendix A: Vector Assembly Code Examples

The following are provided as non-normative text to help explain the vector ISA.

A.1. Vector-vector add example

link:example/vvaddint32.s[]

A.2. Example with mixed-width mask and compute.

```
# Code using one width for predicate and different width for masked
# compute.
#  int8_t a[]; int32_t b[], c[];
#  for (i=0; i<n; i++) { b[i] = (a[i] < 5) ? c[i] : 1; }
#
# Mixed-width code that keeps SEW/LMUL=8
loop:
    vsetvli a4, a0, e8,m1,ta,ma    # Byte vector for predicate calc
    vle8.v v1, (a1)                # Load a[i]
    add a1, a1, a4                  # Bump pointer.
    vmslt.vi v0, v1, 5              # a[i] < 5?

    vsetvli x0, a0, e32,m4,ta,mu    # Vector of 32-bit values.
    sub a0, a0, a4                  # Decrement count
    vmv.v.i v4, 1                   # Splat immediate to destination
    vle32.v v4, (a3), v0.t          # Load requested elements of C, others undisturbed
    sll t1, a4, 2                   # Bump pointer.
    add a3, a3, t1
    vse32.v v4, (a2)                # Store b[i].
    add a2, a2, t1                  # Bump pointer.
    bnez a0, loop                   # Any more?
```

A.3. Memcpy example

link:example/memcpy.s[]

A.4. Conditional example

```
# (int16) z[i] = ((int8) x[i] < 5) ? (int16) a[i] : (int16) b[i];
#
loop:
    vsetvli t0, a0, e8,m1,ta,ma    # Use 8b elements.
    vle8.v v0, (a1)                # Get x[i]
    sub a0, a0, t0                  # Decrement element count
    add a1, a1, t0                  # x[i] Bump pointer
    vmslt.vi v0, v0, 5              # Set mask in v0
    vsetvli t0, a0, e16,m2,ta,mu    # Use 16b elements.
    slli t0, t0, 1                  # Multiply by 2 bytes
    vle16.v v1, (a2), v0.t          # z[i] = a[i] case
    vlnot.m v0, v0                  # Invert v0
    add a2, a2, t0                  # a[i] bump pointer
    vle16.v v1, (a3), v0.t          # z[i] = b[i] case
    add a3, a3, t0                  # b[i] bump pointer
    vse16.v v1, (a4)                # Store z
    add a4, a4, t0                  # z[i] bump pointer
    bnez a0, loop
```

A.5. SAXPY example

`link:example/saxpy.s[]`

A.6. SGEMM example

`link:example/sgemm.S[]`

A.7. Division approximation example

`# v1 = v1 / v2 to almost 23 bits of precision.`

```
vfrece7.v v3, v2      # Estimate 1/v2
    li t0, 0x40000000
vmv.v.x v4, t0         # Splat 2.0
vfnmsac.vv v4, v2, v3  # 2.0 - v2 * est(1/v2)
vfmul.vv v3, v3, v4    # Better estimate of 1/v2
vmv.v.x v4, t0         # Splat 2.0
vfnmsac.vv v4, v2, v3  # 2.0 - v2 * est(1/v2)
vfmul.vv v3, v3, v4    # Better estimate of 1/v2
vfmul.vv v1, v1, v3    # Estimate of v1/v2
```

A.8. Square root approximation example

`# v1 = sqrt(v1) to almost 23 bits of precision.`

```
    fmv.w.x ft0, x0     # Mask off zero inputs
vmfne.vf v0, v1, ft0   # to avoid div by zero
vfrsq7e7.v v2, v1, v0.t # Estimate 1/sqrt(x)
vmfne.vf v0, v2, ft0, v0.t # Additionally mask off +inf inputs
    li t0, 0xbf000000
    fmv.w.x ft0, t0     # -0.5
vfmul.vf v3, v1, ft0, v0.t # -0.5 * x
vfmul.vv v4, v2, v2, v0.t # est * est
    li t0, 0x3fc00000
vmv.v.x v5, t0, v0.t   # Splat 1.5
vfmadd.vv v4, v3, v5, v0.t # 1.5 - 0.5 * x * est * est
vfmul.vv v1, v1, v4, v0.t # estimate to 14 bits
vfmul.vv v4, v1, v1, v0.t # est * est
vfmadd.vv v4, v3, v5, v0.t # 1.5 - 0.5 * x * est * est
vfmul.vv v1, v1, v4, v0.t # estimate to 23 bits
```

Appendix B: Calling Convention

In the RISC-V psABI, the vector registers `v0-v31` are all caller-saved. The `v1` and `vtype` CSRs are also caller-saved.

Procedures may assume that `vstart` is zero upon entry. Procedures may assume that `vstart` is zero upon return from a procedure call.

Application software should normally not write `vstart` explicitly. Any procedure that does explicitly write `vstart` to a nonzero value must zero `vstart` before either returning or calling another procedure.

The `vxrm` and `vxsat` fields of `vcsr` have thread storage duration.

Executing a system call causes all caller-saved vector registers (`v0-v31`, `v1`, `vtype`) and `vstart` to become unspecified.

This scheme allows system calls that cause context switches to avoid saving and later restoring the vector registers.

Most OSes will choose to either leave these registers intact or reset them to their initial state to avoid leaking information across process boundaries.

Appendix C: Vector Quad-Widening Integer Multiply-Add Instructions (Extension Zvqmac)

This is only a proposal for a future extension after v1.0 and might change substantially before ratification.

The quad-widening integer multiply-add instructions add a SEW-bit*SEW-bit multiply result to (from) a 4*SEW-bit value and produce a 4*SEW-bit result. All combinations of signed and unsigned multiply operands are supported.

These instructions are currently not planned to be part of the base V extension.

On ELEN=32 machines, only $8b * 8b = 16b$ products accumulated in a 32b accumulator would be supported. Machines with ELEN=64 would also add $16b * 16b = 32b$ products accumulated in 64b.

```
# Quad-widening unsigned-integer multiply-add, overwrite addend
vqmaccu.vv vd, vs1, vs2, vm    # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vqmaccu.vx vd, rs1, vs2, vm    # vd[i] = +(x[rs1] * vs2[i]) + vd[i]

# Quad-widening signed-integer multiply-add, overwrite addend
vqmacc.vv vd, vs1, vs2, vm     # vd[i] = +(vs1[i] * vs2[i]) + vd[i]
vqmacc.vx vd, rs1, vs2, vm     # vd[i] = +(x[rs1] * vs2[i]) + vd[i]

# Quad-widening signed-unsigned-integer multiply-add, overwrite addend
vqmaccsu.vv vd, vs1, vs2, vm   # vd[i] = +(signed(vs1[i]) * unsigned(vs2[i])) + vd[i]
vqmaccsu.vx vd, rs1, vs2, vm   # vd[i] = +(signed(x[rs1]) * unsigned(vs2[i])) + vd[i]

# Quad-widening unsigned-signed-integer multiply-add, overwrite addend
vqmaccus.vx vd, rs1, vs2, vm   # vd[i] = +(unsigned(x[rs1]) * signed(vs2[i])) + vd[i]
```

Appendix D: Divided Element Extension (Extension Zvediv)

The EDIV extension is currently not planned to be part of the base "V" extension, and will change substantially from this current sketch.

This section has not been updated to account for new mask format in v0.9.

The divided element extension allows each element to be treated as a packed sub-vector of narrower elements. This provides efficient support for some forms of narrow-width and mixed-width arithmetic, and also to allow outer-loop vectorization of short vector and matrix operations. In addition to modifying the behavior of some existing instructions, a few new instructions are provided to operate on vectors when $EDIV > 1$.

The divided element extension adds a two-bit field, `vediv[1:0]` to the `vtype` register.

Table 27. `vtype` register layout

Bits	Name	Description
XLEN-1	vill	Illegal value if set
XLEN-2:10		Reserved (write 0)
9:8	vediv[1:0]	Used by EDIV extension
7	vma	Mask agnostic
6	vta	Tail agnostic
5:3	vsew[2:0]	Selected element width (SEW) setting
2:0	vlmul[2:0]	Vector register group multiplier (LMUL) setting

The `vediv` field encodes the number of ways, *EDIV*, into which each SEW-bit element is subdivided into equal sub-elements. A vector register group is now considered to hold a vector of sub-vectors.

vediv [1:0]			Division EDIV
0	0	1	(undivided, as in base)
0	1	2	two equal sub-elements
1	0	4	four equal sub-elements
1	1	8	eight equal sub-elements

The assembly syntax for `vsetvli` has additional options added to encode the EDIV options.

```
d1  # EDIV 1, assumed if d setting absent
d2  # EDIV 2
d4  # EDIV 4
d8  # EDIV 8
```

```
vsetvli t0, a0, e32,m2,d4  # SEW=32, LMUL=2, EDIV=4
```


SEW	EDIV	Sub-element	Integer accumulator		FP sum/dot accumulator		
			sum	dot	FLEN=32	FLEN=64	FLEN=128
8b	2	4b	8b	8b	-	-	-
8b	4	2b	8b	8b	-	-	-
8b	8	1b	8b	8b	-	-	-
16b	2	8b	16b	16b	-	-	-
16b	4	4b	8b	16b	-	-	-
16b	8	2b	8b	8b	-	-	-
32b	2	16b	32b	32b	32b	32b	32b
32b	4	8b	16b	32b	-	-	-
32b	8	4b	8b	16b	-	-	-
64b	2	32b	64b	64b	32b	64b	64b
64b	4	16b	32b	64b	32b	32b	32b
64b	8	8b	16b	32b	-	-	-
128b	2	64b	128b	128b	32b	64b	128b
128b	4	32b	64b	128b	32b	64b	64b
128b	8	16b	32b	64b	32b	32b	32b
256b	2	128b	256b	256b	32b	64b	128b
256b	4	64b	128b	256b	32b	64b	128b
256b	8	32b	64b	128b	32b	64b	64b

Each implementation defines a minimum size for a sub-element, *SELEN*, which must be at most 8 bits.

While *SELEN* is a fourth implementation-specific parameter, values smaller than 8 would be considered an additional extension.

D.1. Instructions not affected by EDIV

The vector start register *vstart* and exception reporting continue to work as before.

The vector length *v1* control and vector masking continue to operate at the element level.

Vector masking continues to operate at the element level, so sub-elements cannot be individually masked.

SEW can be changed dynamically to enabled per-element masking for sub-elements of 8 bits and greater.

Vector load/store and AMO instructions are unaffected by EDIV, and continue to move whole elements.

Vector mask logical operations are unchanged by EDIV setting, and continue to operate on vector registers containing element masks.

Vector mask population count (*vpopc*), find-first and related instructions (*vfirst*, *vmsbf*, *vmsif*, *vmsof*), *iota* (*viota*), and element index (*vid*) instructions are unaffected by EDIV.

Vector integer bit insert/extract, and integer and floating-point scalar move instruction are unaffected by EDIV.

Vector slide-up/slide-down are unaffected by EDIV.

Vector compress instructions are unaffected by EDIV.

D.2. Instructions Affected by EDIV

D.2.1. Regular Vector Arithmetic Instructions under EDIV

Most vector arithmetic operations are modified to operate on the individual sub-elements, so effective SEW is SEW/EDIV and effective vector length is $v1 * EDIV$. For example, a vector add of 32-bit elements with a *v1* of 5 and EDIV of 4, operates identically to a vector add of 8-bit elements with a vector length of 20.

```

vsetvli t0, a0, e32,m1,d4 # Vectors of 32-bit elements, divided into byte sub-elements
vadd.vv v1,v2,v3           # Performs a vector of 4*v1 8-bit additions.
vsll.vx v1,v2,x1           # Performs a vector of 4*v1 8-bit shifts.

```

D.2.2. Vector Add with Carry/Subtract with Borrow Reserved under EDIV>1

For EDIV > 1, vadc, vmadc, vsbc, vmsbc are reserved.

D.2.3. Vector Reduction Instructions under EDIV

Vector single-width integer sum reduction instructions are reserved under EDIV>1. Other vector single-width reductions and vector widening integer sum reduction instructions now operate independently on all elements in a vector, reducing sub-element values within an element to an element-wide result.

The scalar input is taken from the least-significant bits of the second operand, with the number of bits equal to the number of significant result bits (i.e., for sum and dot reductions, the number of bits are given in table above, for non-sum and non-dot reductions, equal to the element size).

```

# Sum each sub-vector of four bytes into a 16-bit result.
vsetvli t0, a0, e32,d4 # Vectors of 32-bit elements, divided into byte sub-elements
vwredsum.vs v1, v2, v3 # v1[i][15:0] = v2[i][31:24] + v2[i][23:16]
                        #               + v2[i][15:8] + v2[i][7:0] + v3[i][15:0]

# Find maximum among sub-elements
vredmax.vs v5, v6, v7 # v5[i][7:0] = max(v6[i][31:24], v6[i][23:16],
                        #               v6[i][15:8], v6[i][7:0], v7[i][7:0])

```

Integer sub-element non-sum reductions produce a final result that is max(8,SEW/EDIV) bits wide, sign- or zero-extended to full SEW if necessary.

Integer sub-element widening sum reductions produce a final result that is max(8,min(SEW,2*SEW/EDIV)) bits wide, sign- or zero-extended to full SEW if necessary.

Single-width floating-point reductions produce a final result that is SEW/EDIV bits wide.

Widening floating-point sum reductions produce a final result that is min(2*SEW/EDIV,FLEN) bits wide, NaN-boxed to the full SEW width if necessary.

D.2.4. Vector Register Gather Instructions under EDIV

Vector register gather instructions under non-zero EDIV only gather sub-elements within the element. The source and index values are interpreted as relative to the enclosing element only. Index values \geq EDIV write a zero value into the result sub-element.

```

|       |       | SEW = 32b, EDIV=4
7 6 5 4 3 2 1 0 bytes
d e a d b e e f v1
0 1 9 2 0 2 3 2 v2
                        vrgather.vv v3, v1, v2
d a 0 e f e b e v3
                        vrgather.vi v4, v1, 1
a a a a e e e e v4

```

Vector register gathers with scalar or immediate arguments can "splat" values across sub-elements within an element.

Implementations can provide fast implementations of register gathers constrained within a single element width.

D.3. Vector Integer Dot-Product Instruction

The integer dot-product reduction `vdot.vv` performs an element-wise multiplication between the source sub-elements then accumulates the results into the destination vector element. Note the assembler syntax uses a `.vv` suffix since both inputs are vectors of elements.

Sub-element integer dot reductions produce a final result that is $\max(8, \min(\text{SEW}, 4 * \text{SEW} / \text{EDIV}))$ bits wide, sign- or zero-extended to full SEW if necessary.

```
# Unsigned dot-product
vdotu.vv vd, vs2, vs1, vm # Vector-vector

# Signed dot-product
vdot.vv vd, vs2, vs1, vm # Vector-vector
```

```
# Dot product, SEW=32, EDIV=1
vdot.vv vd, vs2, vs1, vm # vd[i][31:0] += vs2[i][31:0] * vs1[i][31:0]

# Dot product, SEW=32, EDIV=2
vdot.vv vd, vs2, vs1, vm # vd[i][31:0] += vs2[i][31:16] * vs1[i][31:16]
                                + vs2[i][15:0] * vs1[i][15:0]

# Dot product, SEW=32, EDIV=4
vdot.vv vd, vs2, vs1, vm # vd[i][31:0] += vs2[i][31:24] * vs1[i][31:24]
                                + vs2[i][23:16] * vs1[i][23:16]
                                + vs2[i][15:8] * vs1[i][15:8]
                                + vs2[i][7:0] * vs1[i][7:0]
```

D.4. Vector Floating-Point Dot Product Instruction

The floating-point dot-product reduction `vfdot.vv` performs an element-wise multiplication between the source sub-elements then accumulates the results into the destination vector element. Note the assembler syntax uses a `.vv` suffix since both inputs are vectors of elements.

```
# Signed dot-product
vfdot.vv vd, vs2, vs1, vm # Vector-vector
```

```
# Dot product. SEW=32, EDIV=2
vfdot.vv vd, vs2, vs1, vm # vd[i][31:0] += vs2[i][31:16] * vs1[i][31:16]
                                + vs2[i][15:0] * vs1[i][15:0]

# Floating-point sub-vectors of two half-precision floats packed into 32-bit elements.
vsetvli t0, a0, e32,m1,d2 # Vectors of 32-bit elements, divided into 16b sub-elements
vfdot.vv v1, v2, v3 # v1[i][31:0] += v2[i][31:16]*v3[i][31:16] + v2[i][16:0]*v3[i][16:0]

# Floating-point sub-vectors of four half-precision floats packed into 64-bit elements.
vsetvli t0, a0, e64,m1,d4 # Vectors of 64-bit elements, divided into 16b sub-elements
vfdot.vv v1, v2, v3
# v1[i][31:0] += v2[i][31:16]*v3[i][31:16] + v2[i][16:0]*v3[i][16:0] +
#               v2[i][63:48]*v3[i][63:48] + v2[i][47:32]*v3[i][47:32];
# v1[i][63:32] = ~0 (NaN boxing)
```
